

Approval Without Amplification: Gender and Visibility in Algorithmically Mediated Entrepreneurship

Katie Apker
Cornell University, ILR School

Draft – September 28, 2026

Abstract

Can negative evaluations help entrepreneurs? Research commonly emphasizes that favorable stakeholder evaluations help ventures secure critical resources, including visibility, and that women’s relatively unfavorable evaluations contribute to persistent gender gaps in outcomes. Yet negative, intense reactions may also drive visibility on engagement-optimizing platforms, where algorithmic recommendation systems allocate attention based on audience engagement rather than approval. This raises questions about how gendered evaluation relates to entrepreneurial outcomes in such systems. I analyze approximately 9,000 randomly sampled U.S. YouTube creator channels, combining platform metrics, 5.7 million audience comments, creator gender, and content measures. Regression, mediation, and decomposition analyses show that negative and profane audience engagement are associated with greater visibility, what I term a “negative arousal bonus.” Women receive more positive and less profane comments and approximately 20% fewer views than comparable men in channels of the same age, video count, and content category, with audiences’ comment valence and profanity accounting for the majority of this gap. These findings highlight how even ostensibly positive evaluation can reproduce gender inequality wherever a valued resource is allocated based on the intensity of the evaluation rather than its favorability, processes that are especially relevant in increasingly common algorithmically mediated settings.

INTRODUCTION

Success within entrepreneurship is often understood to depend on positive stakeholder evaluations (Zuckerman, 1999; Hsu, 2006; Pollock and Rindova, 2003; Roccapriore and Pollock, 2023; Botelho, 2024). In extant research focusing on gendered evaluation processes, women’s

disadvantage is thus commonly explained in terms of their receiving less favorable evaluations than comparable men, with the understanding that more favorable evaluation helps entrepreneurs secure greater access to funding, endorsements, and other key resources, while less positive evaluations can lead to fewer opportunities and resources (Abraham, 2019; Botelho and Abraham, 2017; Eagly and Karau, 2002; Kanze et al., 2018; Lee and Huang, 2018; Snellman and Solal, 2023; Tonoyan et al., 2024). This dynamic communicates an assumed “favorability logic” of evaluative processes that are often used to help explain persistent disparities in entrepreneurial outcomes, which commonly fall along gender lines (Jennings and Brush, 2013).

With the proliferation of digital technologies, platforms, and media, entrepreneurs increasingly engage various audience, customer, partner, and investor stakeholders through online platforms (Nambisan, 2017; Pollock et al., 2023). In these growing settings, venture outcomes often depend on *visibility*—the volume of attention a venture attracts—as a key resource often observed through audience clicks and views (Cutolo and Kenney, 2021; Nambisan, 2017; Rindova et al., 2007; Pollock et al., 2023; Vanacker et al., 2020). This visibility is understood to accumulate as a form of attention-based capital that shapes reputation and can convert directly into revenue, sponsorships, deal flow, and other opportunities (Bhargava, 2022; Verwiebe et al., 2026; Franck, 2018; Petkova et al., 2013; Vanacker et al., 2020; Pillemer et al., 2025). Within the existing literature on audience evaluations, the common expectation is that visibility is dependent on favorable evaluation and that ventures and entrepreneurs who are evaluated less favorably, women among them, will receive less (Brooks et al., 2014; Kanze et al., 2018; Ewens and Townsend, 2020; Snellman and Solal, 2023; Tonoyan et al., 2024).

However, the digital platform settings of this increasing share of entrepreneurship are often mediated by analytic AI, with engagement-optimizing algorithms allocating venture visibility (Nambisan, 2017; Cutolo and Kenney, 2021; Covington et al., 2016; Huszár et al., 2021). Unlike funding or hiring, visibility is not handed over by an evaluator who must judge a venture favorably. Instead, it is commonly distributed algorithmically according to how much engagement a venture draws, favorable or not. This dynamic has long been true of press coverage and word of mouth

(Rindova et al., 2006; Pollock and Rindova, 2003; Kovács and Sharkey, 2014) and is increasingly the dominant structure on digital platforms, whose recommendation systems rank on engagement (Covington et al., 2016; Huszár et al., 2021). Furthermore, a growing literature on virality suggests the engagement that generates visibility is often *negative* rather than positive, with negative and emotionally intense reactions associated with more sharing and algorithmic amplification (Berger and Milkman, 2012; Brady et al., 2017; Milli et al., 2025). Intense reactions are understood to propagate further both through audience sharing and algorithmic amplification that prioritizes behavioral engagement signals such as clicks, watch time, and comment counts (Berger and Milkman, 2012; Covington et al., 2016; Huszár et al., 2021; Milli et al., 2025), communicating an “engagement logic” of evaluative processes. If attention and visibility are allocated this way, then what matters for digital entrepreneurs is not necessarily favorable evaluation per se, but engagement that marks their venture as consequential (e.g., content that audiences dwell on, respond to, and argue about), thereby supplying the kind of high-arousal reactions that engagement-optimizing systems are most likely to amplify. Despite this shift, we have limited research on how audience stakeholders’ evaluations of digital entrepreneurs relate to their visibility in these algorithmically-mediated, attention-driven contexts, and therefore little warrant to assume that favorable evaluation will garner entrepreneurs valuable attention in digital settings.

Furthermore, existing theoretical accounts of gendered evaluation have yet to examine how gender differences in evaluation may function in widespread entrepreneurial settings where visibility is allocated algorithmically through engagement rather than direct approval. Academic and industry evidence indicates that women tend to do less well in various settings where visibility is the key resource (Verwiebe et al., 2026; Döring and Mohseni, 2019; Collabstr, 2025), but a simple explanation for this pattern using existing accounts isn’t clear. While established favorability logic accounts might largely predict that women who receive less favorable evaluations will obtain fewer resources (Botelho and Abraham, 2017; Kanze et al., 2018; Snellman and Solal, 2023; Tonoyan et al., 2024), the engagement-based logic of algorithmic allocation would predict that women may garner more visibility if they indeed tend to receive less favorable audience reception. Additionally,

evidence from digital entrepreneurship extends this complexity, with some studies documenting hostility and sexist comments toward women, particularly in samples drawn from highly visible channels or targeted individuals (Wotanis and McMillan, 2014; Döring and Mohseni, 2019, 2020; Sobieraj, 2017), while others find less negativity or softer engagement toward women in broader or random samples (Thelwall et al., 2009; Weißmann et al., 2025, preprint). Because these patterns diverge, it remains unclear whether women in algorithmically mediated markets generally encounter harsher disapproval or softer, less critical engagement. Given that engagement-based allocation can reward intensity, either pattern may translate into systematically different visibility by gender. Taken together, both the direction of women's reception and its impact on their visibility remain open questions.

In light of these tensions, I therefore ask two related questions. First, in entrepreneurial contexts where visibility matters, how does the nature of audience evaluation and engagement relate to entrepreneurial visibility? Second, are these relationships gendered? i.e. do men and women receive systematically different patterns of audience reception, and do these differences help explain any observed gender gaps in entrepreneurial visibility? These questions are consequential both practically and theoretically. Practically, these forms of entrepreneurship are a large and growing domain where visibility often translates directly into revenue (Cutolo and Kenney, 2021; Rahman, 2021; Bhargava, 2022). Theoretically, resources allocated based on the engagement they elicit coexist alongside those resources that are allocated based on favorability within the same evaluative ecosystems, making it necessary to understand how gendered evaluation plays out when intense engagement, rather than approval, may be a critical gateway to entrepreneurial opportunity.

To address whether gendered patterns of audience reception relate to visibility, I draw on research examining gendered evaluation by audiences in entrepreneurship and organizational contexts (Abraham et al., 2024; Botelho and Abraham, 2017; Botelho and Poskanzer, 2026). This literature provides a theoretical foundation because it shows that the form gendered evaluation takes varies by setting, and under engagement-based allocation, each form implies distinct patterns of engagement with different implications for visibility. In high-stakes selection and funding

settings, women often face harsher evaluations, being seen as less competent or less suited to entrepreneurial roles than comparable men (Botelho and Abraham, 2017; Eagly and Karau, 2002; Snellman and Solal, 2023). By contrast, in more routine or relational settings, women may receive softer, subjectively positive treatment that withholds status, stretch opportunities, and candid feedback (Bareket and Fiske, 2023; Glick and Fiske, 2001; Heilman, 2012; Jampol and Zayas, 2020; King et al., 2010). Under the usual assumption that more favorable evaluation secures more resources, this softer treatment appears advantageous. However, where audiences and recommendation systems reward intense, critical engagement, softer treatment can reduce public conflict, harsh language, and moralized outrage—signals that mark work as consequential and attract the engagement these systems prioritize (Brady et al., 2017; Milli et al., 2025; Rathje et al., 2021). In such settings, apparently positive treatment may suppress the engagement that fuels visibility, potentially disadvantaging women in markets that place a premium on being seen.

I analyze a random sample of 9,096 established U.S. YouTube creator-entrepreneurs, combining YouTube Data API metrics obtained through the YouTube Researcher Program with hand-coded creator gender and other demographics, over 5.7 million audience comments sampled across channel uploads, and text-based measures of creator-produced content. Total views at the creator level serves as the main indicator of visibility. I operationalize audience evaluation using LIWC-22 comment tone and profanity and VADER sentiment measure, capturing valence and intensity in ways that are consistent with prior work on negativity and arousal in sharing and algorithmic reach (Berger and Milkman, 2012; Brady et al., 2017; Rathje et al., 2021; Milli et al., 2025; Lafreniere et al., 2022). My analysis proceeds in four phases. First, it examines the reception-to-visibility relationship both across creators and within individual creators' video content. Second, it measures the gender gap in visibility and examines what remains of it after accounting for different potentially driving factors, including channel tenure, production volume and polish, efforts to attract new viewers, content categories and topics, and comment volume, leveraging three waves of longitudinal data to help address reverse causality. Third, it tests whether comment tone and profanity mediate the gender gap and whether the reception-to-visibility association differs by creator

gender. Fourth, it evaluates alternative explanations and robustness, including the content women produce, concentration in low-visibility categories, audience willingness to interact, commenter gender composition, and sensitivity analyses.

My findings reveal several patterns in how evaluation and visibility relate. Creators receiving more negative and more profane comments tend to have higher views, both across creators and within individual creators' uploads, while more positive comment tone is associated with fewer views, consistent with what I term a "negative arousal bonus." Even when women have been on the platform as long and have produced as many videos as men, they receive about 19 percent fewer views at the same channel age and video count, and about 20 percent fewer when channels of the same content category are also compared. This gap is not accounted for by the content women produce, their concentration in lower-visibility categories, or less audience willingness to interact. Notably, women's comment sections show a softer reception than men's. The average positive comment tone on women's channels is about 63 on LIWC-22's 1-to-99 Tone scale, compared to about 53 on men's channels. Profanity, while relatively infrequent in comments overall, is about 120 percent more common in men's channels, consistent with benevolent sexism (Glick and Fiske, 2001; Jampol and Zayas, 2020). The visibility gap does not appear to arise from women's reception being amplified differently and Gelbach decomposition of the gender view gap (controlling for channel age, video count, and content category) shows comment tone accounts for 88.0 percent and profanity for 6.7 percent. Similarly, positive comment tone and profanity together account for approximately 93 percent of the gender gap in views at equal comment volume, channel age, video count, and content category in joint mediation tests. These patterns hold across a series of alternative explanations, robustness, and sensitivity checks.

My findings make several contributions. First, to the literature on audience evaluation, which has shown that favorable evaluation typically unlocks resources (Zuckerman, 1999; Hsu, 2006; Pollock and Rindova, 2003; Botelho, 2024), this study adds that a resource such as visibility can instead be garnered by drawing negative, intense audience reactions rather than a simply favorable evaluation. This has important implications for entrepreneurial settings and resource

allocation processes, but also for society more broadly, highlighting how engagement-optimizing recommendation systems may contribute to the polarization of public discourse (Rathje et al., 2021, 2024; Huszár et al., 2021; Robertson et al., 2023; Milli et al., 2025). Second, my findings contribute to research on gendered evaluation and benevolent sexism. Prior work has shown that women evaluated less favorably in high-stakes settings receive fewer resources, while the softer treatment women often receive in relational settings has been invoked more often in traditional organizational rather than entrepreneurial settings (Glick and Fiske, 2001; Jampol and Zayas, 2020; King et al., 2010). This study adds that the softer treatment itself can both occur and carry a material cost for entrepreneurs, and, notably, that the standard remedy of more favorable evaluation would likely not close this gap. Third, these findings contribute to the literature on algorithmic management, which has shown how algorithms control and evaluate opaquely and through constrained choices (Kellogg et al., 2020; Rahman, 2021; Curchod et al., 2019; Cameron, 2024; Rahman et al., 2024). This study adds that an algorithmically mediated setting using a gender-neutral engagement allocation rule can produce gendered outcomes when applied to gendered inputs, as observed with other seemingly neutral design features of evaluation processes (Botelho and Poskanzer, 2026; Castilla and Benard, 2010; Lambrecht and Tucker, 2019). Taken together, my study highlights that gendered audience reception and engagement-optimized visibility allocation together reproduce the gendered distribution of resources and rewards, as long observed in organizations (Acker, 1990), within a new technological context. This underscores how platform-mediated systems can reconfigure, rather than eliminate, familiar forms of entrepreneurial inequality.

THEORY AND HYPOTHESES

Stakeholder evaluation, gender, and visibility

Entrepreneurial ventures rely on external stakeholders to access critical resources such as investment, customers, and legitimacy. Favorable evaluations from these stakeholders are closely associated with unlocking these resources (Zuckerman, 1999; Zimmerman and Zeitz, 2002; Pollock

and Rindova, 2003; Hsu, 2006; Botelho, 2024). Empirical research documents that ventures' sensegiving activities attract specialized media attention, which precedes resource provision such as venture capital funding (Petkova et al., 2013). Media coverage not only reflects legitimacy but contributes to it, influencing investor behavior and market valuation in IPO contexts (Pollock and Rindova, 2003). Firms lacking coverage from securities analysts who specialize in their industry experience an illegitimacy discount, indicating that analyst attention shapes valuation rather than simply mirroring it (Zuckerman, 1999). These findings highlight the pivotal role of human intermediaries, journalists, analysts, and experts, in conferring attention, which carries consequential resource implications.

Research examining gender disparities in entrepreneurship builds directly on the premise that favorable evaluation is associated with resource acquisition. Women receive less favorable evaluations than comparable men from investors, judges, and employees, which correlates with their reduced access to resources (Brooks et al., 2014; Kanze et al., 2018; Lee and Huang, 2018; Abraham, 2019; Botelho and Abraham, 2017; Snellman and Solal, 2023; Kacperczyk et al., 2023). These studies share the underlying assumption that more favorable evaluation is associated with greater resource access, particularly when the evaluator also controls the allocation. Yet, attention as a resource does not fit this assumption cleanly.

Management research distinguishes the volume of attention a venture attracts from the valence or favorability of that attention (Pollock and Rindova, 2003; Rindova et al., 2007; Petkova et al., 2013; Vanacker et al., 2020; Graf-Vlachy et al., 2019). Attention can be drawn not by approval but by dramatic or nonconforming actions. For example, ventures engaging in unconventional behaviors receive more media coverage (Rindova et al., 2006), prizewinning books attract more readers despite declining ratings (Kovács and Sharkey, 2014), and negative reviews can increase sales of unknown products by raising awareness rather than signaling approval (Berger et al., 2010). This body of work generally treats attention as a precursor to favorable judgment by an evaluator who controls resources, effectively stopping at the point where attention is conferred (Petkova et al., 2013; Pollock and Rindova, 2003; Zuckerman, 1999). For a growing share of entrepreneurs,

however, attention does not feed a later judgment. Instead, attention itself becomes the return (Franck, 2018).

Visibility functions as a primary resource in contexts where audiences and advertisers directly fund entrepreneurial work. Platforms exemplify this condition, as creator income is directly indexed to views and advertising revenue sharing (Bhargava, 2022), with documented gender gaps in earnings linked to visibility (Verwiebe et al., 2026). In these settings, reach and esteem can diverge (Mears, 2023), and digital markets characterized by heavy-tailed distributions amplify early advantages, compounding disparities (Gala et al., 2024). The criteria by which visibility is allocated serve as a core distribution mechanism shaping resource access, rather than a mere technical detail. The criteria for allocating visibility matter greatly. Research on gendered evaluation has so far examined how evaluation is structured while keeping favorability as the criterion for resource allocation. Understanding how these criteria operate in practice is essential to advancing this work.

Recent scholarship has highlighted that features of the evaluation process itself beyond evaluator bias shape gendered outcomes. Examples include differential conversion of ratings into pay (Castilla, 2008), gendered language in assessments (Correll et al., 2020), funding rates obscuring unequal evaluations (Bol et al., 2022), escalating stakes across evaluation stages (Botelho and Poskanzer, 2026), and algorithmic delivery producing gendered exposure without explicit gender inputs (Lambrecht and Tucker, 2019; Milli et al., 2025). These studies show that structural features influence how consistently favorable evaluation translates into resource access, but favorability remains the underlying criterion. An additional feature that could add to this literature is an account of allocation where the criterion is not favorability but the volume and intensity of engagement. Engagement-based criteria represent a distinct feature of evaluation processes that may reconfigure how visibility and resource access are determined. This paper introduces engagement-based criteria as central to understanding entrepreneurial visibility on digital platforms.

From audience evaluation to expressed reception and engagement-based visibility

Private audience evaluation represents an internal judgment that cannot be directly observed. On digital platforms, this private evaluation becomes visible only through public expressions such as comments, likes, and replies. These expressions are produced in view of earlier engagement and shaped by social constraints, which often cause public expressions to differ from private judgments (Orlikowski and Scott, 2014; Karunakaran et al., 2022; Botelho, 2024; Bitektine and Haack, 2015). Commenters on platforms hold no formal authority over creators and do not supply capital, yet their public expressions carry meaningful consequences. Advertisers, other audience members, and platform ranking systems interpret these expressions as signals when allocating visibility (Curchod et al., 2019; Karunakaran et al., 2022; Roccapriore and Pollock, 2023; Seigner et al., 2023). These publicly observable expressions thus serve as a trace of audience evaluation that platforms and stakeholders rely on to assess entrepreneurial content.

Platforms design their allocation logic to optimize watch time and advertising revenue, making behavioral engagement signals such as clicks, replies, and expected watch time per impression the core criteria for allocating visibility (Covington et al., 2016; Bhargava, 2022; Milli et al., 2025). These criteria produce uneven visibility distributions by amplifying certain content relative to a chronological baseline (Huszár et al., 2021). Content producers experience these scoring mechanisms as opaque and adapt their behavior without full understanding of the criteria (Rahman, 2021; Orlikowski and Scott, 2014; Curchod et al., 2019). The platform system functions as an allocative context rather than an intentional evaluator. Consequently, content that generates more engagement per impression receives greater visibility regardless of whether it is positively evaluated by audiences. This implies that the allocation of attention on digital platforms depends more on the intensity and nature of audience engagement than on favorable evaluation alone.

Even when engagement-based criteria do not explicitly observe gender, they can still produce gendered outcomes if the inputs they read differ by gender. For instance, a gender-neutral advertisement was delivered more frequently to men than women because cost-based optimization crowded out women's impressions without any actor targeting gender (Lambrecht and Tucker,

2019). Similarly, the ranking objective influences how unevenly rewards fall across complementors (Zhang et al., 2026; Huszár et al., 2021). These findings suggest that gendered outcomes can emerge solely from engagement criteria if engagement signals differ systematically by gender. This possibility opens the door for gendered patterns to arise in attention allocation on platforms like YouTube despite the system's formal neutrality, making it essential to examine gender differences in audience engagement and visibility.

Negative information tends to weigh more heavily than positive information (Baumeister et al., 2001; Rozin and Royzman, 2001). Negative words in headlines increase click-through rates (Robertson et al., 2023), and out-group animosity generates more engagement than in-group warmth (Rathje et al., 2021). Recommendation systems amplify divisive and polarizing content beyond users' stated preferences (Milli et al., 2025). Beyond valence, high-arousal content is shared more than low-arousal content (Berger and Milkman, 2012), and moral-emotional language travels further through networks (Brady et al., 2017). These findings indicate that content eliciting negative and high-arousal engagement produces stronger behavioral signals that platforms prioritize.

On YouTube, a more negative and more profane comment section at a given comment volume signals more intense engagement, such as reply chains, repeat visits, and longer dwell times, which engagement-based criteria interpret and amplify. Comment tone captures the valence of expressed reception, while profanity captures arousal; these are separable dimensions of word meaning, with taboo words forming a distinct category (Warriner et al., 2013; Janschewitz, 2008). Intense negative or profane reception is thus expected to carry a visibility premium relative to positive reception alone. This visibility premium, termed a negative arousal bonus, arises because the platform's primary allocator of visibility, the recommendation system, responds to the intensity embedded in audience engagement. Audiences passing content along off-platform provide a secondary amplification channel. A rival interpretation posits that negativity may follow reach rather than produce it, as widely surfaced videos attract viewers who are not fans and comment more negatively. This rival reading is addressed in the study design by using first-wave reception to predict later views while holding first-wave views and comment volume constant.

By this logic, if the criteria read behavioral engagement signals, and if negative and profane comment sections carry more intense engagement at a given comment volume, then more negative and more profane expressed reception should be associated with greater visibility, and more positive, less profane reception with less.

Hypothesis 1a. More negative expressed audience reception (harsher, more critical comment tone) is positively associated with channel visibility (views).

Hypothesis 1b. More profane expressed audience reception is positively associated with channel visibility, consistent with profanity capturing high-arousal engagement.

Gendered audience evaluation and the softer register

Gendered disadvantage in evaluation and resource allocation appears in multiple forms shaped by the stakes and power dynamics of the setting. In high-stakes, accountable, and male-typed environments such as venture funding, promotion, and leadership selection, women are associated with less favorable evaluations or different evaluative standards than comparable men (Brooks et al., 2014; Kanze et al., 2018; Lee and Huang, 2018; Eagly and Karau, 2002; Snellman and Solal, 2023). These disparities include documented double standards in analysts' ratings and attention allocation (Botelho and Abraham, 2017). The gender gaps tend to emerge when evaluators face accountability to third parties and operate within male-typed domains but largely diminish in private assessments or female-typed occupations (Abraham, 2019). Beyond evaluation scores, gendered disadvantage also manifests in the translation of evaluations into tangible rewards. Women receive smaller salary increases despite equivalent performance ratings (Castilla, 2008), early favorable evaluations of female-led startups may fail to convert into resource allocation at later, higher-stakes stages (Botelho and Poskanzer, 2026), and the same board position yields different outcomes by gender (Frankort et al., 2026). This harsher form of evaluation reflects a pattern tied to the stakes and power cues in the setting.

Gendered evaluation also takes a contrasting form in lower-stakes, more relational contexts, where gendered treatment often appears as softer, more protective engagement. Benevolent sexism

preserves gender hierarchies through subjectively positive, protective, and paternalistic treatment and operates alongside hostile sexism as a coordinated system. Hostile sexism responds to power and competence cues with overt hostility, while benevolent sexism responds to role and relational cues with ostensibly positive treatment (Glick and Fiske, 2001; Bareket and Fiske, 2023; Heilman, 2012). The costs of benevolent sexism are well documented. Women receive fewer challenging developmental assignments (King et al., 2010), inflated and less candid feedback (Jampol and Zayas, 2020), and board chairs demonstrate softer orientations toward newly appointed female CEOs (Oliver et al., 2017). These patterns do not stem from women's preferences, as women desire candid and critical feedback as much as men do (Coffman and Klinowski, 2025). Both hostile and benevolent sexism represent forms of disadvantage that vary depending on the stakes and power cues present. Whether commenters, mostly strangers to the creator, carry gendered treatment norms into public comment sections and which form predominates there remains an open question.

Status beliefs prescribe not only how women should behave but also how others should treat them (Ridgeway, 2011). Benevolent treatment of women is a prescriptive norm across cultures, prescribing more supportive and less confrontational interaction toward women than men (Glick and Fiske, 2001; Jampol and Zayas, 2020). Norms for men permit a more assertive and contentious register (Eagly, 1987; Ridgeway, 2001). These norms operate through two channels simultaneously. First, personal stereotypes shape individual responses to gender cues (Fiske et al., 2002; Heilman, 2012). Second, shared status beliefs create coordination pressure even among those without personal bias, as people orient to perceived expectations of others (Correll, 2017; Ridgeway, 2011). Both channels produce similar observable patterns of gendered treatment.

Gendered treatment norms extend into digital markets designed to remove face-to-face contact, such as peer-to-peer lending (Kuwabara and Thébaud, 2017) and online labor markets (Galperin, 2019). On YouTube, the activating conditions for these norms are present. Creator gender is legible through profile images and video content. Commenters are not anonymous: usernames, profile images, and comment histories are visible, and comments are posted publicly in view of other viewers and earlier comments (Karunakaran et al., 2022; Botelho, 2024). This combination of

legible target gender and partially accountable commenter identity creates conditions under which offline treatment norms operate similarly online. If these conditions hold, online evidence should show both hostile and benevolent forms of gendered treatment, contingent on the cues each setting conveys.

Research on gendered audience reception online reveals a complex mix of hostile and softer responses toward women. Hostile findings show that women creators receive more sexist and appearance-focused comments, especially in male-typed contexts such as STEM channels, where female-hosted channels accumulate more comments per view and a higher proportion of critical and hostile remarks than male-hosted channels (Wotanis and McMillan, 2014; Amarasekara and Grant, 2018). Incivility and slurs disproportionately target women, functioning to police and constrain their visibility (Sobieraj, 2017; Döring and Mohseni, 2020, Study 1). Conversely, other research finds softer reception. Women creators on YouTube receive less negative communication than men (Weißmann et al., 2025, preprint), and hostility effects are small in a second study (Döring and Mohseni, 2020, Study 2). Women-led crowdfunding projects more often reach funding goals (Greenberg and Mollick, 2016), and women receive more replies to identical pitches (Gornall and Strebulaev, 2025). This pattern suggests that hostile reception arises in contexts marked by policing, visibility struggles, and male-typed domains, representing backlash and penalties for violating behavioral norms (Rudman and Glick, 2001; Rudman et al., 2012; Heilman, 2012). Softer reception appears in more routine, relational online interactions. This mirrors the stakes-and-power pattern of gendered evaluation described earlier.

Commenting on a creator's video is a routine, relational, public act with low formal stakes and no evaluative power (Karunakaran et al., 2022; Botelho, 2024; Pillemer et al., 2025). This cue profile aligns with the conditions benevolent sexism responds to (Bareket and Fiske, 2023; Botelho and Poskanzer, 2026; Abraham, 2019). Benevolent sexism produces inequity without depressing evaluations of women (Nguyen et al., 2023). Across a random sample of creators spanning many content categories, the central tendency of audience reception in comment sections should be a softer register. If commenting is a low-stakes relational act performed publicly before an audience,

and if benevolent sexism responds to those cues, then women's comment sections should carry a softer register, more positive in tone and less profane. The rival hypothesis is that hostile sexism extends into this setting, so women's comment sections would be more negative and more profane. Rejecting hypotheses H2a and H2b would support that rival.

Hypothesis 2a. Women's channels receive more positive comment tone than men's channels.

Hypothesis 2b. Women's channels receive less profane commentary than men's channels.

Under criteria that reward favorability, a softer expressed reception appears advantageous because the evaluator who is favorable typically controls access to the resource. In contrast, under engagement-based criteria, the same softer register is less charged and less conflictual. It thus generates fewer of the behavioral signals that these criteria interpret and amplify, such as high-arousal engagement. The value of expressed reception depends on the type of resource being allocated and the mechanism translating reception into resource access.

Inequality in visibility involves two linked components: whether women receive favorable reception, and whether that reception converts into visibility under engagement-based criteria. Women can occupy the favorable side of the first component, receiving more positive and less profane comments, while simultaneously being disadvantaged in the second, because the criteria prioritize intensity rather than positivity. This dynamic means that even when the criteria themselves do not explicitly incorporate gender as an input, gendered outcomes can emerge. Since the criteria read engagement intensity equivalently across creators, differences in visibility arise from differences in the input the creators receive rather than from differential weighting of gender. These patterns together imply the following: if the criteria read engagement intensity with no gender input, and if softer reception carries less intensity at a given comment volume, then comment tone and profanity should account for part of the gender gap in visibility.

Hypothesis 3. Comment tone and profanity statistically mediate the association between creator gender and visibility, net of production and content controls.

DATA AND METHODS

Research Setting

YouTube provides a setting well suited to examining the relationship between audience evaluation and entrepreneurial visibility because creator-entrepreneurs operate channels competing for attention and monetization. Visibility on this platform is allocated primarily through engagement signals rather than centralized evaluator approval. Recommendation and search systems interpret behavioral indicators such as click-through rates, watch time, and comment activity, while audiences contribute further by clicking, sharing, and commenting. The most recent public technical description of YouTube’s recommendation system (Covington et al., 2016) specifies optimization for expected watch time per impression using implicit feedback and hundreds of partially disclosed features; I assume this system continued to rely on behavioral engagement signals during the 2025 to 2026 study period. The economic stakes are real: creators receive 55 percent of ad revenue on monetized videos (Verwiebe et al., 2026), and views translate into advertising revenue, sponsorship opportunities, and discoverability. The sample frame’s 2,000-subscriber floor exceeds YouTube’s Partner Program eligibility threshold, ensuring that all sampled creators can convert views into income. Thus, the empirical requirements are met: both expressed evaluation, observable through comment text, and visibility, measured by views, coexist publicly within the same engagement-based allocation system.

Commenting on YouTube is a public, post-viewing behavior. Viewers access content via search, subscription feeds, external links, or recommendations, but the majority do not comment, with passive viewing predominating (Khan, 2017; Nonnecke and Preece, 2000). Therefore, comments represent a visible behavioral trace of expressed reception rather than a comprehensive census of audience evaluation; these comments are legible to other viewers and accessible to platform systems. This empirical distinction aligns with the theoretical separation between private evaluation and expressed reception. Creator gender is socially legible in this context, as individual creators are identifiable through profile images, channel names, creator content, biographical information, and

linguistic cues (Ridgeway, 2008), a necessary condition for gendered audience reception. Publicly available channel-level data include cumulative views, subscriber counts, video counts, channel age, comment activity, video titles, descriptions, transcripts when available, and thumbnails, making the outcome variable, proposed mediators, and a broad set of creator-controlled content features observable from public sources. The design addresses potential confounds through three waves of data collection for temporal separation, a within-creator video-level panel that controls for stable creator traits, controls for channel and comment volume, content-category fixed effects, creator-controlled content elements as rival mediators, and sensitivity analyses for unmeasured confounding.

Sample

The sample frame was defined using Infludata, a commercial influencer-analytics platform accessed under an academic license. Six criteria delimit the frame: the platform was restricted to YouTube; language set to English; geographic location within the United States; presence of a male or female gender assignment in Infludata's database used solely to screen for individual rather than brand accounts; at least one post within 90 days prior to data collection; and at least 2,000 subscribers. Within this frame, the draw was random with no additional design elements. Infludata delivered a list of 14,169 channels meeting these criteria. Channels were kept when they belonged to an individual creator with an identifiable woman or man gender presentation, yielding 9,232 channels: 3,186 women creators (34.5 percent) and 6,046 men creators (65.5 percent). The largest exclusion was channels that did not belong to an individual creator, such as brand, group, or organizational accounts. The analytic sample comprises 9,096 channels with views, channel age, video count, and scored comment language. The 136 sampled channels outside this analytic sample (1.47 percent of 9,232) are no more common among women's channels than men's (1.54 versus 1.44 percent, $p = .71$; Appendix Table A28). Table 1 reports the analytic sample (Panel A) and compares women's and men's channels (Panel B), with arithmetic means, geometric means, and geometric means at the same channel age and video count. The sample represents established, individually identifiable creators indexed by Infludata and does not constitute a census of YouTube. Very new

creators, those below the 2,000-subscriber minimum, non-individual accounts, and creators with illegible gender are outside the scope of this study.

Creator gender was hand-coded from multiple public profile sources including profile images, channel names, creator content, biographical information, and linguistic cues. Coding was verified by a research assistant, achieving an interrater reliability of $\kappa = 0.96$. Infludata's gender field served only as a screen for individual creators during sampling and was not used as an independent variable or coding source.

< INSERT TABLE 1 HERE >

Data

Data were collected using the YouTube Data API v3 under the YouTube Researcher Program, which provided an expanded quota for retrieval. Channel metadata fields included cumulative channel views, subscriber counts, video counts, channel creation dates, video-quality indicators, topic classifications, and branding settings. Data collection occurred in three waves verified by retrieval timestamps: the first wave (T1) from May 20 to June 2, 2025; the second wave (T2) on February 2, 2026, approximately 8.3 months after T1; and the third wave (T3) on June 18, 2026, roughly twelve to thirteen months after T1. At T3, 9,114 of the 9,232 channels were successfully retrieved; 118 were dead or private.

The comment corpus was constructed by first retrieving each channel's uploads playlist. Available videos were randomly shuffled, and up to 200 comments per video were collected until either 1,000 comments were reached or videos were exhausted. Videos with comments disabled were skipped. Comment retrieval was time-ordered, so the sampled comments represent a recent window of audience reception rather than initial reception. Comment and title samples derive from different video sets by design to provide broader coverage of audience response across the creator's body of work. The corpus comprises over 5.7 million comments across 9,100 channels and 451,875 videos, of which 5,638,823 comments were successfully scored by LIWC-22; 144,985 emoji-only comments contained no word count. Using video identifiers in the raw comment pulls,

supplemented in September 2026 by a per-video pull of up to 100 comments on each sampled video, a video-level panel was constructed including 135,148 videos across 7,157 channels (2,512 women creators and 4,645 men creators), applying floors of at least ten comments per video and at least two qualifying videos per channel. Videos not returned by the API because of deletion, privacy settings, or region blocking since collection are absent, so the panel is a selected sample and is used only for the within-creator analyses.

Creator-controlled content elements were collected across four domains: video titles, video descriptions, transcripts where captions were available, and thumbnails. Up to twelve videos per channel were sampled for titles, including the three oldest, three newest, three most popular, and up to three sponsored videos where identifiable. Thumbnails were collected for the newest-video cohort, which provides full coverage and is not outcome-selected; the most-popular cohort is outcome-selected and used descriptively only. LIWC-22 scoring was applied to comment text and to textual content elements including titles, descriptions, transcripts, and thumbnail overlay text extracted by optical character recognition. Commenter gender was inferred from display names using two methods: first, gender-word detection over a fixed dictionary of honorifics, familial terms, identity words, titles, and slang with strong-versus-weak word criteria; second, matching names against the U.S. Social Security Administration historical names file applying minimum-length, low-frequency, and ambiguity cutoffs. The liberal coding variant covers roughly 77 percent of commenters and is used in the reported results.

The next section details how these data sources operationalize the theoretical constructs and measurement strategies.

Measures

Visibility is operationalized as the natural logarithm of cumulative channel views. This log transformation addresses the heavy right skew in the distribution, with the standard deviation of logged views at 2.124, corresponding to roughly an eightfold difference per standard deviation. The measure aggregates views across all public videos on a creator's channel, representing a portfolio-

level visibility outcome. Due to limitations of the public YouTube Data API, it is not possible to decompose views by video format, such as Shorts versus long-form content. Views reflect a combination of content exposure and viewer choice, whereas impressions and click-through rates remain proprietary metrics accessible only to channel owners and are therefore unobserved in this study. This measure serves as the empirical counterpart to the theoretical construct of visibility volume as a resource and corresponds to the “total views at the creator level” described in the introduction. Creator gender is coded as a binary indicator, with one assigned to women creators and zero to men creators, based on the detailed hand coding procedure described earlier. Having established visibility as the outcome, the next step is to define the key theoretical inputs capturing audience response, namely expressed reception in its dual dimensions.

Expressed reception is captured through two mediators derived from language analysis of audience comments. Both comment tone and comment profanity are measured at the individual comment level using LIWC-22 and then averaged to the channel level. Comment tone is operationalized via the LIWC-22 Tone summary variable, which reflects net valence with higher values indicating more positive tone. Each comment’s Tone runs from 1 to 99; a comment with no positive or negative tone words (or with equal amounts of each) scores 20.23, LIWC-22’s own value for that case, and enters the channel mean at that value. Comments with no words at all (emoji only; 144,985 comments) are excluded. Fifty is not a neutral point on this scale. Comment profanity is measured using the LIWC-22 Swear category and is treated as an indicator of emotional arousal or intensity rather than simple negative valence. This distinction rests on the observation that taboo words appear in both positive and negative contexts and convey emotional activation beyond polarity. Because valence and arousal are recognized as distinct semantic dimensions (Warriner et al., 2013), and taboo words constitute a separate lexical category from general valenced words (Janschewitz, 2008), tone and profanity are conceptualized as separable features of expressed reception. Hypothesis H1a concerns the role of tone, while H1b addresses profanity. Both variables are length-normalized word proportions by design. Models include logged total comment volume as a control to isolate the character of reception (its positivity or negativity) rather than the quan-

tity of comments. A robustness check further incorporates mean comment length to account for comment verbosity. Having defined these constructs, the next paragraph addresses the validity and comparability of the measurement approach across data sources and levels. LIWC-22 was selected for its transparent dictionary-based methodology and extensive validation record (Tausczik and Pennebaker, 2009; Boyd et al., 2022), as well as its established use in leading management journals (Malhotra et al., 2017; Gamache et al., 2015). A key advantage is the ability to apply the same dictionary uniformly across audience comments and creator-controlled content elements, facilitating direct comparability. It is important to clarify that LIWC summary variables are not Likert scales; construct validity is assessed through convergence with VADER, a lexicon-based sentiment analysis tool used independently to cross-validate tone measures. The channel-level LIWC-22 Tone and the channel-level VADER compound score, each averaged over the channel's comments, correlate at $r = 0.87$ ($N = 9,096$; Appendix Table A19), supporting convergent validity. LIWC-22 leaves Tone blank when a text contains no positive or negative tone words. The sampled video titles of 719 channels and the video descriptions of 239 channels contain no tone words, and 105 channels have no description text. These are scored at LIWC-22's value of 20.23 (profanity 0 for empty descriptions), each case flagged by an indicator that every model with creator content includes, so title and description language is available for all 9,096 analytic channels (Appendix Table A02). The flags differ by gender: titles without tone words on 6.5 percent of women's and 8.7 percent of men's channels ($p < .001$), descriptions without tone words on 2.0 versus 3.0 percent ($p = .007$), and empty descriptions on 1.5 versus 1.0 percent ($p = .026$). The corpus holds 5,783,808 comments, of which the 5,638,823 that contain words were each scored with LIWC-22 in September 2026. The video panel carries the same per-comment LIWC-22 Tone and Swear measures as the channel level, averaged to the video and standardized. These are the comparable measures and Table 7 reports them (Models 3 and 4). The video panel is also scored with the VADER compound score per comment averaged to the video, and with a hand-built profanity token list that includes mild intensifiers (hell, damn), insults and ableist terms LIWC does not classify as profanity, and coarse or hostile lexical content. Table 7 reports these too (Models 1 and 2), as

a convergent check; the VADER pipeline reproduces the canonical per-comment VADER scores exactly ($r = 1.0000$). The two measure sets give the same within-creator result (tone -0.23 with either measure). With outcome, treatment, and mediators defined, the final measurement task is to specify the controls included to account for confounding factors and their conceptual rationale.

Every model includes two production controls. Channel tenure is captured by channel age in years, reflecting accumulated opportunity for visibility. Production volume is measured by the logged video count. Logged total comment volume, the natural log of the API comment counts on each channel's three newest and three oldest videos, is a post-outcome quantity and enters as its own step rather than as a background control, because women's channels receive more comments per view (Table 1, Panel B). Including comment volume enables models to compare channels at the same amount of reception and isolate the character of reception rather than its quantity. HD video (97 percent of both women's and men's channels; $p = .59$ for the difference) and the non-subscriber trailer (61.6 percent of men's channels, 55.7 percent of women's) are not in the main model progression; both enter the Gelbach decomposition (Table 6) and the rival tests (Table 8; Appendix Table A16). Content topic is controlled through 40 content-category fixed effects derived from YouTube's granular topic classifications, restricted to categories with at least 30 channels. Category sizes range from 32 to 865 channels, with a median of 126 in the 9,096-channel comment-language sample (Appendix Table A17, Panel E). Creator gender varies within every category.

Furthermore, secondary outcomes such as subscriber count and engagement rate are examined but deliberately excluded from the control set because engagement rate's numerator includes commenting and its denominator includes views, placing it on the pathway between gender and views rather than among pre-treatment covariates. Creator-controlled content elements, including video titles, descriptions, transcripts, and thumbnails, are measured using LIWC-22 tone and profanity metrics or visual features (for thumbnails). Titles and video descriptions enter the main model progression as a creator-content block, before comment volume (Table 5, Model 3); transcripts and thumbnails enter the appendix models as rival mediators (Appendix Table A04). Titles and descriptions cover every analytic channel once the no-tone-word cases are flagged. Transcripts

are missing for 1,014 of the 9,096 channels (11.5 percent of men's, 10.6 percent of women's; $p = .19$; Appendix Table A28): 372 channels have captions only in languages other than English, 447 have no captions, 157 have every sampled video unavailable, and 38 were not fetched. Thumbnail visual features are missing for 269 channels (2.9 versus 3.0 percent; $p = .87$) and thumbnail overlay text for 1,107 channels (12.5 versus 11.6 percent; $p = .26$). None of these losses taken one at a time differs by gender, but the 6,727 channels with all four content elements lose more men's than women's channels (27.4 versus 23.5 percent of each group; $p < .001$), which is why the four-element model sits in the appendix. Variable names for all controls are provided in the table notes. With all constructs mapped to measures, the following section details how these measures are integrated into the analytic models following the four-step progression outlined in the introduction.

Analytic Approach

The unit of analysis throughout this study is the creator channel. Estimation employs ordinary least squares with heteroskedasticity-robust standard errors following White (1980). Creators were sampled individually rather than in groups, and creator gender varies within every one of the 40 content categories, so robust standard errors are the main specification; standard errors clustered on the 40 content categories (Appendix Table A08, Model 2) and a wild cluster bootstrap (Appendix Table A24) are reported as robustness checks, and the production-adjusted and volume-adjusted gaps are significant at the .05 level under all three. The interpretive stance is consistent across analyses: all estimates represent associations between expressed reception measured through comment tone and profanity and visibility, operationalized as logged cumulative channel views. Temporal ordering from the lagged design and the within-creator video-level panel provide supporting evidence for reading these associations as directional from reception to visibility, but they do not establish causal effects. The first phase of data analysis examines the cross-channel relationship between reception and visibility by regressing logged views on standardized comment tone and profanity. This is done with the two production controls, then adding comment volume, then comment tone and profanity, and finally adding 40-category fixed effects to account for content-topic

heterogeneity. Creator gender is excluded to estimate the general association across all channels. Reference specifications with single mediators appear in Appendix Table A01. Complementing this, the within-creator design uses a video-level panel with channel fixed effects to compare videos within the same creator, thereby removing stable creator characteristics including gender from the reception-visibility association. This approach addresses the first research question concerning reception and visibility but does not address the gender gap or reverse ordering.

Having established the general reception-to-visibility association, the second phase of data analysis focuses on the gender gap in visibility and the extent to which it persists after accounting for obvious confounders. Every channel-level model is estimated on the 9,096 channels with complete comment-language measures, the analytic sample; the raw gap is estimated by regressing logged views on creator gender alone (Model 1). The production-adjusted gap adds channel age and video count, the production controls (Model 2); this is the headline gap, reported two ways: the production-adjusted gap (about 19 percent fewer views) and the same gap within the 40 content categories (about 20 percent fewer views), because category sorting can itself reflect discrimination, so holding category constant risks over-controlling (Blau and Kahn, 2017). Creator content enters next, adding the tone and profanity of the channel's own video titles and video descriptions (Model 3). Comment volume then enters as the first step of the reception block (Model 4), followed by comment profanity (Model 5) and comment tone (Model 6), in a no-fixed-effects panel and a 40-category fixed-effects panel (Table 5; Figure 4 plots the creator-gender coefficient across the models). The creator-gender coefficient moves from -0.413 (raw) to -0.216 when channel age and video count enter (Table 5). In the Gelbach decomposition, channel age and video count account for 56.4 percent of the raw gap (Table 6, column 1). The Gelbach decomposition (Gelbach, 2016) decomposes three starting gaps: the raw gap (column 1), the production-adjusted gap (column 2), and the production-adjusted gap within the 40 content categories (column 3). In columns 1 and 2, content category is its own block (39 category indicators, reported as one block total). The blocks are channel age and video count (column 1 only), content category, creator content, comment volume, comment profanity, comment tone, HD video, and the non-subscriber trailer. All three

columns end at the same full model with every block and the 40 category indicators. Each block's contribution is the creator-gender difference in that block multiplied by the block's coefficient in the full model, so it does not depend on the order of entry. Appendix Table A27 reports further starting points and versions with transcripts. To test for sorting into content categories, the test is now the content-category block of the Gelbach decomposition (Table 6, columns 1 and 2); the Oaxaca-Blinder decomposition (Appendix Table A03) corroborates it. Further refinement uses a 128-subtopic fixed-effect ladder to tighten topic comparison beyond the 40 categories (Appendix Table A20). Addressing reverse causality, a two-wave lagged-dependent-variable design predicts second-wave logged views from first-wave comment language, conditional on first-wave logged views and controls, approximately eight months apart (Appendix Table A05), with a third wave repeating this at twelve to thirteen months (Appendix Table A06). Attrition across waves is examined against creator gender (Appendix Table A07). These lagged models provide the sole support for temporal ordering from reception to later visibility. For inference on the headline gap, category-clustered standard errors (40 clusters) are reported in the appendix, supplemented by wild cluster bootstrap inference with Webb six-point weights and 9,999 replications to address the small cluster count (Appendix Table A24).

If the gender gap remains after adjusting for production, topic, and comment volume, the theory predicts that women's channels receive a softer expressed reception, an association that corresponds with lower visibility under engagement-based criteria. The third phase tests this mediation hypothesis. Hypotheses H2a and H2b are first tested descriptively and through regressions of comment tone and profanity on creator gender, controlling for channel and comment volume (Table 1, Panel B; Table 4; commenter-gender breakdown in Appendix Table A12). The analytic progression continues by adding comment profanity, then comment tone alongside profanity, to the model that already holds channel age, video count, creator content, and comment volume constant, in both panels (Table 5, Models 5 and 6; $N = 9,096$). The mediation analysis is estimated at equal comment volume, holding channel age, video count, and comment volume constant: the Gelbach decomposition assigns shares of the raw, production-adjusted, and within-category production-

adjusted gaps (Table 6); the mediation analysis tests the path from creator gender through comment tone and profanity to views for the volume-adjusted gap (channel age, video count, and comment volume held constant), with and without category fixed effects (Appendix Table A21). Under OLS with the same controls, each block's Gelbach share equals the product-of-coefficients indirect association of the mediation model. Mediation analysis follows Baron and Kenny (1986) logic, with bootstrap percentile intervals for indirect effects computed using 5,000 bootstrap samples (Preacher and Hayes, 2008). Indirect paths for tone and profanity are estimated separately and jointly, with and without category fixed effects (Appendix Table A21). A generalized structural equation model reproduces the no-fixed-effects indirect effects precisely. As a check rather than a hypothesis test, a creator-gender-by-comment-tone interaction at the channel level assesses whether the reception-to-visibility association differs by gender (Appendix Table A15). A parallel within-creator interaction is tested at the video level (Appendix Table A22). Mediation results are interpreted as statistical evidence regarding whether the gender-to-visibility association is carried by the proposed mediators, not as identified causal effects.

I also address rival explanations to bolster my mediated results. The fourth phase of data analysis systematically tests alternative explanations in the order presented in the introduction, summarizing each with its respective test. Table 8 provides an overview of rival explanations and their test outcomes, with full model details in the appendix. First, the possibility that gender differences in comment language reflect differences in creator content is tested by including the four creator-controlled content elements as control blocks and rival mediators in a parallel multiple-mediator model (Appendix Table A02 for titles and descriptions; Appendix Table A04 for all four elements combined). Titles and video descriptions also enter the main model progression as a block (Table 5, Model 3) and the Gelbach decomposition (Table 6). Second, concentration in low-visibility content categories is addressed through category fixed effects, subtopic fixed effects, and the Oaxaca-Blinder decomposition described above. Additionally, category gender-typing is examined via triple interactions of creator gender, comment tone, and female-typed categories (Appendix Table A15). Third, audience willingness to interact is controlled for throughout by

comment volume, with subscriber count and engagement rate analyzed as alternative outcomes (Table 1, Panel B; Table 8; Appendix Table A17, Panel F). Fourth, commenter gender composition is assessed by computing tone and profanity separately for male-coded and female-coded commenters and including the share of female-coded commenters as a covariate (Appendix Tables A12, A13, A14). Fifth, off-platform multi-homing is tested by including outbound links to external platforms as a covariate and interacting them with creator gender (Appendix Table A16). Sixth, unmeasured confounding is evaluated using the Cinelli and Hazlett (2019) sensitivity analysis framework for both the gender gap and mediator pathways (Appendix Table A10). Finally, a robustness battery addresses alternative outcomes, profanity transformations (Appendix Table A09), mean comment length, VADER sentiment replication, additional comment-language components, models with and without category fixed effects, category-clustered standard errors, trimming outcome extremes, restricting to channels active at least twelve months, tone-by-comment-volume interactions, quantile regression, exclusion of gaming categories, missing-data diagnostics (Appendix Table A18), and within-creator panel robustness (Appendix Table A22).

Completing the methods section, it is critical to delineate clearly what the design observes and what it does not, ensuring that results are interpreted within appropriate boundaries. Observed variables include expressed reception through comment tone and profanity derived from comment text, visibility as cumulative channel views, creator gender hand-coded from public profiles, and production and content characteristics captured by the channel and comment volume controls, content-category indicators, and the four creator-controlled content elements. Not observed are the audience's private judgments, which are mental states distinct from expressed reception; expressed reception itself is a combined trace of evaluation and engagement intensity that cannot be disentangled in comment data. The platform's internal allocation criteria are not observed; the design sees inputs (expressed reception) and outputs (views) but not the recommendation system's internal criteria. Additionally, the two amplification channels, audiences passing content along off-platform and the platform's recommendation systems, are not separately identifiable with these data. Moderated hostility, such as comments removed prior to observation through deletion, review

holds, or blocked words, is not tested and acknowledged as a limitation that may underrepresent hostility in the corpus. The lagged design provides the sole support for temporal ordering from reception to later visibility. The within-creator panel demonstrates that the association is not explained by stable creator traits. Public technical accounts describe YouTube’s recommendation system optimization during the study window (Covington et al., 2016). The interpretive frame is that evaluation shapes expressed reception in valence and arousal, and that recommendation systems alongside audiences allocate visibility based on the intensity of reception rather than its favorability. This paper documents the pattern linking these observable quantities and tests whether it differs by creator gender.

Descriptive Patterns

The channels were drawn at random from the Infludata frame of U.S. individual creators with at least 2,000 subscribers, with creator gender hand-coded (inter-rater agreement kappa = 0.96), and Table 1 describes the analytic sample of 9,096 channels with views, channel age, video count, and scored comment language, of whom 34.5 percent (3,137) are women creators (Table 1, Panel A). Cumulative channel views average 58.0 million and are heavily right-skewed. Comment tone, measured on the LIWC-22 1 to 99 scale where higher values indicate more positive comment tone, averages 56.2, which is above the 20.23 baseline assigned to comments without tone words, indicating that typical comment sections are mildly positive. Profanity appears in 0.35 percent of comment words on the average channel (Table 1, Panel A). To anchor this scale, Table 2 presents illustrative comments: the negative extreme includes phrases like “You suck SO BAD STOP” and “I hate your videos. You’re so dumb.”; the middle consists of neutral queries such as “What wig is this?”; and the positive extreme features comments like “AMAZING WORK! You inspire me so much.” Pairwise correlations foreshadow the paper’s central argument: positive comment tone correlates negatively with logged views ($r = -0.29$), profanity correlates positively with logged views ($r = 0.12$), being a woman creator correlates negatively with logged views ($r = -0.09$), positively with positive comment tone ($r = 0.31$), and negatively with profanity ($r = -0.21$); positive tone and

profanity themselves are negatively correlated ($r = -0.33$) (Appendix Table A23).

< INSERT TABLE 2 HERE >

Comparing women's and men's channels across focal and control measures reveals the core puzzle. Women's channels receive markedly softer comment sections, with mean positive tone at 62.8 versus 52.7 on men's channels, and profanity at 0.19 percent of comment words versus 0.43 percent (Table 1, Panel B). Further, 44.3 percent of men's channels but only 20.6 percent of women's channels have mean positive comment tone below the scale midpoint of 50 ($p < .0001$; Appendix Table A17, Panel B), locating the two distributions and not meaning that men's comment sections are net negative. Audience interaction is higher on women's channels, with engagement rates of 3.90 percent compared to 3.25 percent on men's channels, and total comments on the six sampled videos at a geometric mean of 67 versus 58 (Table 1, Panel B). Despite this, women's channels have fewer views, with views at a geometric mean of 2.76 million compared to 4.18 million for men, 34 percent fewer, and median views per subscriber at 122 versus 173 ($p < .0001$; Appendix Table A17, Panel B). Subscriber counts differ less markedly, with geometric means of 23,319 on women's channels versus 26,017 on men's, representing 10.4 percent fewer subscribers ($p = .005$). Controlling for channel age and video count (Table 1, Panel B), women's channels have 19.5 percent fewer views, their subscriber counts are not statistically distinguishable from men's ($p = .87$), and they exhibit higher engagement: 24.7 percent more likes plus comments per view and 14.2 percent more per subscriber on the six sampled videos (both $p < .001$). Men's channels are approximately 0.9 years older (10.0 versus 9.1 years), have more videos (geometric mean 589 versus 466 videos), and more frequently run trailers for non-subscribers (62 versus 56 percent; Table 1, Panel B). These disparities justify controlling for channel age and video count in subsequent models, with comment volume entered as its own step.

The Reception-to-Visibility Association Across Creators

The first phase tests hypotheses H1a and H1b by examining the association between audience reception and visibility across creators without including creator gender terms. Table 3 builds the

model up across 9,096 channels: the production controls, channel age and video count (Model 1); comment volume (Model 2); standardized positive comment tone and profanity (Model 3); and 40 content-category fixed effects (Model 4). With channel age, video count, and comment volume held constant, a one standard deviation increase in positive comment tone associates with a 0.38 decrease in logged views, while a one standard deviation increase in profanity associates with a 0.08 increase (Model 3; both $p < .001$). The comment volume control isolates the type of reception at a fixed volume of commentary. Adding 40 content-category fixed effects further strengthens the positive tone association to -0.43 and reduces profanity's to 0.05 (Model 4; $p < .001$ and $p = .008$, respectively). Profanity's association diminishes substantially once positive tone is included: entered without tone, profanity's coefficient is 0.18, but it falls by about 56 percent to 0.08 when positive tone enters (Table 5, Panel A, Models 5 and 6). These results support both H1a and H1b, confirming a “negative arousal bonus” at the channel level.

< INSERT TABLE 3 HERE >

To verify that the positive tone association is not an artefact of the language measure or a proxy for other comment features, nine additional LIWC-22 comment-language components moral, positive emotion, affect, social, polite, prosocial, anger, conflict, and swear were added one at a time and then collectively. The positive tone coefficient remains between -0.37 and -0.54 across these specifications, with the full model estimate at -0.54, all highly significant ($p < .001$; $N = 9,096$; Appendix Table A11). This confirms that positive comment tone is not simply capturing anger or conflict words. Figures 1 and 2 visualize the relationship. Figure 1 compares two Health category channels with similar comment volumes (2,000 and 2,300 comments). The channel with more negative comments (tone 39) has logged views of 18.9, whereas the channel with more positive comments (tone 81) has 17.7 logged views. Figure 2 plots mean logged views by decile of positive comment tone across all 9,096 channels, showing a monotonic decline from 16.02 in the most negative decile to 14.04 in the most positive.

< INSERT FIGURES 1 AND 2 HERE >

The Association Within Creators

The second phase tests H1a and H1b within creators by examining whether videos with harsher comment reception receive more views than videos with softer reception on the same channel. Table 7, Panel A, analyzes 135,148 videos across 7,157 channels, controlling for video age, its square, and log duration, with standard errors clustered at the channel level. Within a creator's uploads, a one standard deviation increase in positive comment tone associates with a 0.23 decrease in logged views ($p < .001$), while a one standard deviation increase in profanity associates with a 0.03 increase in logged views ($p < .001$; Model 2, VADER measures). Scoring the same panel with the LIWC-22 measures used at the channel level gives -0.23 for positive tone and 0.02 for profanity (Model 4; both $p < .001$). These signs align with the channel-level association. Pooled models without channel fixed effects yield a positive tone association of -0.43 and a profanity association of -0.02 ($p < .05$) (Model 1). The within-creator estimate is preferred as it isolates variation net of stable creator traits, supporting both hypotheses.

< INSERT TABLE 7 HERE >

To exclude the creator's own language as the source, Table 7, Panel B, assesses own video tone and profanity on logged views within creators (20,872 videos on 7,425 channels). Neither own tone nor own profanity shows significant associations within creators; confidence intervals include zero. Pooled across creators, own positive tone is negatively and own profanity positively associated with views, but neither survives within-creator comparison (Table 7, Model 5). This indicates visibility differences among a creator's videos correspond to audience comment language rather than the creator's own language. The within-creator positive tone association persists under several robustness checks, including a comment floor of 30, exclusion of Shorts, restriction to channels with at least 10 qualifying videos, and alternative tone measures (Appendix Table A22, columns 2 to 5). A video-level interaction of creator gender with positive comment tone, measured with LIWC-22 on the same panel, is not distinguishable from zero (0.000, $p = .997$; joint test of the tone and profanity interactions $F = 0.52$), so the within-creator slope is the same on women's and men's

videos. Limitations include the video panel's non-representativeness of the full channel sample and inability to rule out video-level reverse causality; temporal ordering is addressed subsequently. Phase one thus establishes that audience reception, characterized by positive comment tone and profanity, associates with visibility both across creators and within creators' own content. The next section addresses gender differences in visibility and their mediators.

The Gender Gap in Visibility and What It Survives

Table 5 presents a progression of models estimating the gender gap in logged views on the same sample of 9,096 channels, shown in Panel A without fixed effects and Panel B with 40 content-category fixed effects. The models enter controls in this order: creator gender alone (Model 1); channel age and video count as production controls (Model 2); creator content measured by the tone and profanity of video titles and descriptions (Model 3); comment volume (Model 4); comment profanity (Model 5); and comment tone (Model 6). Figure 4 plots the creator-gender coefficient across these six models in both panels. Without controls, women creators' channels have logged views 0.41 lower than men's, equivalent to about 34 percent fewer views ($p < .001$; Panel A, Model 1). Adjusting for channel age and video count reduces the gap to -0.22, about 19 percent fewer views ($p < .001$; Model 2), and including category fixed effects leaves it essentially unchanged at -0.23, about 20 percent fewer views (Panel B, Model 2). This production-adjusted gap is the headline measure of the gender visibility gap, reported both ways, about 19 percent at the same channel age and video count and about 20 percent within content categories, because category sorting can itself reflect discrimination, so holding category constant risks over-controlling (Blau and Kahn, 2017). Women's channels average 834 videos versus 1,495 for men (the difference in means is not significant at the .05 level, $p = .054$) and 122.3 versus 175.1 videos per year of tenure ($p = .061$), while logged videos per year of tenure are lower for women ($p < .0001$; Appendix Table A17, Panel A). An order-free Gelbach decomposition of the raw gap, with content category as its own block, shows channel age and video count account for 56.4 percent of the raw gap (Table 6, column 1). A wild cluster bootstrap test using 40 content categories supports the statistical

significance of the headline gap, with p-values of .048 without and .029 with fixed effects, and the volume-adjusted gap at .006 and .010 (Appendix Table A24).

< INSERT TABLE 5 HERE >

< INSERT FIGURE 4 HERE >

Adding creator-controlled content variables, specifically the tone and profanity of video titles and descriptions, moves the gender coefficient to -0.13, about 12 percent fewer views ($p = .004$; Model 3), and to -0.17, about 15 percent fewer views with category fixed effects (Panel B, Model 3). Comment volume enters next because assessing differences in audience reception requires holding the amount of reception constant. Women's channels receive more comments on average (geometric mean 67 versus 58 across sampled videos; Table 1, Panel B), so controlling for comment volume enlarges the gap to -0.24, about 21 percent fewer views (Model 4), and to -0.20, about 18 percent fewer views with category fixed effects. The volume-adjusted gap, holding channel age, video count, and comment volume constant but excluding creator content, is -0.26, about 23 percent fewer views without fixed effects and -0.23, about 20 percent fewer views with fixed effects (Appendix Table A24). The size of the content step depends on its order of entry: when entered before comment volume, titles and descriptions move the gap from -0.22 to -0.13; when entered after comment volume, they move the gap from -0.26 to -0.24 (Appendix Table A02). Women's titles and descriptions are more positive than men's (title tone 55.8 versus 42.9; description tone 66.2 versus 55.7; Table 1, Panel B), and more positive creator content correlates with fewer views (Table 5, Model 3). Because the share of the gap explained by any block depends on the order of entry, Table 6 reports an order-free decomposition (Gelbach, 2016) presented alongside mediation results. The next step is to test whether sorting into content categories and subtopics explains the remaining gap.

Adding the category indicators to the volume-adjusted model moves the coefficient from -0.263 to -0.226 (Table 8), but in the Gelbach decomposition, which holds every block equal at once, the category block widens the gap, even though women's categories average fewer views

than men's on raw averages (category mean logged views, women minus men, -0.11). Holding content category equal while every other block is held equal leaves the gap larger, not smaller (the category block's share is -22.9 percent of the 19 percent gap and -12.5 percent of the raw gap; Table 6, columns 2 and 1). Further refining to 128 subtopics, constructed by clustering sentence embeddings of channel video titles within categories, reduces the gap from -0.226 to -0.178 ($p < .001$ under both robust and subtopic-clustered standard errors; Appendix Table A20, columns 2 to 4). Importantly, the gendered pattern in audience reception is not an artefact of topic sorting. Within subtopics, women's channels continue to receive more positive comment tone (0.38 standard deviations) and less profanity (-0.30 standard deviations), both differences significant at $p < .001$ (Appendix Table A20, columns 5 and 6). This persistence of softer reception within tightly defined content niches suggests that neither broad nor fine-grained topical clustering explains the gendered differences in visibility or expressed reception. Having ruled out production volume and topic sorting as full explanations, the next step is to examine temporal ordering to address rival interpretations of the reception-visibility association.

To address reverse causality concerns, I employ a lagged design using three waves of data. First-wave comment language measures predict logged views at the second wave (eight months later) and third wave (twelve to thirteen months later), controlling for first-wave logged views and the channel and comment volume controls (Appendix Tables A05 and A06). At eight months ($N = 9,054$), first-wave positive comment tone is negatively associated with later logged views (-0.03, $p < .001$), while first-wave profanity and creator gender are not statistically distinguishable from zero. At twelve to thirteen months ($N = 8,993$), positive tone remains negatively associated (-0.04, $p < .001$), slightly stronger than at eight months, with profanity and creator gender again indistinguishable from zero. In models without comment language on the same channels as those with it, the creator-gender coefficient on later logged views is not distinguishable from zero at eight months ($p = .52$; Appendix Table A05, Model 1) or at twelve to thirteen months ($p = .09$; Appendix Table A06, Model 1). The coefficient remains non-significant in all lagged models, whether comment language is included or excluded. These results indicate that the gender gap resides

primarily in the accumulated stock of views rather than emerging newly during the observation windows, and that comment language measured prior to the outcome window retains its association net of baseline viewership. Attrition between waves is minimal and unrelated to creator gender (retention 99.6 percent; Appendix Table A07). While this design supports temporal ordering from expressed reception to visibility, it does not eliminate the possibility of unmeasured confounding and does not establish causality.

Together, these findings demonstrate that the gender gap in visibility persists after accounting for production volume, content topic, and temporal ordering. The next analytical step will test the gendered mechanisms proposed: first, whether women's channels exhibit a softer register in audience reception; second, whether the association between reception and visibility is comparable across genders; and third, whether expressed reception statistically mediates the gender gap in visibility.

Gendered Reception and Mediation

Figure 3 outlines the chain to be estimated: from creator gender to expressed reception, from expressed reception to amplification via audiences and the recommendation system, and from amplification to visibility; the direct path from gender to visibility is the focal test here. Women creators' channels receive a softer reception on both key dimensions of expressed reception. The mean positive comment tone on women's channels is 62.8 compared to 52.7 on men's channels, while swear words constitute 0.19 percent of comment words on women's channels versus 0.43 percent on men's channels, about 55 percent less frequent (both $p < .001$; Table 1, Panel B). These findings support hypotheses H2a and H2b and reject the hostile rival that women's comment sections are more negative and profane. The softer register originates from the entire commenting audience: male-coded commenters write 10.0 tone points more positively on women's channels than on men's, and female-coded commenters 9.8 points more positively; both groups also use less profanity on women's channels (all $p < .001$; Appendix Table A12). This pattern holds within subtopics as established in Part 2. This softer treatment represents a gendered evaluative

act consistent with benevolent sexism, rather than an absence of evaluation. The study does not interpret more positive tone as stronger approval; measures of approval-type engagement such as engagement rate, subscribers, and comment volume are treated separately below.

< INSERT FIGURE 3 HERE >

The association between comment reception and views is examined as a check rather than a hypothesis. On the within-creator video panel with channel fixed effects and LIWC-22 measures, the woman-by-tone interaction is 0.000 ($p = .997$), and the joint test of tone and profanity interactions yields $F = 0.52$, indicating the slope relating comment tone to views is the same across women's and men's uploads (Appendix Table A22, column 6). Across creators, a model controlling for channel age, video count, and comment volume shows a woman-by-tone interaction of -0.10 ($p = .005$; $N = 9,096$) (Appendix Table A15, Model 1), reflecting steeper negative associations between positive comment tone and views within female-typed content categories than outside them for both genders. Within female-typed categories, the slope is -0.56, and outside them it is -0.24 (Model 2). Because 72 percent of women's channels and 34 percent of men's channels are in these categories, the pooled interaction reflects where women's channels sit rather than a gendered slope. When allowing each of the 40 content categories its own tone slope, the woman-by-tone interaction is +0.04 ($p = .39$) (Model 3). The subsequent mediation analysis does not require equal slopes.

Testing Hypothesis 3 within the model progression of Table 5 builds on the logic established earlier: since softer expressed reception correlates with fewer views and women's channels receive softer reception, incorporating comment profanity and comment tone into the gender gap model should substantially reduce the observed visibility gap if these reception differences carry it. Beginning with Model 4, which controls for channel age, video count, creator content, and comment volume, the creator-gender coefficient stands at -0.24 without fixed effects and -0.20 with 40 category fixed effects. Introducing comment profanity in Model 5 moves this coefficient closer to zero, to -0.17 and -0.16 respectively. Further adding comment tone in Model 6 reduces the coefficient to values not distinguishable from zero, with p -values of 0.90 without fixed effects and 0.47 with fixed effects. In this full model, comment tone shows a strong negative association with

logged views (coefficients of -0.40 without and -0.44 with fixed effects, both significant at $p < .001$), while comment profanity exhibits a positive association (+0.08 without and +0.05 with fixed effects, significant at $p < .001$ and $p < .01$ respectively). Figure 4 visually depicts this progression. All models use the same analytic sample of 9,096 channels. The 136 sampled channels outside the analytic sample are not disproportionately women's channels (1.54 percent versus 1.44 percent, $p = .71$), and missingness is unrelated to creator gender when adjusting for controls ($p = .55$; Appendix Tables A18 and A28).

To apportion the reduction in the creator-gender coefficient independently of the order in which variables enter the model, the Gelbach decomposition (Gelbach, 2016) is applied as shown in Table 6. This method attributes the total change in the gender coefficient between a baseline model and the fully adjusted model across all blocks simultaneously. Each block's share equals the product of the gender difference in that block and its coefficient in the full model; the shares of all blocks plus what is left over sum to 100 percent of the starting gap. The blocks included are channel age and video count, content category, creator content, comment volume, comment profanity, comment tone, HD video quality, and non-subscriber trailer presence. The decomposition of the 20 percent gap at the same channel age, video count, and content category shows that comment tone accounts for 88.0 percent and comment profanity for 6.7 percent (both significant); creator content accounts for -7.3 percent (significant at the .05 level; holding content equal widens the gap, because women's titles and descriptions are more positive and more positive content correlates with fewer views); comment volume, HD video, and the trailer are not distinguishable from zero; and 10.5 percent is left over. The decomposition of the 19 percent gap at the same channel age and video count, with content category as its own block, shows content category accounts for -22.9 percent, creator content -11.3 percent, comment profanity 9.5 percent, comment tone 130.0 percent, and the trailer 4.3 percent (all significant); comment volume and HD video are not distinguishable from zero; and 11.0 percent is left over. Tone exceeds 100 percent because category, creator content, and comment volume offset it. The decomposition of the raw 34 percent gap shows channel age and video count account for 56.4 percent, comment tone 67.9 percent, comment profanity 5.0 percent, the trailer 3.6

percent, content category -12.5 percent, creator content -5.9 percent, and comment volume -20.0 percent (all significant); HD video is not distinguishable from zero; and 5.8 percent is left over. Comment volume's share disappears within categories because its full-model coefficient is fixed, so its share moves only with the women-minus-men difference in comment volume; women's channels receive 15.0 percent more comments raw (significant), but the difference is not distinguishable from zero at the same channel age and video count ($p = .077$) or within content categories ($p = .99$). The fully adjusted creator-gender coefficient is not distinguishable from zero ($p = .53$). These results demonstrate that at equivalent channel age, video count, and content category, the visibility gap is primarily accounted for by the reception women's channels receive, while creator content contributes to widening rather than explaining the gap, video quality has no explanatory role, and the trailer's contribution is minimal.

< INSERT TABLE 6 HERE >

The mediation analysis directly tests the theorized pathway from creator gender through comment tone and profanity to logged views, employing bootstrap percentile intervals for inference. This analysis is conducted at equal comment volume, isolating the volume-adjusted gap while holding channel age, video count, and comment volume constant, which is -0.26 without fixed effects (equivalent to 23 percent fewer views) and -0.23 with category fixed effects (20.3 percent fewer views). Under ordinary least squares with the same controls, each Gelbach decomposition share corresponds to the product-of-coefficients estimate of indirect association. With category fixed effects, comment tone and profanity jointly account for 93 percent of this volume-adjusted gap, with a 95 percent bootstrap interval ranging from 70 to 137 percent based on 5,000 replications ($N = 9,096$; Appendix Table A21, Panel A). Without fixed effects, the mediated share is 107 percent with an interval of 85 to 143 percent (Panel B). Intervals exceeding 100 percent reflect a direct association near zero rather than a change of sign. The direct association between creator gender and views is not distinguishable from zero in either specification. The mediation paths show that, with category fixed effects, women's channels receive comment tone 0.46 standard deviations more positive and comment profanity 0.32 standard deviations lower than men's channels (Table

4, Models 3 and 6). Comment tone carries approximately 93 percent of the indirect association, with comment profanity accounting for the remaining 7 percent. The profanity path's interval also excludes zero. Across five alternative transformations of the profanity measure including raw, logarithmic, inverse hyperbolic sine, winsorized, and binary forms, every indirect effect interval excludes zero, and the creator-gender coefficient in the views model remains not distinguishable from zero (Appendix Table A09).

Women's own content, specifically video titles and descriptions, is more positive than men's, suggesting that part of the softer reception in comments could be a response to creators' posted content. To quantify this, models holding constant the tone and profanity of creators' titles and video descriptions estimate that creator content accounts for approximately 20 percent of the gender difference in comment tone at the same channel age, video count, and comment volume and within content categories (23 percent without category fixed effects) and 7 percent of the difference in comment profanity (11 percent without), controlling for channel age, video count, and comment volume (Appendix Table A29, Panel A). Including transcripts for 8,082 channels increases the accounted share of the tone difference to 22 percent with category fixed effects (28 percent without) (Appendix Table A29, Panel B). Consequently, between 77 and 80 percent of the gender difference in comment tone remains unexplained by creator content, indicating that most of the softer reception women receive is not merely a reaction to softer content posted by women creators.

Alternative Explanations and Robustness

Three alternative explanations for a softer register in women's comment sections do not involve gendered treatment. First, reciprocity: women's own content, including titles, transcripts, and thumbnails, could be softer. If this is the source, controlling for creator language should remove the comment difference. Second, moderation: hostile comments could be removed before observation through deletion or filtering. This remains untested here and constitutes a limitation. Third, self-selection: differences could arise from who chooses to comment on women's videos, predicting distinct comment patterns by commenter gender. This is tested by comparing positive

comment tone and profanity from male-coded and female-coded commenters.

Table 8 presents each rival explanation, the corresponding test, the estimate, and the verdict. The following paragraphs address the rows sequentially. Women creators' own content differs from men's: title tone is higher with a Cohen's d of 0.40, transcript tone is higher by 0.35, thumbnails show more positive facial expressions ($d = 0.47$) and less clickbait ($d = -0.23$), all $p < .0001$ (Appendix Table A17, Panel C). Therefore, the reciprocity rival, which posits that softer comments reflect softer creator content, is active and requires testing. Adding creator title and description language to the controlled model on all 9,096 channels reduces the gap from -0.26 to -0.24; the coefficient remains significant at $p < .001$ (Appendix Table A02). In the Gelbach decomposition of the 20 percent gap, creator content's share is -7.3 percent, so holding content equal widens the gap (Table 6, column 3). When all four creator-controlled content elements (titles, descriptions, transcripts, thumbnail packaging) enter a single equation alongside comment reception as parallel mediators, with category fixed effects and channel and comment volume controls ($N = 6,727$), comment reception accounts for 92.9 percent of the volume-adjusted gap within content categories (joint mediation; Appendix Table A04), while the content elements are either null or act as suppressors. Thus, creator content is not the primary source of the gap.

< INSERT TABLE 8 HERE >

The second rival concerns concentration in low-visibility categories and gender-typed category dynamics. Holding content category equal widens the 19 percent gap (category share -22.9 percent; Table 6, column 2), although women's categories average fewer views than men's on raw averages (category mean logged views, women minus men, -0.11). The coefficient for the volume-adjusted gap moves from -0.263 to -0.226 with the inclusion of the 40 category indicators. Excluding six gaming categories, the gender gap at the same channel age, video count, and comment volume is -0.27 ($p < .001$; $N = 8,471$), compared to -0.26 on all 9,096 channels (Appendix Table A17, Panel D). An Oaxaca-Blinder decomposition of the raw gap across 40 content categories attributes 0.07 to category composition, which is not statistically distinguishable from zero ($p = .16$). The same decomposition over channel age, video count, and comment volume attributes 0.15

to composition and leaves 0.26 unexplained (both $p < .001$) (Appendix Table A03), corroborating the Gelbach result. Testing interactions among creator gender, positive comment tone, and an indicator for female-typed categories yields a three-way interaction coefficient of -0.04 ($p = .61$, $N = 9,096$), and the woman-by-female-typed category term is not distinguishable from zero (Appendix Table A15, Model 4). No evidence supports a backlash dynamic specific to women creators in female-typed categories. The positive tone slope itself is steeper in female-typed categories (positive-tone-by-typed coefficient -0.32, $p < .001$; Appendix Table A15, Model 2), reflecting a feature of the association rather than creator gender.

The third and fourth rivals concern audience factors: whether audiences simply value women's channels less and whether commenter gender composition explains the softer register. Audiences are not less willing to interact with women creators. The raw subscriber gap is about 10 percent (-0.11 logged subscribers); holding channel age and video count constant closes it entirely ($p = .87$; Table 8; Appendix Table A17, Panel F). Engagement rates are higher on women's channels (3.90 versus 3.25 percent), as is comment volume. The visibility disadvantage is specific to reach. Women's channels receive shorter comments (19.1 versus 20.7 words per comment, $p < .0001$). Adding mean comment length to the controlled model increases the gap from -0.26 to -0.30, indicating comment length acts as a suppressor rather than an explanation (Table 8; Appendix Table A08). Commenter gender composition does not explain the softer register: both male-coded and female-coded commenters write about 10 tone points more positively and swear less on women's channels (Appendix Table A12). Controlling for the share of female-coded commenters leaves the creator-gender coefficient at -0.19 ($p < .001$); it becomes indistinguishable from zero only when positive comment tone and profanity are added alongside the composition control ($N = 9,091$; Appendix Table A14). Mediation occurs through the comment language of either commenter group: entering positive tone and profanity from male-coded commenters alone, female-coded commenters alone, or both simultaneously results in a creator-gender coefficient not distinguishable from zero in each case (Appendix Table A13).

The remaining rivals concern off-platform audience diversion, robustness of estimates,

sensitivity to unmeasured confounding, and measurement validity. Women creators link out to other platforms more frequently (20.7 versus 16.4 percent, $p < .0001$; Appendix Table A16). Including multi-homing as a control leaves the controlled gap at -0.26, and its interaction with creator gender is not distinguishable from zero (Appendix Table A16). The non-subscriber channel trailer, a promotional feature shown to first-time visitors, is used by 55.7 percent of women's channels and 61.6 percent of men's channels. Running a trailer is associated with 0.28 more logged views. Adding the trailer indicator to the volume-adjusted model moves the creator-gender coefficient from -0.26 to -0.25. The return to a trailer does not differ by creator gender ($p = .43$). The verdict is that the gap is not a matter of promotional tactics. The full mediation model remains stable across clustered standard errors, trimming logged views at the 1st and 99th percentiles ($N = 8,915$), restricting to channels active at least twelve months ($N = 9,017$), and category fixed effects (Table 5, Panel B, Model 6); in every model, the creator-gender coefficient is not distinguishable from zero and positive comment tone remains near -0.40 ($p < .001$) (Appendix Table A08). The median channel's comment tone and profanity measures are based on 772 scored comments, with a split-half reliability of 0.95 for channel tone (Appendix Table A25). Across multiple robustness checks including minimum comment floors of 20, 50, 100, 200, and 500 scored comments per channel, weighting channels by comment count, and recomputing tone via different methods, the creator-gender coefficient in the full model remains not distinguishable from zero in all twelve specifications. Additionally, controlling for each channel's like rate, the comment tone coefficient retains 84.5 percent of its size without fixed effects and 83.6 percent with category fixed effects (Appendix Table A26). Within creators, on the video panel with channel fixed effects, the negative association between tone and views persists at 53 percent of its size (-0.12 versus -0.23, $p < .001$), indicating that tone is not a proxy for approval measured by likes. With like rate held constant, the creator-gender coefficient is +0.08 ($p = .014$) without fixed effects and not distinguishable from zero with them. The association between positive tone and views is stronger where comment volume is larger (positive-tone-by-volume interaction coefficient -0.04, $p < .001$). The volume-adjusted gap is stable under winsorized logged views and at the 25th, 50th, and 75th percentiles

of the view distribution (-0.26, -0.27, -0.26, -0.22; Appendix Table A17, Panel D). Influence and collinearity diagnostics show low values (maximum Cook's distance 0.009; mean variance inflation factor 1.04; Appendix Table A17, Panel E). An unmeasured confounder would need to explain 8 percent of the residual variance in both creator gender and logged views to reduce the volume-adjusted gap to zero (robustness value 0.080), and 20 percent for the positive tone association and 4.5 percent for profanity in the mediation model; none of the observed controls approaches this explanatory power (Appendix Table A10). LIWC-22 tone correlates 0.87 with the independent VADER compound score and -0.33 with LIWC-22 swear (N = 9,096; Appendix Table A19). Creator gender (hand coding), comment language (LIWC-22 on comments), and views (API metrics) derive from independent sources, so a shared measurement artefact cannot account for the observed pattern.

Theoretical Implications

The paper began by situating itself within the prevailing consensus that entrepreneurs secure critical resources when external stakeholders evaluate them favorably (Zuckerman, 1999; Hsu, 2006; Pollock and Rindova, 2003), with women's entrepreneurial disadvantage attributed largely to receiving less favorable evaluations. The evidence presented challenges the straightforward extension of this view to visibility on engagement-optimized platforms. Specifically, the association between audience evaluation and visibility runs counter to favorability: channels with more negative and more profane comment sections exhibit greater visibility both across creators and within a creator's own uploads. Conversely, more positive comment tone corresponds with fewer views, a pattern identified here as the negative arousal bonus. Comment tone accounts for most of this association, while profanity contributes a smaller but meaningful share. Women creators experience a distinctly softer reception, with average comment tone on their channels at 62.8 compared to 52.7 on men's channels, and swear words constituting 0.19 percent of comment words versus 0.43 percent on men's channels. This softer reception emerges as a shared treatment norm across commenters of both genders, rather than an artifact of audience composition. Importantly, within

a creator's own uploads, where channel fixed effects absorb every stable creator characteristic, the relationship between comment tone and visibility does not differ by creator gender (interaction 0.000, $p = .997$), indicating that the visibility gap arises not because women's reception is amplified differently but because their comment sections lie toward the more positive end of a slope common to all channels. Quantitatively, women creators receive 34 percent fewer views unconditionally and approximately 19 percent fewer views at the same channel age and video count, holding comment volume constant, which favors women, widens the gap to 23 percent, and about 20 percent fewer views at the same channel age, video count, and content category, reported alongside the 19 percent because category sorting can itself reflect discrimination, so holding category constant risks overcontrolling (Blau and Kahn, 2017). About half of the raw gap reflects women posting fewer videos and channels having been around less time. The Gelbach decomposition of the 20 percent gap (same channel age, video count, and content category) attributes 88 percent to comment tone and 7 percent to profanity. In the joint mediation test of the path, tone and profanity account for 93 percent of the volume-adjusted gap within content categories. The creator-gender coefficient that remains is not distinguishable from zero. Taken together, these findings suggest that gendering occurs one step prior to allocation, in how women creators are received, and that the allocation criteria price reception according to its intensity rather than its favorability. The discussion now proceeds by addressing each contribution in the order promised in the introduction, beginning with the conceptualization of the criteria.

Prior research has demonstrated that evaluation and reward can diverge due to processes occurring after evaluation. For example, equal performance ratings may yield unequal pay (Castilla, 2008), reward criteria may be inconsistently applied across groups (Correll et al., 2020), funding outcomes may depend more on procedural features than assessed quality (Bol et al., 2022), and the stakes attached to each evaluation stage change who advances (Botelho and Poskanzer, 2026). Additionally, ostensibly neutral features of evaluation processes can produce gendered outcomes when applied to gendered inputs, as shown in cases of meritocratic framing (Castilla and Benard, 2010), cost-neutral advertisement delivery (Lambrecht and Tucker, 2019), and engagement-based

ranking systems (Milli et al., 2025). This paper contributes by foregrounding the criterion itself as a design feature. Whereas prior accounts treat favorable evaluation as the input rewarded faithfully or with bias, here the criteria operate in the opposite direction, privileging intensity over favorability. Consequently, a process may convert an input that appears favorable into fewer resources. The finding that the slope of the reception-to-visibility association is the same for women and men within creators supports this as a criteria story rather than a bias-in-conversion story: the criteria are formally gender-neutral but produce gendered outcomes because the input differs by gender. Disadvantage thus can emerge at the allocation stage even when evaluation appears favorable. Within creators, where channel fixed effects absorb every stable creator characteristic, the tone-to-views association does not differ between women's and men's videos (interaction 0.000, $p = .997$).

Readers familiar with research on negativity and spread or attention as a resource may question what this criteria finding adds. Prior work has established that negative and high-arousal content spreads further (Baumeister et al., 2001; Rozin and Royzman, 2001; Berger and Milkman, 2012; Brady et al., 2017; Rathje et al., 2021), and management research has long shown that attention can be allocated based on audience engagement rather than approval, allowing popularity and favorable judgment to diverge (Rindova et al., 2006, 2007; Kovács and Sharkey, 2014; Graf-Vlachy et al., 2019). Three features distinguish the present findings. First, the outcome is a producer's cumulative resource access, total channel views, rather than the spread of a single message. Second, the signal is observed on the receiving end, in what audiences write, while controlling extensively for creator-produced content. Within creators, comment negativity associates with views, whereas the creator's own video language does not. Third, in prior attention literature, attention generally precedes favorable judgment by investors or experts; here, visibility itself constitutes the resource that converts into revenue (Nason et al., 2024; Ørmen and Gregersen, 2022). The interpretation offered is that comment negativity likely marks contestation (arguing, replying, dwelling) that audiences propagate and recommendation systems interpret as engagement intensity (Covington et al., 2016). Profanity serves as a coarser signal of arousal. Although profanity is rare in this

corpus (median channel has 0.16 percent swear words; 14 percent of channels have none; Appendix Table A17), a one standard deviation increase, approximately five swear words per 1,000 comment words, is associated with about 8 percent more views. The corpus mean of 0.35 percent falls within the 0.3 to 0.7 percent base rate of everyday spoken swearing (Jay, 2009). Taboo words capture attention (MacKay et al., 2004) and produce measurable autonomic arousal (Bowers and Pleydell-Pearce, 2011), supporting the interpretation of profanity as indexing arousal rather than valence.

Having established the criteria as a general feature applying to all channels, attention turns to the gendered aspect of the findings, which resides in the input. Research on gendered evaluation in high-stakes settings finds that women are often judged more harshly or deemed less suited for entrepreneurial roles (Botelho and Abraham, 2017; Kanze et al., 2018; Snellman and Solal, 2023). Conversely, research on benevolent and ambivalent sexism identifies a contrasting pattern in routine and interpersonal encounters, where women receive warmer, more protective, and less confrontational treatment (Glick and Fiske, 2001; Heilman, 2012; Bareket and Fiske, 2023). This benevolent treatment carries developmental costs, including less candid and inflated feedback (Jampol and Zayas, 2020), fewer challenging assignments (King et al., 2010), and positive stereotypes that evade traditional bias detection (Czopp et al., 2015). The comment sections examined here resemble this second pattern rather than the first: women creators receive more positive and less profane comments from commenters of both genders. This softer reception constitutes a gendered evaluative act in its own right. It may reflect benevolent regard or a reduced willingness to engage women critically and intensely. As the data do not allow adjudication between these interpretations, I would not read reception as ubiquitous approval. The nuanced reading here may be instead that benevolent treatment, previously shown to impose developmental costs, here carries a material cost in views and consequently revenue because the criteria price intensity. The penalty operates through two linked steps. First, softer treatment at evaluation and, second, the conversion of that treatment into visibility. Both steps produce gendered effects, although only the first is gendered in form. Because the slope of the reception-to-visibility association is the same

for women's and men's videos within creators (interaction 0.000, $p = .997$), and the channel-level difference reflects category composition, women's disadvantage arises from differential exposure to a common slope rather than from differential conversion of reception into visibility. This boundary contrasts with settings where gendered conversion has been documented at later or more prominent stages (Botelho and Poskanzer, 2026; Frankort et al., 2026). Benevolent sexism in startup evaluation has also been found to advantage men rather than be neutral (Nguyen et al., 2023), consistent with interpreting softer reception as a disadvantage in disguise.

If softer treatment is the typical experience for women creators, the reader will recall that a substantial literature documents hostility toward women online. Some studies document hostility, especially where stakes, policing, and male-typed domains are salient (Döring and Mohseni, 2020; Sobieraj, 2017), while others find softer, less negative treatment in more routine, relational interactions (Weißmann et al., 2025, preprint). The hostility-focused studies often sample on hate speech or interview targeted individuals, capturing specific encounters rather than typical interactions. This paper's random sample of established creators, encompassing over 5.7 million everyday comments, captures typical interactions in which women more commonly experience softer treatment. Independent corroboration comes from a study of 3,695 German-speaking channels with 40 million comments, where women receive less negative communication (Weißmann et al., 2025, preprint). These findings are compatible because both receptions occur. The resource cost to women depends on the allocation criteria: under favorability-based criteria, hostility imposes costs; under engagement-based criteria, softer treatment associates with lower visibility. The criteria, not the direction of treatment alone, determine which form of reception carries resource consequences in a given setting.

The introduction mentioned analogs such as a venture's press coverage alongside its funding, a scholar's work being cited and debated alongside its grant score, and organizational feedback that is inflated and softer for women despite equivalent work (Jampol and Zayas, 2020). Extending this, two conditions must hold for the mechanism identified here to travel to other resources. First, the resource must be allocated based on the volume or intensity of engagement rather than on a central

evaluator's favorable judgment. Second, treatment norms must differ by group while the criteria apply equally to all. Where both conditions hold, softer treatment toward a group will associate with less of the resource without any actor explicitly withholding it. Where the first condition fails, the same softer treatment will associate with more, not less, of the resource. These two criteria can coexist within a single career and pull in opposite directions: a favorability-allocated resource (such as funding, a grant score, or promotion) and an engagement-allocated resource (such as press, citations, or visibility) may respond oppositely to the same softer treatment. This explains why the favorability account requires a scope condition rather than replacement. Lowering barriers to entry does not eliminate disadvantage; it shifts it from access to amplification (Cutolo and Kenney, 2021). Women can enter platforms, attract higher engagement rates, convert viewers into subscribers at rates indistinguishable from men's, and yet still be watched by fewer people. In sum, gendered treatment norms combined with engagement-based allocation reproduce, in a new technological setting, the gendered distribution of resources and rewards long documented in organizations (Acker, 1990). Audiences following benevolent norms act as biased actors without animus; the criteria are unbiased actors; yet the outcome remains gendered.

Implications for Creators and Platforms

Advising women creators to close the visibility gap by altering their content misinterprets the source of the disparity. In the Gelbach decomposition of the 20 percent gap, creator content's share is -7.3 percent, so holding it equal widens the gap (Table 6); in the joint mediation of the volume-adjusted gap within content categories (6,727 channels), the three creator-side elements together account for -21.0 percent (Appendix Table A04). In contrast, comment treatment explains 92.9 percent of the volume-adjusted gap within content categories (joint mediation, 6,727 channels; Appendix Table A04). In the Gelbach decomposition, holding content category equal leaves the gap larger, not smaller (category share -22.9 percent of the 19 percent gap; Table 6); finer subtopics move the volume-adjusted gap within categories from -0.226 to -0.178 without eliminating it (Appendix Table A20). This evidence indicates that the margin creators can move through production choices

is relatively small. The cost of this visibility gap is material because visibility translates into monetization. Creators receive 55 percent of advertising revenue on monetized videos, and earnings scale with views (Verwiebe et al., 2026; Bhargava, 2022). Therefore, a 19 percent view gap at the same channel age and video count and about 20 percent fewer views at the same channel age, video count, and content category approximates a comparable gap in advertising income before accounting for sponsorships and off-platform opportunities. Independent evidence from German-speaking creators documents women earning significantly less than men on the same platform (Verwiebe et al., 2026). This disadvantage thus appears to be a demand-side treatment and allocation problem rather than a supply-side production problem. The lever that could address it primarily belongs to the platform rather than the creators. This shifts attention from content adjustments to the mechanisms governing attention allocation.

The platform's recommendation system was last publicly described as optimizing expected watch time per impression (Covington et al., 2016). Engagement-based ranking has been shown to amplify divisive, high-arousal content beyond stated user preferences, with rerankings altering what content is amplified (Milli et al., 2025). Amplification operates both through the recommendation system and through audiences passing content along. The data do not separate the contributions of these two channels, but both respond to the same intensity signals embedded in audience engagement. These criteria translate a gendered treatment norm into unequal resource access without any actor explicitly deciding to disadvantage women. Transparency about the platform's objective would not reveal this dynamic because the objective contains no gender-related terms; the gendered input originates from audience reception. Content moderation would not address this issue either, because the input that disadvantages women is the absence of intense engagement, not the presence of harmful or intense engagement. The intervention point lies in which behavioral signals the recommendation system values and the weight it assigns to them. Reducing the weight on comment intensity or increasing the weight on favorable but low-arousal engagement would alter how softer reception converts into views. Platform-dependent creators have limited bargaining power over these design choices (Cutolo and Kenney, 2021). Therefore, the design feature of the

platform’s attention allocation system, rather than creator coaching, constitutes the policy-relevant object for addressing the visibility gap.

Limitations and Future Directions

All estimates presented are associations derived from cross-sectional data. The only temporal support comes from the lagged design, where first-wave comment language predicts logged views eight and twelve to thirteen months later, controlling for baseline views. Attrition is unrelated to creator gender (Appendix Tables A05, A06, A07). Channel fixed effects absorb stable creator traits at the video level. However, neither approach establishes causality. Residual confounds remain despite supply-side tests. Comment tone could proxy for unmeasured creator-side content that audiences respond to but that is not captured by title, description, transcript, or thumbnail measures. Viewer composition beyond commenter gender is unobserved. Click-through and impression delivery data are unavailable. At the video level, reverse causality where views draw harsher comments is not excluded. The sensitivity analysis bounds the first concern, indicating that an unmeasured confounder would need to explain more than 8 percent of the residual variance in both creator gender and logged views to eliminate the volume-adjusted gap, which exceeds the explanatory power of any observed control (Appendix Table A10). Nonetheless, these confounds are not removed. The two amplification channels, audiences passing content along and the recommendation system, are not separately identified. The evidence establishes that intensity associates with visibility but does not clarify which channel contributes more to this association.

The sample comprises established, individually identifiable, U.S.-based, English-language creators drawn at random within the Infludata frame, with a subscriber threshold of approximately 2,000 and above. Creator gender is hand-coded as binary from public profiles. Very new creators, accounts below coverage thresholds, non-individual accounts, and creators with illegible or nonbinary gender presentation fall outside the estimates. Early-stage creators may experience different treatment because recommendation systems handle new channels through exploration rather than exploitation of known engagement. Consequently, the intensity-to-visibility association observed

here need not hold in the same form before a channel establishes an engagement history. The estimates describe creators who have already passed that early stage. Selection into the frame is gendered: 65.5 percent of the sample are men, reflecting survival to the threshold. If women who survive to this stage are positively selected relative to those who exit earlier, the observed gap likely represents a conservative lower bound on the gap in the full population of entrants. This selection is distinct from within-sample missingness, which is unrelated to gender (Appendix Table A18). The study focuses on one platform and one time window. Recommendation systems evolve, and platforms differ in their reliance on recommendation algorithms versus follower networks, as well as in how legible creator gender is. Therefore, the criteria observed here represent one platform's criteria at one point in time.

Comments capture expressed reception, which is public and socially regulated. They combine audience judgment and communicative register, and these two components cannot be separated here. Private evaluation of women's and men's content may differ from what is publicly expressed in comments. Comment moderation is untested here; hostile comments that platform or creator moderation removed before observation, by deletion, review holds, or blocked words, would leave hostility underrepresented in the corpus, and the design cannot see them. Moderation of the observed pattern is largely untested. The paper establishes that the slope of the reception-visibility association does not differ by creator gender within creators, and that the channel-level difference reflects category composition (Appendix Table A15). However, it does not test which creator, category, or audience conditions strengthen or weaken the softer-treatment pattern or its conversion into visibility. Intersectional characteristics such as race and age are outside the scope. Evidence from German-speaking creators suggests that creators of color face more negative sentiment, while women face less (Weißmann et al., 2025, preprint), indicating that the softer pattern may be specific to gender rather than to marked identities generally. Visibility is treated as a resource because it converts into revenue, but reach and status can diverge (Mears, 2023). Visibility should not be equated with esteem, and attention is not always desired by the creator who receives it. The paper measures a resource, not a reputation. Cumulative views conflate impression delivery with viewer

choice. Public data do not separate Shorts from long-form video formats, limiting granularity.

Each limitation identifies a quantity or relationship that future research designs could target. Exogenous variation in the criteria such as platform redesigns or ranking policy changes would allow tests of whether changing which signals are valued alters how softer reception converts into visibility. Cross-platform comparisons across systems differing in reliance on recommendation algorithms, follower networks, and gender legibility would test the portability of these findings beyond the studied platform. Collecting private audience evaluations alongside public comments would separate judgment from communicative register. Access to traffic-source and impression-level data would apportion the relative contributions of the two amplification channels. A longitudinal panel following creators from their first upload would test early-stage treatment and gendered survival into the established population within a single design. The same logic may apply to other categorical distinctions such as race, age, or class wherever treatment norms differ by group and the resource is allocated based on engagement. It may also apply to other engagement-allocated resources such as press coverage, citations, and debate. These directions are offered as avenues for research rather than claims.

This study addressed two central questions about entrepreneurial visibility on engagement-optimized platforms. First, it found that more negative and more profane audience reception is associated with greater visibility both across creators and within individual creators' uploads. This indicates that favorable reception and visibility operate as different currencies in these contexts. Second, the relationship between reception and visibility is gendered not because the criteria convert women's reception differently (within creators the slope is the same for women's and men's videos) but because women creators receive a distinct, softer reception consistent with benevolent sexism. This softer reception is associated with fewer views despite women's higher engagement rates and a controlled subscriber gap that is statistically indistinguishable from zero.

The broader implication is that opening a market to all participants does not guarantee equal rewards when the criteria that distribute resources convert gendered treatment into gendered resource access without any actor explicitly deciding to do so. The consensus that favorable

evaluation secures resources holds primarily for resources conferred directly by a central evaluator. Where resource allocation depends on the intensity of engagement, softer treatment, however well intended, can withhold the recognition that critical attention confers. My findings highlight how gendered audience reception, combined with engagement-based allocation can reproduce, even in a new technological setting, the continued gendered distribution and inequality of entrepreneurial resources and rewards.

REFERENCES

- Abraham, M. (2019). Gender-role incongruity and audience-based gender bias: An examination of networking among entrepreneurs. *Administrative Science Quarterly*, 65(1):151–180.
- Abraham, M., Botelho, T. L., and Lamont-Dobbin, G. (2024). The (re)production of inequality in evaluations: A unifying framework outlining the drivers of gender and racial differences in evaluative outcomes. *Research in Organizational Behavior*, 44:100207.
- Acker, J. (1990). Hierarchies, jobs, bodies: A theory of gendered organizations. *Gender & Society*, 4(2):139–158.
- Amarasekara, I. and Grant, W. J. (2018). Exploring the YouTube science communication gender gap: A sentiment analysis. *Public Understanding of Science*, 28(1):68–84.
- Bareket, O. and Fiske, S. T. (2023). A systematic review of the ambivalent sexism literature: Hostile sexism protects men’s power; benevolent sexism guards traditional gender roles. *Psychological Bulletin*, 149(11-12):637–698.
- Baron, R. M. and Kenny, D. A. (1986). The moderator-mediator variable distinction in social psychological research: Conceptual, strategic, and statistical considerations. *Journal of Personality and Social Psychology*, 51(6):1173–1182.
- Baumeister, R. F., Bratslavsky, E., Finkenauer, C., and Vohs, K. D. (2001). Bad is stronger than good. *Review of General Psychology*, 5(4):323–370.
- Berger, J. and Milkman, K. L. (2012). What makes online content viral? *Journal of Marketing Research*, 49(2):192–205.
- Berger, J., Sorensen, A. T., and Rasmussen, S. J. (2010). Positive effects of negative publicity: When negative reviews increase sales. *Marketing Science*, 29(5):815–827.

- Bhargava, H. K. (2022). The creator economy: Managing ecosystem supply, revenue sharing, and platform design. *Management Science*, 68(7):5233–5251.
- Bitektine, A. and Haack, P. (2015). The “macro” and the “micro” of legitimacy: Toward a multilevel theory of the legitimacy process. *Academy of Management Review*, 40(1):49–75.
- Blau, F. D. and Kahn, L. M. (2017). The gender wage gap: Extent, trends, and explanations. *Journal of Economic Literature*, 55(3):789–865. Added 2026-09-28 (Katie approved; Decision 82): reporting the gap with and without category controls, since sorting can itself reflect discrimination. Verified via OpenAlex. Zotero 7HCXHS7J.
- Bol, T., de Vaan, M., and van de Rijt, A. (2022). Gender-equal funding rates conceal unequal evaluations. *Research Policy*, 51(1):104399.
- Botelho, T. L. (2024). From audience to evaluator: When visibility into prior evaluations leads to convergence or divergence in subsequent evaluations among professionals. *Organization Science*, 35(5):1682–1703.
- Botelho, T. L. and Abraham, M. (2017). Pursuing quality: How search costs and uncertainty magnify gender-based double standards in a multistage evaluation process. *Administrative Science Quarterly*, 62(4):698–730.
- Botelho, T. L. and Poskanzer, E. J. (2026). Getting in the door versus winning it all: How gendered outcomes change across evaluation stages in entrepreneurship. *Administrative Science Quarterly*, 71(3):613–656.
- Bowers, J. S. and Pleydell-Pearce, C. W. (2011). Swearing, euphemisms, and linguistic relativity. *PLoS ONE*, 6(7):e22341.
- Boyd, R. L., Ashokkumar, A., Seraj, S., and Pennebaker, J. W. (2022). The development and psychometric properties of liwc-22. Technical report, University of Texas at Austin.

- Brady, W. J., Wills, J. A., Jost, J. T., Tucker, J. A., and Van Bavel, J. J. (2017). Emotion shapes the diffusion of moralized content in social networks. *Proceedings of the National Academy of Sciences*, 114(28):7313–7318.
- Brooks, A. W., Huang, L., Kearney, S. W., and Murray, F. E. (2014). Investors prefer entrepreneurial ventures pitched by attractive men. *PNAS*, 111(12):4427–4431.
- Cameron, L. D. (2024). The making of the “good bad” job: How algorithmic management manufactures consent through constant and confined choices. *Administrative Science Quarterly*, 69(2):458–514.
- Castilla, E. J. (2008). Gender, race, and meritocracy in organizational careers. *American Journal of Sociology*, 113(6):1479–1526.
- Castilla, E. J. and Benard, S. (2010). The paradox of meritocracy in organizations. *Administrative Science Quarterly*, 55(4):543–576.
- Cinelli, C. and Hazlett, C. (2019). Making sense of sensitivity: Extending omitted variable bias. *Journal of the Royal Statistical Society Series B (Statistical Methodology)*, 82(1):39–67.
- Coffman, K. B. and Klinowski, D. (2025). Gender and preferences for performance feedback. *Management Science*, 71(4):3497–3516.
- Collabstr (2025). 2025 influencer marketing report: Trends and statistics for influencer marketing, ugc, and the creator economy.
- Correll, S. J. (2017). SWS 2016 feminist lecture: Reducing gender biases in modern workplaces: A small wins approach to organizational change. *Gender & Society*, 31(6):725–750.
- Correll, S. J., Weisshaar, K. R., Wynn, A. T., and Wehner, J. D. (2020). Inside the black box of organizational life: The gendered language of performance assessment. *American Sociological Review*, 85(6):1022–1050.

- Covington, P., Adams, J., and Sargin, E. (2016). Deep neural networks for youtube recommendations. In *Proceedings of the 10th ACM Conference on Recommender Systems*, pages 191–198. ACM.
- Curchod, C., Patriotta, G., Cohen, L., and Neysen, N. (2019). Working for an algorithm: Power asymmetries and agency in online work settings. *Administrative Science Quarterly*, 65(3):644–676.
- Cutolo, D. and Kenney, M. (2021). Platform-dependent entrepreneurs: Power asymmetries, risks, and strategies in the platform economy. *Academy of Management Perspectives*, 35(4):584–605.
- Czopp, A. M., Kay, A. C., and Cheryan, S. (2015). Positive stereotypes are pervasive and powerful. *Perspectives on Psychological Science*, 10(4):451–463.
- Döring, N. and Mohseni, M. R. (2019). Fail videos and related video comments on YouTube: A case of sexualization of women and gendered hate speech? *Communication Research Reports*, 36(3):254–264.
- Döring, N. and Mohseni, M. R. (2020). Gendered hate speech in YouTube and YouNow comments: Results of two content analyses. *SCM Studies in Communication and Media*, 9(1).
- Eagly, A. H. (1987). *Sex Differences in Social Behavior: A Social-Role Interpretation*. Lawrence Erlbaum Associates, Hillsdale, NJ.
- Eagly, A. H. and Karau, S. J. (2002). Role congruity theory of prejudice toward female leaders. *Psychological Review*, 109(3):573–598.
- Ewens, M. and Townsend, R. R. (2020). Are early stage investors biased against women? *Journal of Financial Economics*, 135(3):653–677.
- Fiske, S. T., Cuddy, A. J. C., Glick, P., and Xu, J. (2002). A model of (often mixed) stereotype content: Competence and warmth respectively follow from perceived status and competition. *Journal of Personality and Social Psychology*, 82(6):878–902.

- Franck, G. (2018). The economy of attention. *Journal of Sociology*, 55(1):8–19.
- Frankort, H. T. W., Fernandez-Mateo, I., and Brands, R. (2026). Unequal in the spotlight: Gender differences in how serving on prominent firms affects directors' new board appointments. *Administrative Science Quarterly*, 71(2):263–305.
- Gala, K., Schwab, A., and Mueller, B. A. (2024). Star entrepreneurs on digital platforms: Heavy-tailed performance distributions and their generative mechanisms. *Journal of Business Venturing*, 39(1):106347.
- Galperin, H. (2019). “this gig is not for women”: Gender stereotyping in online hiring. *Social Science Computer Review*, 39(6):1089–1107.
- Gamache, D. L., McNamara, G., Mannor, M. J., and Johnson, R. E. (2015). Motivated to acquire? the impact of CEO regulatory focus on firm acquisitions. *Academy of Management Journal*, 58(4):1261–1282.
- Gelbach, J. B. (2016). When do covariates matter? And which ones, and how much? *Journal of Labor Economics*, 34(2):509–543.
- Glick, P. and Fiske, S. T. (2001). An ambivalent alliance: Hostile and benevolent sexism as complementary justifications for gender inequality. *American Psychologist*, 56(2):109–118.
- Gornall, W. and Strebulaev, I. A. (2025). Gender, race, and entrepreneurship: A randomized field experiment on venture capitalists and angels. *Management Science*, 71(6):5308–5327.
- Graf-Vlachy, L., Oliver, A. G., Banfield, R., König, A., and Bundy, J. (2019). Media coverage of firms: Background, integration, and directions for future research. *Journal of Management*, 46(1):36–69.
- Greenberg, J. and Mollick, E. (2016). Activist choice homophily and the crowdfunding of female founders. *Administrative Science Quarterly*, 62(2):341–374.

- Heilman, M. E. (2012). Gender stereotypes and workplace bias. *Research in Organizational Behavior*, 32:113–135.
- Hsu, G. (2006). Jacks of all trades and masters of none: Audiences' reactions to spanning genres in feature film production. *Administrative Science Quarterly*, 51(3):420–450.
- Huszár, F., Ktena, S. I., O'Brien, C., Belli, L., Schlaikjer, A., and Hardt, M. (2021). Algorithmic amplification of politics on Twitter. *Proceedings of the National Academy of Sciences*, 119(1):e2025334119.
- Jampol, L. and Zayas, V. (2020). Gendered white lies: Women are given inflated performance feedback compared with men. *Personality and Social Psychology Bulletin*, 47(1):57–69.
- Janschewitz, K. (2008). Taboo, emotionally valenced, and emotionally neutral word norms. *Behavior Research Methods*, 40(4):1065–1074.
- Jay, T. (2009). The utility and ubiquity of taboo words. *Perspectives on Psychological Science*, 4(2):153–161.
- Jennings, J. E. and Brush, C. G. (2013). Research on women entrepreneurs: Challenges to (and from) the broader entrepreneurship literature? *Academy of Management Annals*, 7(1):663–715.
- Kacperczyk, O., Younkin, P., and Rocha, V. (2023). Do employees work less for female leaders? a multi-method study of entrepreneurial firms. *Organization Science*, 34(3):1111–1133.
- Kanze, D., Huang, L., Conley, M. A., and Higgins, E. T. (2018). We ask men to win and women not to lose: Closing the gender gap in startup funding. *Academy of Management Journal*, 61(2):586–614.
- Karunakaran, A., Orlikowski, W. J., and Scott, S. V. (2022). Crowd-based accountability: Examining how social media commentary reconfigures organizational accountability. *Organization Science*, 33(1):170–193.

- Kellogg, K. C., Valentine, M. A., and Christin, A. (2020). Algorithms at work: The new contested terrain of control. *Academy of Management Annals*, 14(1):366–410.
- Khan, M. L. (2017). Social media engagement: What motivates user participation and consumption on YouTube? *Computers in Human Behavior*, 66:236–247.
- King, E. B., Botsford, W., Hebl, M. R., Kazama, S., Dawson, J. F., and Perkins, A. (2010). Benevolent sexism at work: Gender differences in the distribution of challenging developmental experiences. *Journal of Management*, 38(6):1835–1866.
- Kovács, B. and Sharkey, A. J. (2014). The paradox of publicity: How awards can negatively affect the evaluation of quality. *Administrative Science Quarterly*, 59(1):1–33.
- Kuwabara, K. and Thébaud, S. (2017). When beauty doesn't pay: Gender and beauty biases in a peer-to-peer loan market. *Social Forces*, 95(4):1371–1398.
- Lafreniere, K. C., Moore, S. G., and Fisher, R. J. (2022). The power of profanity: The meaning and impact of swear words in word of mouth. *Journal of Marketing Research*, 59(5):908–925.
- Lambrecht, A. and Tucker, C. (2019). Algorithmic bias? An empirical study of apparent gender-based discrimination in the display of STEM career ads. *Management Science*, 65(7):2966–2981.
- Lee, M. and Huang, L. (2018). Gender bias, social impact framing, and evaluation of entrepreneurial ventures. *Organization Science*, 29(1):1–16.
- MacKay, D. G., Shafto, M., Taylor, J. K., Marian, D. E., Abrams, L., and Dyer, J. R. (2004). Relations between emotion, memory, and attention: Evidence from taboo Stroop, lexical decision, and immediate memory tasks. *Memory & Cognition*, 32(3):474–488.
- Malhotra, S., Reus, T. H., Zhu, P., and Roelofsen, E. M. (2017). The acquisitive nature of extraverted CEOs. *Administrative Science Quarterly*, 63(2):370–408.
- Mears, A. (2023). Bringing bourdieu to a content farm: Social media production fields and the cultural economy of attention. *Social Media + Society*, 9(3).

- Milli, S., Carroll, M., Wang, Y., Pandey, S., Zhao, S., and Dragan, A. D. (2025). Engagement, user satisfaction, and the amplification of divisive content on social media. *PNAS Nexus*, 4(3):pgaf062.
- Nambisan, S. (2017). Digital entrepreneurship: Toward a digital technology perspective of entrepreneurship. *Entrepreneurship Theory and Practice*, 41(6):1029–1055.
- Nason, R., Vedula, S., Bothello, J., Bacq, S., and Charman, A. (2024). Sight unseen: The visibility paradox of entrepreneurship in an informal economy. *Journal of Business Venturing*, 39(2):106364.
- Nguyen, N., Hideg, I., Engel, Y., and Godart, F. (2023). Benevolent sexism and the gender gap in startup evaluation. *Entrepreneurship Theory and Practice*, 48(2):506–546.
- Nonnecke, B. and Preece, J. (2000). Lurker demographics: Counting the silent. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, pages 73–80. ACM.
- Oliver, A. G., Krause, R., Busenbark, J. R., and Kalm, M. (2017). BS in the boardroom: Benevolent sexism and board chair orientations. *Strategic Management Journal*, 39(1):113–130.
- Orlikowski, W. J. and Scott, S. V. (2014). What happens when evaluation goes online? exploring apparatuses of valuation in the travel sector. *Organization Science*, 25(3):868–891.
- Petkova, A. P., Rindova, V. P., and Gupta, A. K. (2013). No news is bad news: Sensegiving activities, media attention, and venture capital funding of new technology organizations. *Organization Science*, 24(3):865–888.
- Pillemer, J., Harrison, S., Murphy, C., and Park, Y. (2025). Audience entanglement: How independent creative workers experience the pressures of widespread appeal on digital platforms. *Administrative Science Quarterly*, 71(1):215–260.
- Pollock, N., Chapple, D., Chen, S., and D’Adderio, L. (2023). The valorising pitch: How digital start-ups leverage intermediary coverage. *Journal of Management Studies*, 60(2):346–371.

- Pollock, T. G. and Rindova, V. P. (2003). Media legitimization effects in the market for initial public offerings. *Academy of Management Journal*, 46(5):631–642.
- Preacher, K. J. and Hayes, A. F. (2008). Asymptotic and resampling strategies for assessing and comparing indirect effects in multiple mediator models. *Behavior Research Methods*, 40(3):879–891.
- Rahman, H. A. (2021). The invisible cage: Workers’ reactivity to opaque algorithmic evaluations. *Administrative Science Quarterly*, 66(4):945–988.
- Rahman, H. A., Karunakaran, A., and Cameron, L. D. (2024). Taming platform power: Taking accountability into account in the management of platforms. *Academy of Management Annals*, 18(1):251–294.
- Rathje, S., Van Bavel, J. J., and Matz, S. C. (2024). Out-group animosity increases with target group size and political distance—but only in engagement algorithms. *Science*, 386(6720):431–436.
- Rathje, S., Van Bavel, J. J., and van der Linden, S. (2021). Out-group animosity drives engagement on social media. *Proceedings of the National Academy of Sciences*, 118(26):e2024292118.
- Ridgeway, C. L. (2001). Gender, status, and leadership. *Journal of Social Issues*, 57(4):637–655.
- Ridgeway, C. L. (2008). Framed before we know it: How gender shapes social relations. *Gender & Society*, 23(2):145–160.
- Ridgeway, C. L. (2011). *Framed by Gender: How Gender Inequality Persists in the Modern World*. Oxford University Press, New York.
- Rindova, V. P., Petkova, A. P., and Kotha, S. (2007). Standing out: how new firms in emerging markets build reputation. *Strategic Organization*, 5(1):31–70.
- Rindova, V. P., Pollock, T. G., and Hayward, M. L. A. (2006). Celebrity firms: The social construction of market popularity. *Academy of Management Review*, 31(1):50–71.

- Robertson, C. E., Pröllochs, N., Schwarzenegger, K., Pärnamets, P., Van Bavel, J. J., and Feuerriegel, S. (2023). Negativity drives online news consumption. *Nature Human Behaviour*, 7(5):812–822.
- Roccapriore, A. Y. and Pollock, T. G. (2023). I don't need a degree, I've got abs: Influencer warmth and competence, communication mode, and stakeholder engagement on social media. *Academy of Management Journal*, 66(3):979–1006.
- Rozin, P. and Royzman, E. B. (2001). Negativity bias, negativity dominance, and contagion. *Personality and Social Psychology Review*, 5(4):296–320.
- Rudman, L. A. and Glick, P. (2001). Prescriptive gender stereotypes and backlash toward agentic women. *Journal of Social Issues*, 57(4):743–762.
- Rudman, L. A., Moss-Racusin, C. A., Phelan, J. E., and Nauts, S. (2012). Status incongruity and backlash effects: Defending the gender hierarchy motivates prejudice against female leaders. *Journal of Experimental Social Psychology*, 48(1):165–179.
- Seigner, B. D. C., Milanov, H., Lundmark, E., and Shepherd, D. A. (2023). Tweeting like Elon? provocative language, new-venture status, and audience engagement on social media. *Journal of Business Venturing*, 38(2):106282.
- Snellman, K. and Solal, I. (2023). Does investor gender matter for the success of female entrepreneurs? gender homophily and the stigma of incompetence in entrepreneurial finance. *Organization Science*, 34(2):680–699.
- Sobieraj, S. (2017). Bitch, slut, skank, cunt: Patterned resistance to women's visibility in digital publics. *Information, Communication & Society*, 21(11):1700–1714.
- Tausczik, Y. R. and Pennebaker, J. W. (2009). The psychological meaning of words: Liwc and computerized text analysis methods. *Journal of Language and Social Psychology*, 29(1):24–54.

- Thelwall, M., Wilkinson, D., and Uppal, S. (2009). Data mining emotion in social network communication: Gender differences in MySpace. *Journal of the American Society for Information Science and Technology*, 61(1):190–199.
- Tonoyan, V., Strohmeyer, R., and Jennings, J. E. (2024). Working for Jessica or Michael? Implications of gender stereotypes for job application intentions at technology startups. *Strategic Entrepreneurship Journal*, 19(2):256–282.
- Vanacker, T., Forbes, D. P., Knockaert, M., and Manigart, S. (2020). Signal strength, media attention, and resource mobilization: Evidence from new private equity firms. *Academy of Management Journal*, 63(4):1082–1105.
- Verwiebe, R., Osorio-Krauter, C., Fritsch, N.-S., Weißmann, S., Philipp, A., and Buder, C. (2026). Working on an algorithm-controlled platform as a content creator: What explains unequal earnings on YouTube? *European Societies*.
- Warriner, A. B., Kuperman, V., and Brysbaert, M. (2013). Norms of valence, arousal, and dominance for 13,915 English lemmas. *Behavior Research Methods*, 45(4):1191–1207.
- Weißmann, S., Philipp, A., Verwiebe, R., Osorio Krauter, C., Fritsch, N.-S., and Buder, C. (2025). Clicks, comments, consequences: Are content creators’ socio-structural and platform characteristics shaping the exposure to negative sentiment, offensive language, and hate speech on YouTube? arXiv preprint arXiv:2504.07676.
- White, H. (1980). A heteroskedasticity-consistent covariance matrix estimator and a direct test for heteroskedasticity. *Econometrica*, 48(4):817–838.
- Wotanis, L. and McMillan, L. (2014). Performing gender on YouTube: How Jenna Marbles negotiates a hostile online environment. *Feminist Media Studies*, 14(6):912–928.
- Zhang, X., Tong, S., Luo, X., Lin, Z., and Li, J. (2026). AI orchestrator: How recommendation algorithms shape complementor strategy and market equality. *Strategic Management Journal*.

- Zimmerman, M. A. and Zeitz, G. J. (2002). Beyond survival: Achieving new venture growth by building legitimacy. *Academy of Management Review*, 27(3):414–431.
- Zuckerman, E. W. (1999). The categorical imperative: Securities analysts and the illegitimacy discount. *American Journal of Sociology*, 104(5):1398–1438.
- Ørmen, J. and Gregersen, A. (2022). Towards the engagement economy: interconnected processes of commodification on youtube. *Media, Culture & Society*, 45(2):225–245.

Table 1: Summary Statistics for Dependent and Independent Variables, Analytic Sample and by Creator Gender

<i>Panel A: Analytic sample</i>						
	N	Mean	Geometric mean	SD	Min	Max
Views	9,096	58,029,502	3,623,265	544,974,751	2,287	31,662,697,819
Subscribers	9,096	183,007	25,053	846,400	2,010	31,400,000
Woman creator	9,096	0.345		0.475	0.000	1.000
Comment tone	9,096	56.20		15.55	9.42	99.00
Comment profanity	9,096	0.346		0.539	0.000	9.235
Title tone	9,096	47.32		33.04	1.00	99.00
Title profanity	9,096	0.070		0.422	0.000	10.340
Video-description tone	9,096	59.32		26.30	1.00	99.00
Video-description profanity	9,096	0.033		0.245	0.000	11.180
Engagement per view (%)	9,096	3.47	2.48	3.34	0.00	104.08
Engagement per subscriber (%)	9,096	1.407	0.351	10.831	0.000	799.656
Total comments	9,096	427	68	1,713	0	50,259
Channel age (years)	9,096	9.7		4.9	0.3	19.9
Video count	9,096	1,267	542	21,404	6	2,036,022
HD video	9,096	0.971		0.169	0.000	1.000
Non-subscriber trailer	9,096	0.595		0.491	0.000	1.000

<i>Panel B: By creator gender</i>					
	Men	Women	Difference	p-value	N
Views, arithmetic mean	59,285,183	55,644,229	-6.1%	0.773	9,096
Views, geometric mean	4,178,436	2,763,693	-33.9%***	<0.001	9,096
at the same channel age and video count	3,904,046	3,144,310	-19.5%***	<0.001	9,096
Subscribers, arithmetic mean	193,319	163,418	-15.5%	0.112	9,096
Subscribers, geometric mean	26,017	23,319	-10.4%**	0.005	9,096
at the same channel age and video count	25,106	24,953	-0.6%	0.873	9,096
Engagement per view, twelve videos (%), arithmetic mean	3.25	3.90	20.2%***	<0.001	9,096
Engagement per view, twelve videos (%), geometric mean	2.34	2.79	19.2%***	<0.001	9,087
at the same channel age and video count	2.34	2.78	18.8%***	<0.001	9,087
Engagement per view, six videos (%), arithmetic mean	4.51	5.73	27.1%***	<0.001	9,096
Engagement per view, six videos (%), geometric mean	3.06	3.91	27.6%***	<0.001	9,076
at the same channel age and video count	3.09	3.85	24.7%***	<0.001	9,076
Engagement per subscriber, six videos (%), arithmetic mean	1.430	1.363	-4.7%	0.717	9,096
Engagement per subscriber, six videos (%), geometric mean	0.314	0.435	38.7%***	<0.001	9,076
at the same channel age and video count	0.335	0.383	14.2%***	<0.001	9,076
Comment tone	52.71	62.82	10.10***	<0.001	9,096
Comment profanity	0.427	0.194	-0.233***	<0.001	9,096
Title tone	42.87	55.76	12.89***	<0.001	9,096
Title profanity	0.087	0.038	-0.050***	<0.001	9,096
Video-description tone	55.68	66.23	10.55***	<0.001	9,096
Video-description profanity	0.040	0.020	-0.020***	<0.001	9,096
Total comments (geometric mean)	58	67	15.0%**	0.001	9,096
Channel age (years)	10.03	9.11	-0.91***	<0.001	9,096
Video count (geometric mean)	589	466	-20.8%***	<0.001	9,096
HD video	0.970	0.972	0.002	0.590	9,096
Non-subscriber trailer	0.616	0.557	-0.059***	<0.001	9,096

Notes: Randomly drawn channels from the Infludata frame of U.S.-based individual creators with at least 2,000 subscribers; the analytic sample is the 9,096 channels with views, channel age, video count, and scored comment language (3,137 women, 5,959 men). Creator gender was hand-coded and verified by a second coder ($\kappa = 0.96$); every other measure comes from the YouTube Data API or from text collected through it. Unstandardized channel-level values. Geometric means (exp of the mean of the logged variable; the scale on which the models are estimated) are reported for right-skewed quantities; the geometric-mean difference is the percent by which women's geometric mean differs from men's (t -test on the logged variable, robust standard errors), and the row "at the same channel age and video count" gives the geometric means predicted at the sample means of channel age and ln(video count). Median views: 3,121,150 (men) versus 1,990,970 (women). Comment tone is LIWC-22 Tone (1 to 99; higher is more positive) for each scored comment, averaged to the channel; comments with no positive or negative tone words take LIWC-22's own value of 20.23, and comments with no words are excluded. Comment profanity is LIWC-22 Swear, the percentage of words that are swear words. Title and video-description tone and profanity are the same LIWC-22 measures applied to each channel's sampled video titles and video descriptions; where the text contains no tone words, tone takes LIWC-22's value of 20.23, and empty descriptions take tone 20.23 and profanity 0 (flagged by indicators in the models). Engagement per view is (likes + comments)/views $\times$ 100 across each channel's twelve sampled videos; the six-video rows compute engagement per view, (likes + comments)/views $\times$ 100, and engagement per subscriber, mean (likes + comments) per video divided by subscribers $\times$ 100, on the same three newest and three oldest videos. Total comments is the API comment count on those six videos. Differences are women – men; arithmetic-mean differences for skewed rows are percent differences. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table 2: Example comments across the comment-tone distribution

Comment	Creator	Category	Tone
<i>Panel A: Negative</i>			
“You suck SO BAD STOP.”	Woman	Humour	Negative
“SO BORING BORING LITTLE GIRLS ADVICES”	Woman	Fashion	Negative
“Awful lips. I stopped watching.”	Woman	Health	Negative
“Worst video ever seen”	Woman	Lifestyle	Negative
“Holding coffee mug all the time while interviewing made him look so stupid.”	Man	Technology	Negative
“I hate your videos. You’re so dumb.”	Man	Gaming	Negative
“WTF??? So boring! Get a life man!”	Man	Pets	Negative
<i>Panel B: Neutral</i>			
“What wig is this?”	Woman	Lifestyle	Neutral
“is this Open D?”	Woman	Pop music	Neutral
“Where to buy one?”	Man	Vehicles	Neutral
“15 pro vs 15 s pro?”	Man	Technology	Neutral
<i>Panel C: Positive</i>			
“That bag is awesome!!! You do amazing work!!!”	Woman	Hobbies	Positive
“Just awesome!!! The Best video ever. You are strong !!!”	Woman	Tourism	Positive
“AMAZING WORK! You inspire me so much, I hope one day to reach your level!”	Woman	Hobbies	Positive
“Your channel is fantastic !! The best on YouTube. THANK YOU !”	Man	Television	Positive
“That’s a Masterpiece! Thank you for the wonderful video!”	Man	Tourism	Positive
“Brilliant episode man, your editing is flawless! Absolutely amazing work! Love it”	Man	Tourism	Positive

Notes: Comments from the analysis corpus, drawn from the extremes and the middle of the per-comment tone distribution and screened for brevity and clarity; text is reproduced verbatim, including spelling, except emoji, which are omitted. Tone classifies each comment by its LIWC-22 Tone score (scale 1 to 99): the negative examples score at or near 1, the positive examples score 99, and the neutral examples contain no words from the LIWC-22 tone dictionaries and score approximately zero on the VADER compound sentiment scale. The channel-level measure used in the models averages the LIWC-22 Tone score over all sampled comments on a channel. Creator indicates the gender of the creator whose channel received the comment; Category is the channel’s content category (Gaming groups the game-related categories).

Table 3: OLS Regressions Predicting Logged Channel Views from Audience Reception, Across Creators

	(Model 1)	(Model 2)	(Model 3)	(Model 4)
Comment tone			-0.377*** (0.017)	-0.427*** (0.019)
Comment profanity			0.076*** (0.017)	0.046** (0.017)
Channel age (years)	0.036*** (0.004)	0.036*** (0.003)	0.046*** (0.003)	0.044*** (0.003)
ln(Video count)	0.716*** (0.018)	0.890*** (0.014)	0.829*** (0.014)	0.847*** (0.014)
ln(Total comments)		0.620*** (0.008)	0.597*** (0.008)	0.594*** (0.008)
Constant	10.242*** (0.111)	6.587*** (0.103)	6.972*** (0.102)	6.701*** (0.168)
Content-category fixed effects (40)	No	No	No	Yes
Observations	9,096	9,096	9,096	9,096
R^2	0.174	0.505	0.541	0.560

Notes: Randomly drawn channels from the Infludata frame of U.S.-based individual creators with at least 2,000 subscribers; the analytic sample is the 9,096 channels with views, channel age, video count, and scored comment language (3,137 women, 5,959 men). Creator gender was hand-coded and verified by a second coder ($\kappa = 0.96$); every other measure comes from the YouTube Data API or from text collected through it. Dependent variable: ln(cumulative channel views). Robust standard errors in parentheses. Comment tone is LIWC-22 Tone (1 to 99; higher is more positive) for each scored comment, averaged to the channel; comments with no positive or negative tone words take LIWC-22's own value of 20.23, and comments with no words are excluded. Comment profanity is LIWC-22 Swear, the percentage of words that are swear words. Both are standardized, so a coefficient is the change in ln(views) associated with a one standard deviation change. Channel age (years) and ln(video count) are the production controls. ln(total comments) is the natural log of the API comment counts on each channel's three newest and three oldest videos; it holds constant how much reception a channel receives so that comment tone and profanity capture the type of reception at a given amount. No model includes creator gender. Estimates are associations from cross-sectional channel-level data. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table 4: OLS Regressions Predicting Comment Tone and Comment Profanity from Creator Gender

	Comment tone			Comment profanity		
	(Model 1)	(Model 2)	(Model 3)	(Model 4)	(Model 5)	(Model 6)
Woman creator	0.650*** (0.022)	0.656*** (0.021)	0.461*** (0.023)	-0.432*** (0.019)	-0.429*** (0.019)	-0.317*** (0.021)
Channel age (years)		0.031*** (0.002)	0.024*** (0.002)		-0.001 (0.002)	0.002 (0.002)
ln(Video count)		-0.133*** (0.009)	-0.085*** (0.009)		0.023* (0.011)	0.002 (0.011)
ln(Total comments)		-0.063*** (0.005)	-0.046*** (0.005)		0.012 (0.006)	-0.002 (0.006)
Constant	-0.224*** (0.012)	0.568*** (0.065)	-0.113 (0.085)	0.149*** (0.014)	-0.043 (0.079)	0.488*** (0.144)
Content-category fixed effects (40)	No	No	Yes	No	No	Yes
Observations	9,096	9,096	9,096	9,096	9,096	9,096
R^2	0.095	0.145	0.289	0.042	0.043	0.150

Notes: Randomly drawn channels from the Infludata frame of U.S.-based individual creators with at least 2,000 subscribers; the analytic sample is the 9,096 channels with views, channel age, video count, and scored comment language (3,137 women, 5,959 men). Creator gender was hand-coded and verified by a second coder ($\kappa = 0.96$); every other measure comes from the YouTube Data API or from text collected through it. Dependent variables: standardized comment tone (Models 1 to 3) and standardized comment profanity (Models 4 to 6), so the Woman creator coefficient is the difference in standard deviations. Comment tone is LIWC-22 Tone (1 to 99; higher is more positive) for each scored comment, averaged to the channel; comments with no positive or negative tone words take LIWC-22's own value of 20.23, and comments with no words are excluded. Comment profanity is LIWC-22 Swear, the percentage of words that are swear words. In raw units the difference is 10.1 tone points (62.8 versus 52.7) and 0.23 percentage points of swear words (0.19 versus 0.43 percent). Robust standard errors in parentheses. Channel age (years) and ln(video count) are the production controls. ln(total comments) is the natural log of the API comment counts on each channel's three newest and three oldest videos; it holds constant how much reception a channel receives so that comment tone and profanity capture the type of reception at a given amount. Estimates are associations from cross-sectional channel-level data. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table 5: OLS Regressions Predicting Logged Channel Views from Creator Gender, Adding Production, Creator Content, Comment Volume, and Audience Reception

	(Model 1)	(Model 2)	(Model 3)	(Model 4)	(Model 5)	(Model 6)
<i>Panel A: No fixed effects</i>						
Woman creator	-0.413*** (0.046)	-0.216*** (0.042)	-0.126** (0.043)	-0.237*** (0.034)	-0.169*** (0.035)	-0.004 (0.034)
Comment tone						-0.402*** (0.019)
Comment profanity					0.179*** (0.020)	0.079*** (0.018)
ln(Total comments)				0.622*** (0.008)	0.621*** (0.008)	0.602*** (0.008)
Title tone			-0.155*** (0.023)	0.002 (0.018)	0.014 (0.018)	0.087*** (0.017)
Title profanity			0.020 (0.028)	-0.008 (0.020)	-0.020 (0.019)	-0.020 (0.020)
Video-description tone			-0.100*** (0.023)	-0.074*** (0.018)	-0.068*** (0.018)	-0.004 (0.018)
Video-description profanity			-0.022 (0.020)	-0.007 (0.016)	-0.015 (0.016)	-0.021 (0.017)
Channel age (years)		0.035*** (0.004)	0.035*** (0.004)	0.035*** (0.003)	0.035*** (0.003)	0.047*** (0.003)
ln(Video count)		0.709*** (0.018)	0.708*** (0.018)	0.877*** (0.014)	0.874*** (0.014)	0.826*** (0.014)
Content-blank indicators	No	No	Yes	Yes	Yes	Yes
Content-category fixed effects (40)	No	No	No	No	No	No
Observations	9,096	9,096	9,096	9,096	9,096	9,096
R ²	0.009	0.176	0.186	0.511	0.518	0.543
<i>Panel B: Content-category fixed effects</i>						
Woman creator	-0.400*** (0.053)	-0.226*** (0.049)	-0.166*** (0.049)	-0.203*** (0.038)	-0.163*** (0.039)	-0.027 (0.038)
Comment tone						-0.439*** (0.020)
Comment profanity					0.133*** (0.019)	0.047** (0.018)
ln(Total comments)				0.614*** (0.008)	0.615*** (0.008)	0.597*** (0.008)
Title tone			-0.157*** (0.023)	-0.009 (0.018)	-0.002 (0.018)	0.058*** (0.017)
Title profanity			0.001 (0.027)	-0.020 (0.020)	-0.027 (0.019)	-0.026 (0.020)
Video-description tone			-0.096*** (0.023)	-0.078*** (0.018)	-0.074*** (0.018)	-0.015 (0.017)
Video-description profanity			-0.042 (0.021)	-0.023 (0.016)	-0.026 (0.017)	-0.030 (0.018)
Channel age (years)		0.037*** (0.004)	0.036*** (0.004)	0.034*** (0.003)	0.034*** (0.003)	0.045*** (0.003)
ln(Video count)		0.693*** (0.018)	0.700*** (0.018)	0.882*** (0.015)	0.882*** (0.015)	0.845*** (0.014)
Content-blank indicators	No	No	Yes	Yes	Yes	Yes
Content-category fixed effects (40)	Yes	Yes	Yes	Yes	Yes	Yes
Observations	9,096	9,096	9,096	9,096	9,096	9,096
R ²	0.060	0.218	0.227	0.532	0.535	0.562

Notes: Randomly drawn channels from the Infludata frame of U.S.-based individual creators with at least 2,000 subscribers; the analytic sample is the 9,096 channels with views, channel age, video count, and scored comment language (3,137 women, 5,959 men). Creator gender was hand-coded and verified by a second coder ($\kappa = 0.96$); every other measure comes from the YouTube Data API or from text collected through it. Dependent variable: ln(cumulative channel views). Robust standard errors in parentheses. Models add, in order, the production controls (Model 2), creator content (Model 3: standardized tone and profanity of the channel's sampled video titles and video descriptions, with indicators for titles or descriptions containing no tone words and for empty descriptions), comment volume (Model 4), comment profanity (Model 5), and comment tone (Model 6). Comment tone is LIWC-22 Tone (1 to 99; higher is more positive) for each scored comment, averaged to the channel; comments with no positive or negative tone words take LIWC-22's own value of 20.23, and comments with no words are excluded. Comment profanity is LIWC-22 Swear, the percentage of words that are swear words. Comment tone, profanity, and the content measures are standardized. Channel age (years) and ln(video count) are the production controls. ln(total comments) is the natural log of the API comment counts on each channel's three newest and three oldest videos; it holds constant how much reception a channel receives so that comment tone and profanity capture the type of reception at a given amount. Because the change in the creator-gender coefficient depends on the order in which blocks enter, Table 6 reports an order-invariant (Gelbach) decomposition, in full in Appendix Table A27. Estimates are associations from cross-sectional channel-level data. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table 6: Gelbach Decomposition of the Gender Gap in Logged Channel Views: The Share of the Gap Each Block Accounts For

	(1) From the raw gap	(2) From the production-adjusted gap	(3) From the production-adjusted gap within content categories
Starting model	Creator gender only	+ channel age, video count	+ channel age, video count, 40 category fixed effects
Full model (all columns)	Every block, 40 content-category indicators, robust standard errors		
<i>Contribution of each block to the change in the creator-gender coefficient (log points), robust SE, [share of the starting gap]</i>			
Channel age and video count	-0.233*** (0.022) [56.4%]	(in start)	(in start)
Content category (40)	0.052** (0.019) [-12.5%]	0.050* (0.019) [-22.9%]	(in start)
Creator content (titles, video descriptions)	0.024** (0.009) [-5.9%]	0.024** (0.009) [-11.3%]	0.017* (0.007) [-7.3%]
Comment volume	0.083** (0.025) [-20.0%]	0.045 (0.025) [-20.7%]	0.000 (0.029) [-0.2%]
Comment profanity	-0.021** (0.008) [5.0%]	-0.020** (0.008) [9.5%]	-0.015** (0.006) [6.7%]
Comment tone	-0.281*** (0.017) [67.9%]	-0.281*** (0.017) [130.0%]	-0.199*** (0.014) [88.0%]
HD video	0.001 (0.001) [-0.2%]	0.000 (0.001) [0.0%]	-0.002 (0.001) [0.7%]
Non-subscriber trailer	-0.015*** (0.003) [3.6%]	-0.009** (0.003) [4.3%]	-0.003 (0.003) [1.5%]
Starting creator-gender coefficient	-0.413	-0.216	-0.226
Percent fewer views on women's channels	33.9	19.5	20.2
Ending creator-gender coefficient (full model)	-0.024 ($p = 0.53$)	-0.024 ($p = 0.53$)	-0.024 ($p = 0.53$)
Total accounted for (%)	94.2	89.0	89.5
Left over (%)	5.8	11.0	10.5
Observations	9,096	9,096	9,096

Notes: Gelbach (2016) decomposition of the creator-gender gap in logged cumulative channel views on 9,096 channels. Each entry is the women-minus-men difference in the block (holding the starting model's variables fixed) multiplied by the block's coefficient in the full model, so the entries do not depend on the order in which blocks are added. Robust standard errors are in parentheses and the share of the starting coefficient is in brackets. Positive shares indicate that equalizing the block between women's and men's channels narrows the gap; negative shares indicate widening. Shares sum with the share left over to 100 percent and can exceed 100 percent or be negative if blocks offset each other. Columns differ by starting gap: (1) raw gap with creator gender only (-0.413, 33.9 percent fewer views), where channel age and $\ln(\text{video count})$ form a block; (2) production-adjusted gap adding channel age and $\ln(\text{video count})$ (-0.216, 19.5 percent); and (3) production-adjusted gap within content categories adding 40 content-category fixed effects (-0.226, 20.2 percent), where content category is part of the starting model. In columns 1 and 2, content category is its own block (39 category indicators, reported as one block total). All full models include every block plus category indicators, with the final creator-gender coefficient not distinguishable from zero (-0.024, $p = 0.53$). Blocks include channel age and $\ln(\text{video count})$; content category; creator content (standardized LIWC-22 Tone and profanity of sampled titles and descriptions, plus indicators for no tone words and empty descriptions); comment volume ($\ln$ total comments on newest and oldest videos); standardized comment profanity; standardized comment tone (LIWC-22 Tone, comments with no tone words at LIWC-22's value of 20.23); HD video indicator; and non-subscriber trailer indicator. Standard errors are heteroskedasticity-robust. This is an accounting of coefficient changes, not causal effects. Appendix Table A27 reports further starting points and versions with transcripts. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$ (for each block's contribution).

Table 7: OLS Regressions Predicting Logged Video Views from Audience Reception and from the Creator’s Own Transcript Language, Within Creators

Panel A: The comments a video received

	(Model 1)	(Model 2)	(Model 3)	(Model 4)
Comment tone (VADER compound)	-0.434*** (0.018)	-0.230*** (0.007)		
Comment profanity (token list)	-0.024* (0.010)	0.030*** (0.004)		
Comment tone (LIWC-22 Tone)			-0.447*** (0.018)	-0.231*** (0.007)
Comment profanity (LIWC-22 Swear)			0.050*** (0.009)	0.024*** (0.004)
Videos	135,148	135,148	131,859	131,859
Channels	7,157	7,157	7,153	7,153
Adjusted R^2	0.099	0.680	0.107	0.681

Panel B: The creator’s own transcript language

	(Model 5)	(Model 6)		
Own transcript tone (LIWC-22 Tone)	-0.079*** (0.021)	-0.029 (0.017)		
Own transcript profanity (LIWC-22 Swear)	0.051* (0.021)	-0.015 (0.014)		
Videos	20,872	20,872		
Channels	7,425	7,425		
Adjusted R^2	0.055	0.788		
Channel fixed effects	No	Yes	No	Yes

Notes: Unit of analysis: the video. Dependent variable: $\ln(\text{video views})$. Standard errors clustered on channel in parentheses. Panel A uses the comments each video received, scored two ways on the same panel: the VADER compound score with a fixed profanity token list (Models 1 and 2) and LIWC-22 Tone and Swear per comment (Models 3 and 4), each averaged to the video and standardized; higher tone is more positive. The panel holds every video with at least ten comments on channels contributing at least two such videos, built from the 2025 comment pulls and a 2026 per-video pull of up to 100 comments per video (135,148 videos on 7,157 channels; 2,512 women and 4,645 men); Models 3 and 4 use the videos with at least ten LIWC-scored comments. Models control $\ln(\text{video age in days})$, its square, and $\ln(\text{duration})$. Panel B uses LIWC-22 Tone and Swear on the creator’s own video transcripts, on the transcript sample (20,872 videos on 7,425 channels), controlling $\ln(\text{video age})$ and its square. Odd-numbered models pool across channels; even-numbered models add channel fixed effects, which absorb every time-invariant creator characteristic, creator gender included. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table 8: Rival Explanations for the Gender Gap in Views: Tests and Results

Rival explanation	Test	Result	Verdict
Women make different content: titles and descriptions	Add the creator's own title and video-description tone and profanity to the volume-adjusted model (all 9,096 channels; text with no tone words scored at LIWC-22's value for such text and flagged; Appendix Table A02)	Adding title and description language to the volume-adjusted model moves the gap from -0.263 to -0.237; the gap remains significant ($p < 0.001$).	Not creator content: holding titles and descriptions equal widens the gap (share -7.3 percent of the 20 percent gap in the Gelbach decomposition; Table 6)
Women make different content: all four content elements	Titles and video descriptions, transcripts, and thumbnails entered alongside comment reception as parallel mediators, fourteen in one equation with category fixed effects (6,727 channels; Appendix Table A04)	Comment reception accounts for 92.9% of the volume-adjusted gap within content categories (joint mediation); the three creator-side elements together account for -21.0% (a negative share means holding the element constant would make the gap larger); the remaining direct gap is not distinguishable from zero ($p = 0.102$)	Not the creator's own content: what audiences write carries the gap, not what creators post

Table 8, continued

Rival explanation	Test	Result	Verdict
Women sort into lower-visibility content categories	Content-category block of the Gelbach decomposition (Table 6) is the main test, followed by the addition of 40 content-category indicators to the volume-adjusted model and an Oaxaca–Blinder decomposition of the raw gap over the same categories (Appendix Tables A03 and A20).	Gelbach decomposition shows holding content category equal widens the 19 percent gap (category share -22.9 percent); with category indicators the volume-adjusted gap moves from -0.263 to -0.226; the Oaxaca–Blinder composition term is 0.066 ($p = 0.161$), not distinguishable from zero, corroborating the Gelbach result.	Not sorting: category membership does not account for the gap (holding category equal widens it)
Women’s softer reception reflects who comments, not how commenters treat women	Comment tone from male-coded and from female-coded commenters compared by creator gender; the share of female-coded commenters added as a control (Appendix Tables A12 to A14)	Women’s channels receive 10.0 more tone points from male-coded commenters and 9.8 more from female-coded commenters; with the female share controlled the gap is -0.191, and it falls to -0.020 ($p = 0.625$) once tone and profanity enter	Not commenter composition: both men and women write more positively on women’s channels
Audiences subscribe to and engage with women’s channels less	The subscriber gap raw and with the production controls; engagement per view (likes plus comments per view) by creator gender (Table 1; Appendix Table A17, Panel F)	Subscribers: -10.4% raw, -0.6% once channel age and video count are held constant ($p = 0.873$); engagement per view 3.90% on women’s channels versus 3.25% on men’s	Not fewer subscribers or less engagement: the disadvantage is specific to views

Table 8, continued


Rival explanation	Test	Result	Verdict
Women use promotional tactics less: the non-subscriber trailer	The trailer indicator added to the volume-adjusted model, and interacted with creator gender (Appendix Table A16)	55.7% of women's channels run a trailer versus 61.6% of men's; a trailer is associated with 0.279 more logged views; adding it moves the gap from -0.263 to -0.252; the return does not differ by gender (interaction -0.052, $p = 0.434$)	Not promotional tactics: the same return for both, and the gap barely moves
The views caused the comments	Two-wave lagged design: first-wave comment tone predicting second-wave views, conditional on first-wave views and the controls (9,054 channels; Appendix Tables A05 to A07)	First-wave tone -0.032 on later logged views ($p < 0.001$); the creator-gender coefficient on later views is not distinguishable from zero	Consistent with reception preceding visibility; does not remove unmeasured confounding
Women divert their audience to other platforms	Indicator for links to other platforms in the channel description, added as a control and interacted with creator gender (Appendix Table A16)	20.7% of women's channels link out versus 16.4% of men's; with the indicator the gap is -0.265; the interaction is not distinguishable from zero ($p = 0.321$)	Not off-platform diversion: the gap is unchanged
Backlash against women in female-typed categories	Female-typed categories are those above the median share of female-typed video titles; creator gender interacted with comment tone and the female-typed indicator (Appendix Table A15, Model 4)	Three-way interaction -0.038 ($p = 0.611$)	Not category backlash: the tone slope does not differ for women in female-typed categories

Table 8, continued

Rival explanation	Test	Result	Verdict
Women receive shorter, less substantial comments	ln(mean words per comment) added to the volume-adjusted model (Appendix Table A08)	Comments on women's channels average 19.1 words versus 20.7 on men's; holding length constant moves the gap from -0.263 to -0.296, larger rather than smaller	Not comment length: it works in women's favour, so the gap is not a length artefact


Notes: Every row is estimated on the 9,096 channels with complete comment-language measures (the analytic sample of Tables 1 and 3 to 5), or the subset of them with the measure named in the Test column; the lagged row uses the 9,054 of these channels observed at the second wave. Percentages of channels, engagement per view, and words per comment are computed on the same 9,096 channels as Table 1. Dependent variable: ln(cumulative channel views), except ln(subscribers) in the subscriber row and second-wave ln(views) in the lagged row. The volume-adjusted model includes channel age (years), ln(video count), and ln(total comments); the production controls are channel age (years) and ln(video count). Comment tone is LIWC-22 Tone averaged over each channel's scored comments (comments with no words excluded). Robust standard errors throughout. Each row's full model, N, and standard errors appear in the appendix table named in the Test column. Estimates are associations from cross-sectional channel-level data.

Figure 1: Two real comment sections in the same content category with similar comment volume




Notes: Each bar is the share of a channel's sampled comments at that VADER sentiment score; vertical lines mark each channel's mean. Both are Health channels with about 2,000 sampled comments. Channel A (LIWC-22 comment tone 39) has ln(views) 18.9; Channel B (comment tone 81) has 17.7.

Figure 2: Mean logged views by decile of comment-section tone




Notes: The 9,096 channels of the analytic sample are sorted into deciles of mean LIWC-22 comment tone. Each point is the decile's mean $\ln(\text{views})$; the band is the 95 percent confidence interval of that mean. Descriptive; no model is fitted.

Figure 3: The conceptual mechanism: creator gender, expressed reception, amplification, and visibility



Notes: Conceptual diagram; no estimates plotted. Amplification runs through audiences passing content along and the recommendation system; the direct path from creator gender to visibility is shown as not distinguishable from zero.

Figure 4: The creator-gender coefficient across the models of Table 5



Notes: Point estimates with 95 percent confidence intervals across the six models of Table 5, Panel A and Panel B side by side. $N = 9,096$ in every model.

ONLINE APPENDIX

Supplementary tables, numbered A01 to A29 to match references in the text. All appendix tables use the 9,096-channel analytic sample unless otherwise noted in their table notes for subsets or alternative samples. Appendix Tables A18 and A28 use the drawn sample of 9,232 channels.

Table A01: OLS Regressions Predicting Logged Channel Views from Comment Tone and Comment Profanity Entered Separately and Together

	(Model 1)	(Model 2)	(Model 3)	(Model 4)
Comment tone	-0.403*** (0.016)		-0.377*** (0.017)	-0.427*** (0.019)
Comment profanity		0.197*** (0.019)	0.076*** (0.017)	0.046** (0.017)
Channel age (years)	0.047*** (0.003)	0.036*** (0.003)	0.046*** (0.003)	0.044*** (0.003)
ln(Video count)	0.828*** (0.014)	0.883*** (0.014)	0.829*** (0.014)	0.847*** (0.014)
ln(Total comments)	0.596*** (0.008)	0.618*** (0.008)	0.597*** (0.008)	0.594*** (0.008)
Observations	9,096	9,096	9,096	9,096
R^2	0.539	0.514	0.541	0.560
Content-category fixed effects (40)	No	No	No	Yes

Notes: Randomly drawn channels from the Infludata frame of U.S.-based individual creators with at least 2,000 subscribers; the analytic sample is the 9,096 channels with views, channel age, video count, and scored comment language (3,137 women, 5,959 men). Creator gender was hand-coded and verified by a second coder ($\kappa = 0.96$); every other measure comes from the YouTube Data API or from text collected through it. Dependent variable: ln(cumulative channel views); the 9,096 channels with complete comment-language measures. Robust standard errors in parentheses. Comment tone and profanity are the standardized LIWC-22 Tone and Swear measures from the audience comment corpus; higher tone is more positive. Channel age (years) and ln(video count) are the production controls. ln(total comments) is the natural log of the API comment counts on each channel's three newest and three oldest videos; it holds constant how much reception a channel receives so that comment tone and profanity capture the type of reception at a given amount. No model includes creator gender. Estimates are associations from cross-sectional channel-level data. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table A02: OLS Regressions Predicting Logged Channel Views from Creator Gender, Adding the Creator's Own Title and Video-Description Language

	(Model 1)	(Model 2)	(Model 3)	(Model 4)
Woman creator	-0.263*** (0.033)	-0.256*** (0.034)	-0.236*** (0.034)	-0.237*** (0.034)
Title tone		-0.019 (0.017)		0.002 (0.018)
Title profanity		-0.008 (0.017)		-0.008 (0.020)
Video-description tone			-0.073*** (0.017)	-0.074*** (0.018)
Video-description profanity			-0.010 (0.013)	-0.007 (0.016)
Channel age (years)	0.034*** (0.003)	0.034*** (0.003)	0.035*** (0.003)	0.035*** (0.003)
ln(Video count)	0.882*** (0.014)	0.882*** (0.014)	0.878*** (0.014)	0.877*** (0.014)
ln(Total comments)	0.621*** (0.008)	0.620*** (0.008)	0.622*** (0.008)	0.622*** (0.008)
Observations	9,096	9,096	9,096	9,096
R^2	0.509	0.509	0.511	0.511
Indicators for text with no tone words	No	Yes	Yes	Yes

Notes: Randomly drawn channels from the Infludata frame of U.S.-based individual creators with at least 2,000 subscribers; the analytic sample is the 9,096 channels with views, channel age, video count, and scored comment language (3,137 women, 5,959 men). Creator gender was hand-coded and verified by a second coder ($\kappa = 0.96$); every other measure comes from the YouTube Data API or from text collected through it. Dependent variable: ln(cumulative channel views); the 9,096 channels of the analytic sample. Robust standard errors in parentheses. Title and video-description tone and profanity are standardized LIWC-22 Tone and Swear on the creator's own text. Titles or video descriptions with no tone words are scored at LIWC-22's value for such text (20.23) and flagged with an indicator; empty video descriptions are scored the same way, with profanity 0, and flagged with their own indicator; indicator coefficients are not shown. Channel age (years) and ln(video count) are the production controls. ln(total comments) is the natural log of the API comment counts on each channel's three newest and three oldest videos; it holds constant how much reception a channel receives so that comment tone and profanity capture the type of reception at a given amount. Estimates are associations from cross-sectional channel-level data. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table A03: Oaxaca–Blinder Decomposition of the Raw Gender Gap in Logged Channel Views over Content Categories and Channel Characteristics

	(Model 1)	(Model 2)	(Model 3)
Men (mean ln views)	15.245*** (0.028)	15.245*** (0.028)	15.245*** (0.028)
Women (mean ln views)	14.832*** (0.037)	14.832*** (0.037)	14.832*** (0.037)
Difference (men minus women)	0.413*** (0.046)	0.413*** (0.046)	0.413*** (0.046)
Endowments (composition)	0.066 (0.047)	0.152*** (0.033)	0.243*** (0.046)
Coefficients (unexplained)	0.452*** (0.062)	0.264*** (0.033)	0.259*** (0.045)
Interaction	-0.105 (0.061)	-0.003 (0.009)	-0.089* (0.043)
Observations	9,096	9,096	9,096
Decomposed over	40 categories	Channel controls	Both

Notes: Randomly drawn channels from the Infludata frame of U.S.-based individual creators with at least 2,000 subscribers; the analytic sample is the 9,096 channels with views, channel age, video count, and scored comment language (3,137 women, 5,959 men). Creator gender was hand-coded and verified by a second coder ($\kappa = 0.96$); every other measure comes from the YouTube Data API or from text collected through it. Dependent variable: $\ln(\text{cumulative channel views})$; the 9,096 channels with complete comment-language measures. Robust standard errors in parentheses. Three-fold decomposition by creator gender; the difference is the men-minus-women gap in mean logged views; endowments are the part attributable to differences in the covariates, weighted by women's coefficients; coefficients the part attributable to different returns to them, evaluated at women's covariate means; interaction the remainder. Model 1 decomposes over the 40 content-category indicators, Model 2 over channel age (years), $\ln(\text{video count})$, and $\ln(\text{total comments})$, Model 3 over both. Estimates are associations from cross-sectional channel-level data. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table A04: Parallel Mediation of the Volume-Adjusted Gender Gap in Logged Channel Views Through Audience Comment Language and Three Creator-Controlled Content Elements, Bootstrapped

	Estimate	95% percentile interval	Share of the gap
Indirect through comment tone and profanity (audience)	-0.2382	[-0.2741, -0.2046]	92.9%
Indirect through title and video-description tone and profanity (creator)	0.0091	[-0.0058, 0.0241]	-3.6%
Indirect through transcript tone and profanity (creator)	0.0150	[0.0060, 0.0263]	-5.8%
Indirect through thumbnail expression, clickbait, capitals, and overlay text (creator)	0.0297	[0.0083, 0.0512]	-11.6%
Direct association of creator gender	-0.0719	[-0.1579, 0.0130]	28.0%
Volume-adjusted gap (total)	-0.2564	[-0.3405, -0.1724]	100.0%

Notes: Randomly drawn channels from the Infludata frame of U.S.-based individual creators with at least 2,000 subscribers; the analytic sample is the 9,096 channels with views, channel age, video count, and scored comment language (3,137 women, 5,959 men). N = 6,727 channels with all fourteen measures present. Fourteen standardized mediators enter one view equation with 40 content-category fixed effects: LIWC-22 Tone and Swear on audience comments; on video titles and video descriptions; on the creator's own transcripts; and six thumbnail features (positive and aroused facial expression, clickbait score, share of capital letters, overlay-text Tone and Swear). Each indirect quantity sums the products of the creator-gender coefficient in the mediator equation and the mediator's coefficient in the view equation over the mediators in that block. Nonparametric bootstrap, 5,000 replications, seed 20260615, percentile intervals in brackets. Every equation contains channel age (years), ln(video count), and ln(total comments), the natural log of the API comment counts on each channel's three newest and three oldest videos. A negative share means holding the block constant would make the gap larger. Estimates are associations from cross-sectional channel-level data.

Table A05: OLS Regressions Predicting Second-Wave Logged Channel Views (February 2026) from First-Wave Comment Language, Conditional on First-Wave Views

	(Model 1)	(Model 2)	(Model 3)
Woman creator	-0.008 (0.012)	0.014 (0.014)	0.017 (0.014)
First-wave comment tone		-0.036*** (0.007)	-0.032*** (0.008)
First-wave comment profanity			0.013 (0.008)
First-wave ln(views)	0.964*** (0.005)	0.958*** (0.005)	0.958*** (0.005)
Channel age (years)	-0.008*** (0.001)	-0.006*** (0.001)	-0.007*** (0.001)
ln(Video count)	0.009 (0.009)	0.009 (0.009)	0.010 (0.009)
ln(Total comments)	0.017** (0.005)	0.018*** (0.005)	0.019*** (0.005)
Observations	9,054	9,054	9,054
R^2	0.936	0.937	0.937

Notes: Channels of the 9,096-channel analytic sample with a second-wave view count. Dependent variable: ln(cumulative channel views) at the second wave (February 2, 2026), about 8.3 months after the first wave (May 20 to June 2, 2025). Robust standard errors in parentheses. Comment tone and profanity are the standardized LIWC-22 Tone and Swear measures from the first-wave comment corpus; higher tone is more positive. Channel age (years) and ln(video count) are the production controls; ln(total comments) is the natural log of the API comment counts on each channel's three newest and three oldest videos, all measured at the first wave. Conditioning on first-wave views makes the coefficients associations with views accumulated between the waves. Temporal ordering only; the design does not remove unmeasured confounding and supports no causal claim. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table A06: OLS Regressions Predicting Third-Wave Logged Channel Views (June 2026) from First-Wave Comment Language, Conditional on First-Wave Views

	(Model 1)	(Model 2)	(Model 3)	(Model 4)
Woman creator	-0.024 (0.014)	0.002 (0.015)	0.006 (0.015)	0.017 (0.017)
First-wave comment tone		-0.043*** (0.008)	-0.039*** (0.009)	-0.038*** (0.010)
First-wave comment profanity			0.013 (0.009)	0.009 (0.009)
First-wave ln(views)	0.950*** (0.006)	0.943*** (0.006)	0.942*** (0.006)	0.940*** (0.007)
Channel age (years)	-0.010*** (0.002)	-0.009*** (0.002)	-0.009*** (0.002)	-0.010*** (0.002)
ln(Video count)	0.007 (0.010)	0.007 (0.010)	0.008 (0.010)	0.012 (0.011)
ln(Total comments)	0.023*** (0.007)	0.025*** (0.007)	0.025*** (0.007)	0.026*** (0.007)
Observations	8,993	8,993	8,993	8,993
R^2	0.917	0.918	0.918	0.918
Content-category fixed effects (40)	No	No	No	Yes

Notes: Channels of the 9,096-channel analytic sample with a third-wave view count. Dependent variable: ln(cumulative channel views) at the third wave (June 18, 2026), twelve to thirteen months after the first wave. Robust standard errors in parentheses. Comment tone and profanity are the standardized LIWC-22 Tone and Swear measures from the first-wave comment corpus; higher tone is more positive. Channel age (years) and ln(video count) are the production controls; ln(total comments) is the natural log of the API comment counts on each channel's three newest and three oldest videos, all measured at the first wave. Temporal ordering only; the design does not remove unmeasured confounding and supports no causal claim. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table A07: Retention to the Second Wave by Creator Gender: Linear Probability and Probit Models

	(Model 1)	(Model 2)	(Model 3)	(Model 4)
Woman creator	-0.0019 (0.0014)	-0.0013 (0.0014)	-0.1422 (0.1256)	0.0006 (0.0015)
Channel age (years)		0.0007*** (0.0002)	0.0795*** (0.0174)	0.0008*** (0.0002)
ln(Video count)		0.0004 (0.0006)	0.0163 (0.0565)	0.0011 (0.0008)
ln(Total comments)		0.0003 (0.0003)	0.0253 (0.0316)	0.0005 (0.0004)
Observations	9,096	9,096	9,096	9,096
Estimator	LPM	LPM	Probit	LPM
Outcome	In wave 2	In wave 2	In wave 2	Wave-2 views present

Notes: The 9,096-channel analytic sample. The outcome in Models 1 to 3 is whether the channel appears in the second-wave retrieval (February 2, 2026; 99.6 percent of channels); in Model 4 whether a second-wave view count is present, the stricter condition the lagged models require (99.5 percent). Robust standard errors in parentheses; Model 3 reports probit coefficients (creator gender $p = 0.258$). Channel age (years) and ln(video count) are the production controls; ln(total comments) is the natural log of the API comment counts on each channel's three newest and three oldest videos, all measured at the first wave. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table A08: Robustness of the Full Model Predicting Logged Channel Views from Creator Gender, Comment Tone, and Comment Profanity

	(Model 1)	(Model 2)	(Model 3)	(Model 4)	(Model 5)	(Model 6)	(Model 7)	(Model 8)
Woman creator	0.019 (0.034)	0.019 (0.063)	-0.000 (0.033)	0.017 (0.034)	0.032 (0.034)	-0.033 (0.034)	0.019 (0.034)	0.029 (0.034)
Comment tone	-0.379*** (0.018)	-0.379*** (0.059)	-0.360*** (0.017)	-0.380*** (0.018)	-0.233*** (0.035)	-0.363*** (0.018)	-0.380*** (0.018)	-0.345*** (0.018)
Comment profanity	0.077*** (0.018)	0.077** (0.025)	0.077*** (0.018)	0.076*** (0.018)	0.077*** (0.018)	0.048** (0.017)	0.077*** (0.018)	0.082*** (0.018)
Tone $\times$ ln(total comments)					-0.037*** (0.008)			
ln(mean words per comment)						-0.487*** (0.031)		
HD video quality							0.431*** (0.087)	
Observations	9,096	9,096	8,915	9,017	9,096	9,096	9,096	9,096
R^2	0.541	0.541	0.511	0.540	0.542	0.557	0.542	0.452

Notes: Randomly drawn channels from the Infludata frame of U.S.-based individual creators with at least 2,000 subscribers; the analytic sample is the 9,096 channels with views, channel age, video count, and scored comment language (3,137 women, 5,959 men). Creator gender was hand-coded and verified by a second coder ($\kappa = 0.96$); every other measure comes from the YouTube Data API or from text collected through it. Dependent variable: ln(cumulative channel views) in Models 1 to 7 and ln(views per video), cumulative views divided by video count, in Model 8. Every model contains channel age (years) and ln(total comments); Models 1 to 7 also contain ln(video count), which Model 8 omits because the outcome is already per video. Model 1 adds comment profanity and comment tone to the volume-adjusted model. Model 2 clusters standard errors on the 40 content categories; Model 3 trims logged views at the 1st and 99th percentiles; Model 4 keeps channels at least twelve months old; Model 5 adds the tone-by-comment-volume interaction; Model 6 adds ln(mean words per comment), which is a suppressor (comments on women's channels average 19.1 words versus 20.7 on men's; without it the volume-adjusted gap is -0.263 and with it -0.296); Model 7 adds HD video quality, which 97.0 percent of men's and 97.2 percent of women's channels have and which moves no estimate; in Model 8 the views-per-video gap before the mediators enter is -0.242. Robust standard errors in parentheses except Model 2. Comment tone and profanity are the standardized LIWC-22 Tone and Swear measures from the audience comment corpus; higher tone is more positive. Estimates are associations from cross-sectional channel-level data. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table A09: The Profanity Path Under Five Transformations of the Comment Profanity Measure: Bootstrapped Indirect Associations and the Remaining Creator-Gender Coefficient

Profanity measure	Woman → profanity (<i>a</i>)	Profanity → ln(views) (<i>b</i>)	Indirect (<i>a</i> × <i>b</i>)	95% interval	Woman creator in the view equation
Raw	-0.429	0.077	-0.0332	[-0.0484, -0.0186]	0.019
ln(1 + <i>x</i>)	-0.524	0.102	-0.0535	[-0.0709, -0.0340]	0.031
Inverse hyperbolic sine	-0.511	0.099	-0.0506	[-0.0678, -0.0314]	0.029
Winsorized at the 99th percentile	-0.487	0.093	-0.0452	[-0.0617, -0.0284]	0.026
Any profanity, binary	-0.165	0.067	-0.0111	[-0.0176, -0.0053]	0.011

Notes: Randomly drawn channels from the Infludata frame of U.S.-based individual creators with at least 2,000 subscribers; the analytic sample is the 9,096 channels with views, channel age, video count, and scored comment language (3,137 women, 5,959 men). N = 9,096 channels with complete comment-language measures. The raw profanity measure (LIWC-22 Swear, percent of comment words) has skewness 4.33 and kurtosis 37.5; each variant is standardized before entry. Each row re-estimates the mediation equations with that variant in place of raw profanity, comment tone held in the view equation; *a* is the creator-gender coefficient in the profanity equation, *b* the profanity coefficient in the view equation, and the indirect quantity their product, with a 500-replication nonparametric bootstrap percentile interval (seed 12345). Every equation contains channel age (years), ln(video count), and ln(total comments), the natural log of the API comment counts on each channel's three newest and three oldest videos. Estimates are associations from cross-sectional channel-level data. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$ on the creator-gender coefficient.

Table A10: Sensitivity of the Creator-Gender Gap and the Comment-Language Associations to Unmeasured Confounding (Cinelli and Hazlett Robustness Values)

	Estimate	SE	Partial R^2	$RV_{q=1}$	$RV_{q=1, \alpha=0.05}$	N
Woman creator, volume-adjusted gap model	-0.263	(0.033)	0.007	0.080	0.061	9,096
Comment tone, full model	-0.379	(0.018)	0.047	0.198	0.182	9,096
Comment profanity, full model	0.077	(0.018)	0.002	0.045	0.025	9,096

Notes: Randomly drawn channels from the Infludata frame of U.S.-based individual creators with at least 2,000 subscribers; the analytic sample is the 9,096 channels with views, channel age, video count, and scored comment language (3,137 women, 5,959 men). Creator gender was hand-coded and verified by a second coder ($\kappa = 0.96$). Dependent variable: $\ln(\text{cumulative channel views})$; the 9,096 channels with complete comment-language measures. Robustness values follow Cinelli and Hazlett (2020): $RV_{q=1}$ is the share of residual variance an unmeasured confounder would need to explain in both the treatment and the outcome to reduce the estimate to zero; $RV_{q=1, \alpha=0.05}$ the share needed to make it not distinguishable from zero at the 5 percent level. Benchmarks from the observed controls (partial R^2 with the outcome, with creator gender): $\ln(\text{video count})$ 0.293, 0.006; channel age 0.012, 0.006; $\ln(\text{total comments})$ 0.378, 0.000. Channel age (years) and $\ln(\text{video count})$ are the production controls. $\ln(\text{total comments})$ is the natural log of the API comment counts on each channel's three newest and three oldest videos; it holds constant how much reception a channel receives so that comment tone and profanity capture the type of reception at a given amount. Estimates are associations from cross-sectional channel-level data.

Table A11: OLS Regressions Predicting Logged Channel Views from Creator Gender and Comment Tone, Adding One LJWC-22 Comment-Language Component at a Time and Then All Together

	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)	(10)	(11)
Woman creator	0.001 (0.034)	0.019 (0.034)	-0.005 (0.034)	-0.035 (0.034)	0.001 (0.034)	0.019 (0.034)	0.001 (0.034)	-0.001 (0.034)	0.001 (0.034)	0.002 (0.034)	0.014 (0.035)
Comment tone	-0.403*** (0.017)	-0.379*** (0.018)	-0.417*** (0.018)	-0.509*** (0.022)	-0.528*** (0.024)	-0.372*** (0.018)	-0.390*** (0.019)	-0.390*** (0.019)	-0.409*** (0.018)	-0.415*** (0.018)	-0.543*** (0.030)
Comment profanity		0.077*** (0.018)									0.072*** (0.019)
Moral			-0.061*** (0.014)								-0.040** (0.014)
Positive emotion				0.180*** (0.023)							0.118*** (0.028)
Affect					0.176*** (0.025)						0.116*** (0.032)
Social						-0.089*** (0.016)					-0.100*** (0.023)
Polite							-0.029 (0.016)				0.003 (0.033)
Prosocial								-0.032 (0.018)			0.049 (0.030)
Anger									-0.027 (0.019)		-0.022 (0.018)
Conflict										-0.045** (0.016)	-0.034 (0.017)
Observations	9,096	9,096	9,096	9,096	9,096	9,096	9,096	9,096	9,096	9,096	9,096
R ²	0.539	0.541	0.540	0.544	0.543	0.541	0.540	0.540	0.540	0.540	0.548

Notes: Randomly drawn channels from the Infludata frame of U.S.-based individual creators with at least 2,000 subscribers; the analytic sample is the 9,096 channels with views, channel age, video count, and scored comment language (3,137 women, 5,959 men). Creator gender was hand-coded and verified by a second coder ($\kappa = 0.96$). Dependent variable: $\ln(\text{cumulative channel views})$; the 9,096 channels with complete comment-language measures. Robust standard errors in parentheses. Every comment-language measure is a standardized LJWC-22 channel mean over the audience comment corpus; higher tone is more positive. Every model contains channel age (years), $\ln(\text{video count})$, and $\ln(\text{total comments})$. Estimates are associations from cross-sectional channel-level data. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table A12: Comment Tone and Comment Profanity Received by Men's and Women's Channels, by Commenter Gender: Two-Sample *t*-Tests

	Men's channels	Women's channels	Difference	<i>p</i> -value	N
Comment tone, commenters coded male	52.349	62.363	10.014***	0.0000	9,069
Comment tone, commenters coded female	53.714	63.536	9.823***	0.0000	9,042
Comment profanity, commenters coded male	0.440	0.223	-0.217***	0.0000	9,072
Comment profanity, commenters coded female	0.389	0.171	-0.218***	0.0000	9,048

Notes: Two-sample *t*-tests on all channels with the measure; Difference = women's channels – men's channels. Commenter gender was inferred from commenter display names against Social Security Administration name frequencies under the liberal coding rule; N varies by row because a channel needs at least one assigned commenter of the relevant gender. Comment tone and comment profanity are the LIWC-22 Tone and Swear measures, unstandardized; tone uses LIWC-22's own coding (a comment with no tone words scores 20.23; comments with no words are excluded). * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table A13: OLS Regressions Predicting Logged Channel Views from Creator Gender and the Comment Language of Male-Coded and Female-Coded Commenters

	(Model 1)	(Model 2)	(Model 3)	(Model 4)	(Model 5)
Woman creator	-0.263*** (0.033)	0.019 (0.034)	-0.017 (0.034)	-0.043 (0.034)	0.007 (0.034)
Comment tone		-0.379*** (0.018)			
Comment profanity		0.077*** (0.018)			
Comment tone, commenters coded male			-0.359*** (0.019)		-0.253*** (0.029)
Comment profanity, commenters coded male			0.063*** (0.019)		0.048** (0.018)
Comment tone, commenters coded female				-0.349*** (0.017)	-0.149*** (0.027)
Comment profanity, commenters coded female				0.051** (0.019)	0.027 (0.017)
Channel age (years)	0.034*** (0.003)	0.046*** (0.003)	0.045*** (0.003)	0.045*** (0.003)	0.047*** (0.003)
ln(Video count)	0.882*** (0.014)	0.830*** (0.014)	0.833*** (0.014)	0.839*** (0.014)	0.832*** (0.014)
ln(Total comments)	0.621*** (0.008)	0.596*** (0.008)	0.599*** (0.008)	0.601*** (0.008)	0.598*** (0.008)
Observations	9,096	9,096	9,066	9,039	9,014
R^2	0.509	0.541	0.535	0.534	0.539

Notes: Randomly drawn channels from the Infludata frame of U.S.-based individual creators with at least 2,000 subscribers; the analytic sample is the 9,096 channels with views, channel age, video count, and scored comment language (3,137 women, 5,959 men). Creator gender was hand-coded and verified by a second coder ($\kappa = 0.96$). Dependent variable: ln(cumulative channel views). Robust standard errors in parentheses. Commenter gender is inferred from display names against Social Security Administration name frequencies (liberal coding, which assigns about 84 percent of the scored comments); tone and profanity are standardized LIWC-22 channel means over the comments of each commenter group, with tone on LIWC-22's own coding (comments with no tone words at 20.23; comments with no words excluded). N varies because a channel needs at least one comment from the relevant commenter group. Channel age (years) and ln(video count) are the production controls. ln(total comments) is the natural log of the API comment counts on each channel's three newest and three oldest videos; it holds constant how much reception a channel receives so that comment tone and profanity capture the type of reception at a given amount. Estimates are associations from cross-sectional channel-level data. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table A14: OLS Regressions Predicting Logged Channel Views from Creator Gender, Holding the Gender Composition of Commenters Constant

	(Model 1)	(Model 2)
Woman creator	-0.191*** (0.041)	-0.020 (0.041)
Comment tone		-0.384*** (0.018)
Comment profanity		0.082*** (0.018)
Share of commenters coded female	-0.003** (0.001)	0.002 (0.001)
Channel age (years)	0.034*** (0.003)	0.046*** (0.003)
ln(Video count)	0.883*** (0.014)	0.828*** (0.014)
ln(Total comments)	0.622*** (0.008)	0.596*** (0.008)
Observations	9,091	9,091
R^2	0.509	0.540

Notes: Randomly drawn channels from the Infludata frame of U.S.-based individual creators with at least 2,000 subscribers; the analytic sample is the 9,096 channels with views, channel age, video count, and scored comment language (3,137 women, 5,959 men). Creator gender was hand-coded and verified by a second coder ($\kappa = 0.96$). Dependent variable: ln(cumulative channel views). Robust standard errors in parentheses. The share of commenters coded female is inferred from display names (liberal coding). Comment tone and profanity are the standardized LIWC-22 Tone and Swear measures from the audience comment corpus; higher tone is more positive. Channel age (years) and ln(video count) are the production controls. ln(total comments) is the natural log of the API comment counts on each channel's three newest and three oldest videos; it holds constant how much reception a channel receives so that comment tone and profanity capture the type of reception at a given amount. Estimates are associations from cross-sectional channel-level data. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table A15: OLS Regressions Predicting Logged Channel Views: Does the Reception Slope Differ by Creator Gender, and Is There a Woman-Specific Dynamic in Female-Typed Categories?

	(Model 1)	(Model 2)	(Model 3)	(Model 4)
Woman creator	0.042 (0.037)	-0.051 (0.038)	0.022 (0.039)	-0.110* (0.051)
Comment tone	-0.336*** (0.022)	-0.242*** (0.024)	-0.079 (0.192)	-0.240*** (0.026)
Comment profanity	0.076*** (0.019)	0.060** (0.018)	0.037 (0.019)	0.070*** (0.017)
Woman $\times$ tone	-0.104** (0.037)	0.021 (0.040)	0.036 (0.042)	0.030 (0.053)
Woman $\times$ profanity	0.027 (0.052)	0.047 (0.052)	0.053 (0.048)	
Female-typed category		0.296*** (0.033)		0.269*** (0.041)
Tone $\times$ female-typed		-0.319*** (0.035)		-0.314*** (0.044)
Woman $\times$ female-typed				0.094 (0.073)
Woman $\times$ tone $\times$ female-typed				-0.038 (0.074)
Channel age (years)	0.046*** (0.003)	0.045*** (0.003)	0.044*** (0.003)	0.045*** (0.003)
ln(Video count)	0.833*** (0.014)	0.838*** (0.014)	0.858*** (0.014)	0.837*** (0.014)
ln(Total comments)	0.598*** (0.008)	0.599*** (0.008)	0.601*** (0.008)	0.599*** (0.008)
Content-category intercepts (40)	No	No	Yes	No
Content-category $\times$ tone slopes (40)	No	No	Yes	No
Observations	9,096	9,096	9,096	9,096
R^2	0.541	0.550	0.569	0.550

Notes: Randomly drawn channels from the Infludata frame of U.S.-based individual creators with at least 2,000 subscribers; the 9,096 channels with complete comment-language measures. Dependent variable: ln(cumulative channel views). Robust standard errors in parentheses. Comment tone and profanity are the standardized LIWC-22 Tone and Swear measures from the audience comment corpus; higher tone is more positive. Channel age (years) and ln(video count) are the production controls; ln(total comments) is the natural log of the API comment counts on each channel's three newest and three oldest videos. A category is coded female-typed if it sits above the median of the female-topic content index of its video titles; 72 percent of women's channels and 34 percent of men's channels are in female-typed categories. Model 1 fits one tone slope for all channels; the woman-by-tone term ($p = 0.005$) reflects women's concentration in female-typed categories, where comment tone is associated with views more steeply for creators of both genders (Model 2). Once each of the 40 content categories has its own tone slope (Model 3), the woman-by-tone term is not distinguishable from zero ($p = 0.39$). Model 4 is the backlash check: the three-way interaction that would indicate a woman-specific dynamic in female-typed categories is not distinguishable from zero ($p = 0.61$). The within-creator version of the check is in Appendix Table A22, column 6. Estimates are associations from cross-sectional channel-level data. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table A16: OLS Regressions Predicting Logged Channel Views from Creator Gender, Adding Off-Platform Links in the Channel Description and the Non-Subscriber Channel Trailer

	(Model 1)	(Model 2)	(Model 3)	(Model 4)	(Model 5)
Woman creator	-0.263*** (0.033)	-0.265*** (0.033)	-0.281*** (0.037)	-0.252*** (0.033)	-0.222*** (0.050)
Links to another platform (0/1)		0.039 (0.041)	0.007 (0.051)		
Woman creator $\times$ links to another platform			0.084 (0.085)		
Non-subscriber channel trailer (0/1)				0.279*** (0.032)	0.297*** (0.040)
Woman creator $\times$ non-subscriber trailer					-0.052 (0.066)
Channel age (years)	0.034*** (0.003)	0.034*** (0.003)	0.034*** (0.003)	0.032*** (0.003)	0.032*** (0.003)
ln(Video count)	0.882*** (0.014)	0.881*** (0.014)	0.881*** (0.014)	0.865*** (0.014)	0.865*** (0.014)
ln(Total comments)	0.621*** (0.008)	0.621*** (0.008)	0.621*** (0.008)	0.618*** (0.008)	0.618*** (0.008)
Observations	9,096	9,096	9,096	9,096	9,096
R^2	0.509	0.509	0.509	0.513	0.513

Notes: Randomly drawn channels from the Infludata frame of U.S.-based individual creators with at least 2,000 subscribers; the analytic sample is the 9,096 channels with views, channel age, video count, and scored comment language (3,137 women, 5,959 men). Creator gender was hand-coded and verified by a second coder ($\kappa = 0.96$). Dependent variable: ln(cumulative channel views); the 9,096 channels with complete comment-language measures. Robust standard errors in parentheses. Model 1 is the volume-adjusted model. The links indicator equals one when the channel description names or links to at least one of fifteen other platforms (Instagram, TikTok, X, Facebook, Patreon, Twitch, Discord, link aggregators, Snapchat, Threads, Pinterest, Reddit, Telegram, Spotify, SoundCloud); 20.7 percent of women's and 16.4 percent of men's channels do. The trailer indicator equals one when the channel runs a channel trailer for non-subscribers; 55.7 percent of women's and 61.6 percent of men's channels do. Channel age (years) and ln(video count) are the production controls. ln(total comments) is the natural log of the API comment counts on each channel's three newest and three oldest videos; it holds constant how much reception a channel receives so that comment tone and profanity capture the type of reception at a given amount. Estimates are associations from cross-sectional channel-level data. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table A17: Supplementary Estimates Referenced in the Text

<i>Panel A. Sequential controls on the creator-gender coefficient, and posting volume</i>					
Specification (ln views on woman)	Woman	SE	Implied gap (%)	Δ from prior	N
No controls	-0.4134***	(0.0460)	-33.9	–	9,096
+ channel age (years)	-0.3543***	(0.0457)	-29.8	0.0591	9,096
+ ln(video count) (the production controls)	-0.2164***	(0.0424)	-19.5	0.1378	9,096
+ ln(total comments) (the volume-adjusted model)	-0.2633***	(0.0333)	-23.1	-0.0468	9,096
Posting volume	Men	Women	Difference	<i>p</i>	N
Videos per channel (mean)	1,495	834	-662	0.0539	9,096
Videos per year of tenure (mean)	175.1	122.3	-52.8	0.0605	9,096
ln(videos per year of tenure) (mean)	4.254	4.103	-0.151	< .0001	9,096
<i>Panel B. Descriptive gaps by creator gender</i>					
Measure	Men	Women	Difference	<i>p</i>	N
Engagement rate, newest uploads (%)	6.62	7.62	1.00	< .0001	9,085
Engagement rate, oldest uploads (%)	4.08	5.64	1.56	< .0001	8,873
Engagement rate, Shorts (%)	7.72	9.47	1.76	< .0001	8,710
Views per subscriber (median)	173	122	-52	< .0001	9,096
Channels with mean comment tone below the scale midpoint (%)	44.3	20.6	-23.7	< .0001	9,096

Notes: All panels use the 9,096-channel analytic sample; smaller Ns reflect measures not available for every channel. *Panel A:* OLS of ln(total cumulative channel views) on the creator-gender indicator with channel age (years), ln(video count), and ln(total comments) entered one at a time in the order listed, robust standard errors in parentheses; implied gap is $100(e^{\beta} - 1)$; Δ is the change in the coefficient from the row above; the third row is the production-adjusted model and the final row the volume-adjusted model. Posting-volume rows are unequal-variance two-sample *t*-tests; videos per year is the video count divided by channel age in years. *Panel B:* engagement rate is (likes + comments)/views $\times 100$, with unequal-variance *t*-tests; views per subscriber is total views divided by subscribers, reported as the median with a nonparametric equality-of-medians test; the tone-share row counts channels whose mean comment tone (LIWC-22's own coding; comments with no words excluded) is below the scale midpoint of 50; comments with no tone-dictionary words receive a default LIWC-22 Tone of 20.23, so a mean below the midpoint does not indicate net-negative affect. Across the 9,096 channels, swear words average 0.35 percent of comment words (median 0.16 percent; SD 0.54), and 1,273 channels (14.0 percent) have none. Difference = women – men throughout. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table A17: Supplementary estimates referenced in the text (continued)

<i>Panel C. Creator content elements: standardized gender differences</i>					
Content element	Men (mean)	Women (mean)	Cohen's <i>d</i>	<i>p</i>	N
Title tone (LIWC-22 Tone)	42.871	55.761	+0.40	< .0001	9,096
Title profanity (LIWC-22 Swear)	0.087	0.038	-0.12	< .0001	9,096
Transcript tone (LIWC-22 Tone)	47.265	54.524	+0.35	< .0001	8,082
Transcript profanity (LIWC-22 Swear)	0.087	0.052	-0.11	< .0001	8,082
Thumbnail: positive facial expression	0.256	0.371	+0.47	< .0001	8,830
Thumbnail: warm-colour share	0.435	0.485	+0.37	< .0001	8,830
Thumbnail: clickbait score	0.281	0.227	-0.23	< .0001	8,830
<i>Panel D. Distributional robustness of the controlled gap</i>					
Specification (channel and comment volume controls)	Woman	SE	<i>p</i>	Gap (%)	N
Winsorized ln views, 1st/99th percentiles (OLS)	-0.2642***	(0.0327)	< .0001	-23.2	9,096
Quantile regression, 25th percentile	-0.2725***	(0.0417)	< .0001	-23.9	9,096
Quantile regression, median	-0.2577***	(0.0415)	< .0001	-22.7	9,096
Quantile regression, 75th percentile	-0.2247***	(0.0481)	< .0001	-20.1	9,096
Gaming categories excluded (OLS)	-0.2674***	(0.0342)	< .0001	-23.5	8,471
<i>Panel E. Diagnostics on the controlled model and the category structure</i>					
Diagnostic	Value				N
Maximum Cook's distance	0.0088				9,096
Variance inflation factor, mean (max)	1.04 (1.07)				9,096
Category size, min / median / max	32 / 126 / 865; 40 categories				9,096
Multi-homing rate, men vs. women (%)	16.4 vs. 20.7, <i>p</i> < .0001				9,096
<i>Panel F. Subscribers by creator gender</i>					
Specification (ln subscribers on woman)	Woman	SE	<i>p</i>	Gap (%)	N
No controls	-0.1095**	(0.0390)	0.0050	-10.4	9,096
+ channel age (years) and ln(video count) (the production controls)	-0.0061	(0.0380)	0.8726	-0.6	9,096

Notes: *Panel C:* two-sample *t*-tests by creator gender with Cohen's *d* (women — men; positive means women higher), unstandardized channel-level measures; title and transcript measures are LIWC-22 scores, thumbnail features are machine-coded on the newest-cohort thumbnails, and N varies with the channels on which each content element is available (titles: all 9,096 channels, with titles that contain no tone-dictionary words scored at LIWC-22's value of 20.23, as in Appendix Table A02; transcripts: 8,082; thumbnails: 8,830). *Panel D:* the volume-adjusted model (creator gender, channel age in years, ln(video count), ln(total comments)); winsorized and gaming-excluded rows are OLS with robust standard errors, quantile rows use the default *qreg* standard errors; gaming categories are the 6 content categories whose label contains *gam*. *Panel E:* diagnostics computed on the controlled-model fit; category sizes are counted on the 9,096-channel analytic sample; the multi-homing rate is the share of channels whose description links to at least one other platform (Appendix Table A16). *Panel F:* OLS of ln(subscribers) on the creator-gender indicator, without controls and with channel age (years) and ln(video count), on the 9,096 channels with complete comment-language measures (the sample of Table 8), robust standard errors in parentheses; gap is $100(e^{\beta} - 1)$. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table A18: Missingness of the Comment-Language Measures by Creator Gender

Drawn sample	9,232
Channels with comment tone and profanity present	9,099
Share missing	1.44%
Share present, men's channels	0.986
Share present, women's channels	0.985
p -value, logit of creator gender on presence	0.70
p -value, adding channel age, $\ln(\text{video count})$, $\ln(\text{total comments})$	0.55

Notes: Random sample of 9,232 US-based individual creators with at least 2,000 subscribers, drawn from the Infludata sampling frame (3,186 women, 6,046 men). Creator gender was hand-coded and verified by a second coder ($\kappa = 0.96$). Comment tone and profanity are missing where a channel had no scorable comments on its sampled videos. The test is run on all 9,232 channels, 9,099 of which have both comment measures; presence is unrelated to creator gender with and without the controls, so the comment-language sample is not gender-selected.

Table A19: Pearson and Spearman Correlations of Channel Comment Tone and Comment Profanity (LIWC-22) with VADER Sentiment Scores

Measure pair	Pearson r	Spearman ρ	N
LIWC Tone vs. VADER compound (convergent, valence)	0.87	0.89	9,096
LIWC Swear vs. VADER negative score	0.54	0.69	9,096
LIWC Swear vs. VADER compound	-0.40	-0.49	9,096
LIWC Tone vs. LIWC Swear (discriminant)	-0.33	-0.42	9,096

Notes: Channel-level correlations on the 9,096 channels with complete comment-language measures (3,137 women, 5,959 men), computed on the audience comment corpus. Comment tone is LIWC-22 Tone averaged over each channel's scored comments; comments with no positive or negative tone words take LIWC-22's own value of 20.23, and comments with no words are excluded. Comment profanity is LIWC-22 Swear. VADER is a rule-based sentiment instrument built for social-media text and is independent of the LIWC-22 dictionaries. All reported correlations are significant at $p < 0.001$.

Table A20: OLS Regressions Predicting Logged Channel Views, Comment Tone, and Comment Profanity from Creator Gender Within Content Categories and Within 128 Subtopics

	ln(Views)				Comment tone	Comment profanity
	(Model 1)	(Model 2)	(Model 3)	(Model 4)	(Model 5)	(Model 6)
Woman creator	-0.263*** (0.033)	-0.226*** (0.038)	-0.178*** (0.040)	-0.178*** (0.046)	0.376*** (0.024)	-0.297*** (0.022)
Channel age (years)	0.034*** (0.003)	0.034*** (0.003)	0.033*** (0.003)	0.033*** (0.004)	0.020*** (0.002)	0.002 (0.002)
ln(Video count)	0.882*** (0.014)	0.883*** (0.015)	0.892*** (0.015)	0.892*** (0.021)	-0.073*** (0.009)	-0.003 (0.012)
ln(Total comments)	0.621*** (0.008)	0.614*** (0.008)	0.611*** (0.008)	0.611*** (0.009)	-0.040*** (0.005)	-0.007 (0.006)
Observations	9,096	9,096	9,096	9,096	9,096	9,096
R^2	0.509	0.530	0.557	0.557	0.349	0.181
Content-category fixed effects (40)	No	Yes	No	No	No	No
Subtopic fixed effects (128)	No	No	Yes	Yes	Yes	Yes
Standard errors	Robust	Robust	Robust	Clustered on subtopic	Robust	Robust

Notes: Randomly drawn channels from the Infludata frame of U.S.-based individual creators with at least 2,000 subscribers; every model is estimated on the 9,096 channels with complete comment-language measures. Subtopics are 128 clusters of sentence embeddings of each channel's video titles, formed within the 40 content categories (do-file 116). Channel age (years) and ln(video count) are the production controls; ln(total comments) is the natural log of the API comment counts on each channel's three newest and three oldest videos. Models 5 and 6 show that women's channels still receive more positive and less profane comments within the same subtopic (standardized LIWC-22 Tone and Swear; higher tone is more positive). Estimates are associations from cross-sectional channel-level data.

* $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table A21: Bootstrap Decomposition of the Volume-Adjusted Gender Gap in Logged Channel Views Through Comment Tone and Comment Profanity

	Estimate	95% percentile interval
<i>Panel A: With 40-category fixed effects</i>		
Indirect through comment tone	-0.1960	[-0.2223, -0.1704]
Indirect through comment profanity	-0.0145	[-0.0257, -0.0037]
Total indirect	-0.2105	[-0.2379, -0.1844]
Direct association of creator gender	-0.0160	[-0.0884, 0.0576]
Volume-adjusted gap (total)	-0.2265	[-0.3008, -0.1523]
Share of the volume-adjusted gap accounted for	92.9%	[70.0%, 136.9%]
<i>Panel B: No fixed effects</i>		
Indirect through comment tone	-0.2490	[-0.2781, -0.2221]
Indirect through comment profanity	-0.0332	[-0.0486, -0.0188]
Total indirect	-0.2822	[-0.3136, -0.2544]
Direct association of creator gender	0.0190	[-0.0483, 0.0860]
Volume-adjusted gap (total)	-0.2633	[-0.3289, -0.1977]
Share of the volume-adjusted gap accounted for	107.2%	[85.1%, 143.1%]

Notes: Randomly drawn channels from the Infludata frame of U.S.-based individual creators with at least 2,000 subscribers; N = 9,096 channels with complete comment-language measures in every equation. Nonparametric bootstrap, 5,000 replications, seed 20260610, percentile intervals in brackets. Every equation contains channel age (years), ln(video count), and ln(total comments), so the total is the volume-adjusted gap, and the share is the share of that gap; Panel A additionally absorbs the 40 content-category indicators in every equation. ln(total comments) is the natural log of the API comment counts on each channel's three newest and three oldest videos; it holds constant how much reception a channel receives so that comment tone and profanity capture the type of reception at a given amount. The indirect quantities are the product of the creator-gender coefficient in the equation for each comment-language measure and that measure's coefficient in the view equation containing both measures. Where the interval on the share crosses 100 percent, it should be reported as an interval and not as a point claim that mediation is complete. Estimates are associations from cross-sectional channel-level data; the decomposition is of an association, not an identified causal mechanism.

Table A22: Robustness of the Within-Creator Association Between the Comments a Video Received and Its Logged Views

	Main (1)	Floor 30 (2)	No Shorts (3)	10+ videos (4)	Share negative (5)	Gender $\times$ reception (6)	Women (7)	Men (8)	Bad control (9)
Comment tone (LIWC-22)	-0.231*** (0.007)	-0.283*** (0.014)	-0.210*** (0.007)	-0.226*** (0.007)		-0.231*** (0.008)	-0.226*** (0.011)	-0.234*** (0.008)	-0.108*** (0.005)
Comment profanity (LIWC-22)	0.024*** (0.004)	0.037*** (0.011)	0.022*** (0.005)	0.024*** (0.004)		0.022*** (0.005)	0.030*** (0.008)	0.023*** (0.005)	0.010*** (0.003)
Share of negative comments (VADER)					0.109*** (0.005)				
Comment profanity (token list)					0.036*** (0.004)				
Woman $\times$ comment tone						-0.000 (0.014)			
Woman $\times$ comment profanity						0.010 (0.009)			
ln(comments on the video)									1.055*** (0.007)
Videos	131,859	44,002	115,217	120,006	135,148	131,859	48,945	82,914	131,859
Channels	7,153	5,982	7,002	4,906	7,157	7,153	2,510	4,643	7,153
Adjusted R^2	0.681	0.700	0.697	0.658	0.676	0.681	0.719	0.653	0.814

Notes: Unit of analysis: the video; dependent variable ln(video views); channel fixed effects in every column; standard errors clustered on channel in parentheses. The panel holds every video with at least ten comments on channels contributing at least two such videos, from the 2025 comment pulls and the 2026 per-video pull (135,148 videos on 7,157 channels); columns using LIWC-22 measures use the videos with at least ten LIWC-scored comments. Every column controls ln(video age in days), its square, and ln(duration). Column 1 is Table 7, Model 4. Column 2 keeps videos with at least 30 scored comments; column 3 drops Shorts; column 4 keeps channels with at least ten qualifying videos; column 5 replaces tone with the share of comments scored negative by VADER and profanity with the token-list rate; column 6 interacts creator gender with both reception measures (joint test of the two interactions $F = 0.52$, $p = 0.595$); columns 7 and 8 estimate the model separately for women's and men's channels; column 9 adds the video's own comment count, a post-treatment control shown only to document that it absorbs the association. Higher tone is more positive. Estimates are associations. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table A23: Pearson Correlations Among Logged Channel Views, Creator Gender, Comment Language, and the Controls

	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
1. ln(Views)								
2. Woman creator	-0.09							
3. Comment tone	-0.29	0.31						
4. Comment profanity	0.12	-0.21	-0.33					
5. Channel age (years)	0.16	-0.09	0.10	0.02				
6. ln(Video count)	0.41	-0.10	-0.14	0.04	0.19			
7. HD video	0.04	0.01	0.01	0.00	-0.00	-0.06		
8. Non-subscriber trailer	0.16	-0.06	-0.06	0.02	0.11	0.16	-0.00	
9. ln(Total comments)	0.50	0.03	-0.09	0.01	-0.03	-0.16	0.06	0.02

Notes: Randomly drawn channels from the Infludata frame of U.S.-based individual creators with at least 2,000 subscribers; every correlation is computed on the 9,096 channels of the analytic sample (3,137 women, 5,959 men), so every cell rests on the same channels. Table 1 reports summary statistics for these measures. Comment tone is LIWC-22 Tone averaged over each channel's scored comments (comments with no words excluded); comment profanity is LIWC-22 Swear. Channel age is measured in years, as it enters the models. Channel age (years) and ln(video count) are the production controls; ln(total comments) is the natural log of the API comment counts on each channel's three newest and three oldest videos and holds constant how much reception a channel receives. HD video indicates that the channel's sampled videos are high definition; non-subscriber trailer indicates that the channel runs a trailer for visitors who do not subscribe; both enter the Gelbach decomposition (Table 6) and are not in the main controls. 27 of the 36 correlations are significant at $p < 0.05$ (pairwise t -tests; see the log of record). Estimates are associations from cross-sectional channel-level data.

Table A24: Wild Cluster Bootstrap Inference on the Creator-Gender Gap in Logged Channel Views

	(Model 1)	(Model 2)	(Model 3)	(Model 4)
Woman creator	-0.216	-0.226	-0.263	-0.226
Clustered standard error	(0.097)	(0.089)	(0.084)	(0.077)
Analytic clustered p	0.031	0.015	0.003	0.005
Wild cluster bootstrap p	0.048	0.029	0.006	0.010
Bootstrap 95% confidence interval	[-0.436, -0.001]	[-0.421, -0.024]	[-0.449, -0.063]	[-0.398, -0.055]
Channel age (years), ln(video count)	Yes	Yes	Yes	Yes
ln(total comments)	No	No	Yes	Yes
Content-category fixed effects (40)	No	Yes	No	Yes
Observations	9,096	9,096	9,096	9,096

Notes: Randomly drawn channels from the Infludata frame of U.S.-based individual creators with at least 2,000 subscribers; the 9,096 channels with complete comment-language measures, as in Table 5. Dependent variable: ln(cumulative channel views). Models 1 and 2 are the production-adjusted gap (Table 5, Model 2); Models 3 and 4 add ln(total comments), the natural log of the API comment counts on each channel's three newest and three oldest videos (the volume-adjusted gap). Standard errors are clustered on the 40 content categories. Because the number of clusters is small, the wild cluster bootstrap of Cameron, Gelbach and Miller is reported alongside the analytic clustered p -value, using Webb six-point weights, 9,999 replications, the null imposed, and seed 20260827; Webb weights are preferred to Rademacher at this cluster count because Rademacher admits only 2^{40} distinct draws. Estimates are associations from cross-sectional channel-level data.

Table A25: OLS Regressions Predicting Logged Channel Views from Creator Gender (the Ladder Without Creator Content) under Floors on Scored Comments per Channel, Comment-Count Weights, and Alternative Channel-Level Reception Measures

	N	Woman coefficient			Tone	Profanity	Share
		Production	+ Volume	+ Tone, profanity			
All channels	9,096	-0.216	-0.263	0.019	-0.379	0.077	1.07
Floor: ≥ 20 scored comments	8,742	-0.240	-0.270	0.032	-0.401	0.087	1.12
Floor: ≥ 50 scored comments	8,277	-0.266	-0.280	0.041	-0.408	0.105	1.15
Floor: ≥ 100 scored comments	7,687	-0.303	-0.302	0.038	-0.412	0.116	1.13
Floor: ≥ 200 scored comments	6,876	-0.366	-0.344	0.007	-0.408	0.135	1.02
Floor: ≥ 500 scored comments	5,374	-0.387	-0.354	0.021	-0.414	0.138	1.06
Floor: ≥ 20 scored comments on videos with ≥ 10 scored comments	7,212	-0.306	-0.305	0.031	-0.399	0.123	1.10
Weighted by scored comments (analytic weights)	9,096	-0.364	-0.330	0.022	-0.403	0.130	1.07
Measure: videos with ≥ 10 scored comments only	7,576	-0.272	-0.288	-0.018	-0.361	0.066	0.94
Measure: mean of video means	9,096	-0.216	-0.263	0.004	-0.377	0.074	1.02
Measure: one random half of comments	9,083	-0.215	-0.263	0.002	-0.369	0.078	1.01
Measure: first 50 comments only	9,096	-0.216	-0.263	-0.037	-0.350	0.057	0.86

Notes: Randomly drawn channels from the Infludata frame of U.S.-based individual creators with at least 2,000 subscribers; the 9,096 channels with complete comment-language measures, or the subset named at left. Dependent variable: $\ln(\text{cumulative channel views})$. Each row re-estimates the ladder without creator content (no category fixed effects) on the sample or measure named at left. Production = channel age (years) and $\ln(\text{video count})$, the production controls; + Volume adds $\ln(\text{total comments})$, the natural log of the API comment counts on each channel's three newest and three oldest videos; + Tone, profanity adds standardized comment tone and comment profanity (LIWC-22 Tone and Swear; comments with no words excluded). Tone and Profanity are the coefficients on those two measures in the final model. Share = one minus the final woman coefficient over the volume-adjusted woman coefficient. Floors count the LIWC-22-scored comments behind each channel's tone and profanity means (median 772 per channel); the video-depth floor requires at least 20 scored comments on videos that each carry at least ten scored comments. The weighted row uses the scored-comment counts as analytic weights. Alternative measures recompute the channel means from comments on videos with at least ten scored comments, as the mean of video means, from one random half of each channel's comments, and from the first fifty comments only. Robust standard errors throughout; every woman coefficient in the Production and + Volume columns is significant at $p < 0.001$, and no woman coefficient in the + Tone, profanity column is distinguishable from zero at $p < 0.05$. Estimates are associations from cross-sectional channel-level data.

Table A26: Comment Reception and Logged Views Holding the Like Rate Constant, Across and Within Creators

<i>Panel A: Across creators (channel level; dependent variable $\ln(\text{cumulative channel views})$)</i>				
	(Model 1)	(Model 2)	(Model 3)	(Model 4)
Woman creator	0.019 (0.034)	0.081* (0.033)	-0.016 (0.038)	0.026 (0.037)
Comment tone	-0.379*** (0.018)	-0.321*** (0.017)	-0.426*** (0.020)	-0.356*** (0.019)
Comment profanity	0.077*** (0.018)	0.064*** (0.017)	0.046** (0.017)	0.035* (0.017)
Like rate (likes per 100 views)		-0.410*** (0.015)		-0.382*** (0.016)
Channel controls and comment volume	Yes	Yes	Yes	Yes
Content-category fixed effects (40)	No	No	Yes	Yes
Observations	9,096	9,096	9,096	9,096
R^2	0.541	0.574	0.560	0.588
<i>Panel B: Within creators (video level; dependent variable $\ln(\text{video views})$)</i>				
	(Model 5)	(Model 6)	(Model 7)	(Model 8)
Comment tone (LIWC-22)	-0.231*** (0.007)	-0.123*** (0.006)		
Comment profanity (LIWC-22)	0.024*** (0.004)	0.015*** (0.004)		
Comment tone (VADER)			-0.230*** (0.007)	-0.119*** (0.006)
Comment profanity (token list)			0.030*** (0.004)	0.024*** (0.004)
Like rate (likes per 100 views)		-0.822*** (0.013)		-0.818*** (0.013)
Video age, age squared, log duration	Yes	Yes	Yes	Yes
Channel fixed effects	Yes	Yes	Yes	Yes
Videos	131,859	128,074	135,148	131,222
Channels	7,153	7,031	7,157	7,037
Adjusted R^2	0.681	0.745	0.680	0.744

Notes: Randomly drawn channels from the Infludata frame of U.S.-based individual creators with at least 2,000 subscribers; the analytic sample is the 9,096 channels with views, channel age, video count, and scored comment language (3,137 women, 5,959 men). Panel A: robust standard errors in parentheses; the like rate is likes per 100 views on each channel's three newest and three oldest videos, winsorized at the 99th percentile and standardized (women 4.54, men 3.64; correlation with comment tone 0.12). Channel controls are channel age (years) and $\ln(\text{video count})$; comment volume is $\ln(\text{total comments})$. Panel B: the video panel of Table 7 (135,148 videos on 7,157 channels; LIWC-22 columns use the 131,859 videos with at least ten LIWC-scored comments); standard errors clustered on channel; the like rate is the video's likes per 100 views, standardized; 3926 videos with hidden like counts drop from Models 6 and 8. Because the like rate divides likes by views, part of its within-creator coefficient is arithmetic; it enters to hold audience approval constant and is not a control in the main models. Comment measures are standardized.

Table A27: Gelbach Decomposition of the Change in the Creator-Gender Coefficient on Logged Channel Views: Further Starting Points and Samples

	(1) Raw gap	(2) Raw gap	(3) Production- adjusted gap	(4) Raw gap	(5) Production- adjusted gap
Content category	Not in the model	Fixed effects	Not in the model	Own block	Own block
Transcripts in creator content	No	No	No	Yes	Yes
<i>Contribution of each block to the change in the creator-gender coefficient (log points), robust SE, [share of the starting gap]</i>					
Channel age and video count	-0.231*** (0.022) [55.9%]	-0.208*** (0.026) [52.0%]	(in start)	-0.249*** (0.023) [55.1%]	(in start)
Content category (40)		(in start)		0.041* (0.021) [-9.2%]	0.038 (0.021) [-16.0%]
Creator content (titles, video descriptions)	0.037*** (0.009) [-8.9%]	0.016* (0.007) [-4.1%]	0.037*** (0.009) [-17.1%]	0.044*** (0.011) [-9.7%]	0.044*** (0.011) [-18.2%]
Comment volume	0.084** (0.026) [-20.2%]	0.037 (0.029) [-9.3%]	0.045 (0.025) [-20.8%]	0.086** (0.027) [-19.1%]	0.044 (0.027) [-18.4%]
Comment profanity	-0.035*** (0.008) [8.4%]	-0.015** (0.006) [3.8%]	-0.035*** (0.008) [16.0%]	-0.026** (0.009) [5.8%]	-0.026** (0.009) [10.9%]
Comment tone	-0.257*** (0.015) [62.1%]	-0.198*** (0.014) [49.4%]	-0.257*** (0.015) [118.9%]	-0.329*** (0.020) [72.6%]	-0.324*** (0.020) [135.5%]
HD video	0.001 (0.002) [-0.2%]	-0.001 (0.001) [0.3%]	0.000 (0.002) [0.0%]	0.001 (0.001) [-0.2%]	0.000 (0.001) [-0.1%]
Non-subscriber trailer	-0.015*** (0.003) [3.7%]	-0.008* (0.003) [2.0%]	-0.010** (0.003) [4.5%]	-0.017*** (0.004) [3.9%]	-0.012*** (0.003) [4.9%]
Starting creator-gender coefficient	-0.413	-0.400	-0.216	-0.453	-0.240
Percent fewer views on women's channels	33.9	33.0	19.5	36.4	21.3
Ending creator-gender coefficient (full model)	0.003 ($p = 0.93$)	-0.024 ($p = 0.53$)	0.003 ($p = 0.93$)	-0.004	-0.004
Total accounted for (%)	100.8	94.0	101.4	99.2	98.5
Left over (%)	-0.8	6.0	-1.4	0.8	1.5
Observations	9,096	9,096	9,096	8,082	8,082

Notes: Gelbach (2016) decomposition of the creator-gender gap in logged cumulative channel views on 9,096 channels (except columns 4 and 5 on 8,082 with transcripts). Each entry is the women-minus-men difference in the block (holding the starting model's variables fixed) multiplied by the block's coefficient in the full model, so the entries do not depend on the order in which blocks are added. Robust standard errors are in parentheses and the share of the starting coefficient is in brackets. Positive shares indicate that equalizing the block between women's and men's channels narrows the gap; negative shares indicate widening. Shares sum with the share left over to 100 percent and can exceed 100 percent or be negative if blocks offset each other. Columns 1 to 3 differ by starting gap and content-category inclusion: (1) raw gap without category indicators (full model includes all blocks except category, ending coefficient 0.003, $p = 0.93$); (2) raw gap with 40 content-category fixed effects (ending coefficient -0.024, $p = 0.53$); (3) production-adjusted gap without category indicators (ending coefficient 0.003, $p = 0.93$); (4) raw gap with content category as a block and transcripts added to creator content; (5) production-adjusted gap with content category block and transcripts (ending coefficient -0.004). Blocks and controls as in Table 6. Standard errors heteroskedasticity-robust. This is an accounting of coefficient changes, not causal effects. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$ (for each block's contribution).

Table A28: Sample Flow from the Drawn Sample to the Analytic Sample, by Creator Gender

Step (variables added)	Channels	Lost	Women lost (%)	Men lost (%)	<i>p</i>
Drawn sample	9,232		3,186 women, 6,046 men		
Views, video count, and comment volume from the YouTube Data API	9,230	2	0.00	0.03	0.30
Channel age	9,229	1	0.03	0.00	0.17
Comment language (any scorable comment)	9,097	132	1.51	1.39	0.65
Comment tone defined (at least one comment with words)	9,096	1	0.00	0.02	0.47
Creator content (titles, video descriptions)	9,096	0	0.00	0.00	n/a
Transcripts (robustness models only)	8,082	1,014	10.55	11.46	0.19
Drawn to analytic sample (cumulative)	9,096	136	1.54	1.44	0.71

Notes: Each row adds its variables to every row above it; counts channels present in the row above and missing here, and the percentages are of women's and of men's channels in the row above; *p* is a two-sample test of proportions. Views: the two channels without views returned no YouTube Data API channel data. Channel age: one channel's creation date is blank in the sampling frame. Comment language: no scorable comment was collected for 132 channels (50 of them have no comments on any of the 12 API-sampled videos). Comment tone defined: one channel's scored comments contain only emoji, for which LIWC-22 Tone is undefined. Creator content loses no channel: titles or video descriptions with no tone words take LIWC-22's value of 20.23, and empty descriptions take tone 20.23 and profanity 0, each with an indicator in the models. Transcripts are used only in robustness models; the 1,014 analytic-sample channels without them have only non-English caption tracks (372), no caption track (447), no sampled video still available (157), or no sampled videos in the transcript fetch (38). The analytic sample is the 9,096 channels in the Creator content row.

Table A29: How Much of the Gender Difference in Comment Tone and Comment Profanity Creators' Own Content Accounts For

	Comment tone		Comment profanity	
	(Model 1)	(Model 2)	(Model 3)	(Model 4)
<i>Panel A: Titles and video descriptions (all 9,096 channels)</i>				
Woman creator, without creator content	0.656*** (0.021)	0.461*** (0.023)	-0.429*** (0.019)	-0.317*** (0.021)
Woman creator, with creator content	0.504*** (0.021)	0.369*** (0.022)	-0.383*** (0.020)	-0.295*** (0.021)
Share accounted for by creator content (%)	23.2	19.9	10.7	6.7
Share remaining (%)	76.8	80.1	89.3	93.3
<i>Panel B: Titles, video descriptions, and transcripts (the 8,082 channels with transcripts)</i>				
Woman creator, without creator content	0.689*** (0.021)	0.479*** (0.022)	-0.458*** (0.020)	-0.335*** (0.021)
Woman creator, with creator content	0.498*** (0.020)	0.376*** (0.021)	-0.374*** (0.022)	-0.293*** (0.022)
Share accounted for by creator content (%)	27.6	21.5	18.4	12.8
Share remaining (%)	72.4	78.5	81.6	87.2
Content-category fixed effects (40)	No	Yes	No	Yes

Notes: Randomly drawn channels from the Infludata frame of U.S.-based individual creators with at least 2,000 subscribers; Panel A uses the 9,096 channels of the analytic sample, Panel B the 8,082 of them with transcripts. Dependent variables: standardized comment tone (LIWC-22 Tone; Models 1 and 2) and standardized comment profanity (LIWC-22 Swear; Models 3 and 4), so the Woman creator coefficient is the difference in standard deviations. Every model contains channel age (years), ln(video count), and ln(total comments). Creator content is the standardized LIWC-22 tone and profanity of the channel's sampled video titles and video descriptions, with indicators for text with no tone words and for empty descriptions (Panel A), plus the tone and profanity of the creator's transcripts (Panel B). The share accounted for is the drop in the Woman creator coefficient when creator content is added, as a percentage of the coefficient without it. Robust standard errors in parentheses. Estimates are associations from cross-sectional channel-level data. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.