

Approval Without Amplification: Gender and Entrepreneurial Visibility on Algorithmically Mediated Platforms

Katie Apker
Cornell University, ILR School

Draft – July 28, 2026

Abstract

Research on gender inequality in entrepreneurship often explains women’s disadvantage through unfavorable evaluation: women receive less positive assessments and, as a result, fewer resources. I examine whether this mechanism operates the same way on algorithmically mediated platforms, where audience reception becomes visibility through engagement signals. Using a random sample of 9,232 established U.S. YouTube creator entrepreneurs, I combine channel metrics, creator characteristics, audience comments, and visual, written, and spoken content measures to examine channel views, a key metric of visibility that shapes monetization and venture success. Women creators receive about 33% fewer views than men creators. Production volume differences account for roughly a third of this raw gap, but the remaining 22% gap is not explained by the content women produce. Nor do audiences appear to like women’s content less: women receive more engagement, show no comparable controlled subscriber gap, and receive more positive, less profane comments. The visibility disadvantage instead reflects what I term a “negative arousal bonus”: positivity is associated with fewer views, negativity and profanity with more views, and women’s softer comment sections, a pattern consistent with benevolent sexism, are less likely to attract the visibility-feeding engagement signals that algorithmic recommendation systems promote, and this comment-language channel accounts for most of the remaining gap. These findings show that favorable, lower-intensity gendered reception can coexist with disadvantage when engagement signals feed visibility.

INTRODUCTION

Gender inequality remains one of the most durable patterns in work. Organizational, structural, and social processes sort women and men into different roles, attach different expectations to their competence and authority, and translate those expectations into unequal access to resources, rewards, and advancement (Acker 1990; Ridgeway 2011; Heilman 2012). Against that backdrop, entrepreneurship has often been viewed as an alternative to these constraints, a form of work through which women may bypass organizational gatekeepers and exercise greater control over opportunity, autonomy, and income (Thebaud 2015; Yang, Kacperczyk, and Naldi 2024). Yet even within entrepreneurship, disadvantage persists and women entrepreneurs frequently secure fewer entrepreneurial resources and experience worse entrepreneurial outcomes than men (Jennings and Brush 2013).

Extant research examining women's disadvantage in entrepreneurship has often explained this pattern through evaluation processes involving key demand-side audiences. Across entrepreneurial settings, investors, evaluators, and other audience groups tend to judge women's ventures less favorably, and women in turn receive fewer resources and less favorable treatment from these key resource holders (Kanze et al. 2018; Brooks et al. 2014; Ewens and Townsend 2020; Lee and Huang 2018; Abraham 2020). In these accounts, the party that evaluates is usually also the party that allocates the critical resources that aid entrepreneurs in their success, so unfavorable evaluation is seen to hurt women entrepreneurs while favorable evaluation should help them.

However, a growing share of entrepreneurship now occurs on algorithmically mediated platforms, where the link between audience evaluation and venture outcomes is less direct (Cutolo and Kenney 2021; Nambisan 2017). On such platforms, algorithms are part of the evaluation system: they sit between audience reception and visibility, a critical monetizable resource central to platform entrepreneurship (Covington et al. 2016; Rahman 2021). Recommendation systems distribute visibility partly by reading audience behavior toward content, which means that platforms change the conditions under which audience evaluation becomes reward. There is further reason to

doubt that these dynamics always follow the expectations set out by existing theories of audience evaluation. Online diffusion often follows intense, negative reactions, the same behavioral signals that engagement-optimizing systems reward, so negative audience evaluation can be associated with virality and visibility rather than only with rejection in such settings (Berger and Milkman 2012; Brady et al. 2017; Milli et al. 2025). We know less, however, about what these changed dynamics mean for gender and audience evaluation, where prior research leads us to expect favorable evaluation to help women and unfavorable evaluation to hurt them. Does the evaluation-driven mechanism behind women’s entrepreneurial disadvantage operate the same way when audience reception passes through algorithmic amplification before it becomes visibility? Which creators gain visibility and which do not, and how is that gendered?

I study these questions on YouTube, a large entrepreneurial attention market in which individual creators compete for visibility, audience reception is publicly observable, and views are economically consequential because they shape monetization and audience scale. I analyze a random sample of 9,232 established U.S. creators, combining API channel metrics, hand-coded creator demographic characteristics, audience comments, and text and image measures of creator content. I estimate mediation and robustness specifications that compare women’s and men’s total channel views and characterize audience reception through comment positivity and profanity, the dimensions that virality research identifies as central to what spreads in attention markets. To separate how audiences respond to women from what women produce, I hold constant the content itself in visual, written, and spoken form, as well as production features, channel characteristics, and topics. Finally, to address reverse causality, the concern that greater visibility might drive reception rather than the reverse, I use repeated channel-stat waves to test whether earlier reception is associated with later views conditional on baseline views.

My findings reveal a visibility disadvantage for women alongside favorable audience reception, and they point to a mechanism different from the unfavorable-evaluation account that prior research would suggest. Women creators receive about 33% fewer views on average than men creators, and roughly a third of that raw gap reflects production differences: women’s channels are

slightly younger and have fewer videos on average. The remaining roughly 22% is not explained by the content women produce, including topics, title language, transcripts, and thumbnails. The gap is also not explained by women sorting into lower-view categories or audiences liking women's content less. Audiences engage more with women's channels, women show no comparable controlled subscriber gap, and their comment sections are more positive and less profane. Rather, my findings suggest that the visibility disadvantage women creators experience sits not in how audiences respond to women but in what happens to those responses afterward, as engagement signals become inputs into algorithmically mediated visibility. More positive comment tone is associated with fewer views and profanity with more views, a combined dynamic in which negative, arousing engagement is rewarded with visibility, what I term a "negative arousal bonus." Women's softer comment sections are less likely to attract this bonus, a pattern consistent with benevolent sexism: the everyday default toward women is warmer and gentler, an ostensibly positive orientation (Glick and Fiske 2001; Jampol and Zayas 2021). Once this comment language enters the model, the controlled gap is reduced to statistical non-significance, and mediation specifications suggest that comment language accounts for most of it. Additional robustness analyses rule out reverse ordering, as well as differences in self-promotion, commenter gender composition, and audience backlash against women, and the results are robust to outliers, missing data, comment length, and alternative codings.

The findings contribute to three important conversations. First, they extend gender and entrepreneurship research on demand-side evaluation by showing that women's disadvantage need not arise from unfavorable evaluation (Kanze et al. 2018; Brooks et al. 2014; Lee and Huang 2018; Abraham 2020). In algorithmically mediated attention markets, more positive reception can coexist with worse entrepreneurial outcomes when the allocation of visibility depends on signals other than positivity, and lowering barriers to entry may shift inequality from access to amplification rather than remove it. Second, they contribute to research on platform attention and algorithmic amplification by identifying recommender-mediated amplification of gender gaps, distinct from bias in human evaluators, and by measuring the gap at scale (Lambrecht and Tucker

2019; Huszar et al. 2022; Milli et al. 2025). Finally, my findings also extend research on benevolent and ambivalent sexism by showing that positive, lower-intensity treatment carries material costs outside interpersonal evaluation and organizational assignment (Glick and Fiske 2001; Jampol and Zayas 2021). Gendered reception and engagement-optimizing amplification jointly author this structural disadvantage, reproducing in a new technological setting the gendered allocation of resources and rewards that organizations have long enacted (Acker 1990).

THEORY

Evaluation and Allocation Are Usually Coupled

Research on gender and entrepreneurship has treated audiences, investors, and other market actors as consequential evaluators of entrepreneurial quality. This demand-side turn has been important because it locates women's disadvantage not only in the ventures women choose to start, the sectors they enter, or the growth aspirations they hold, but also in how resource holders interpret and respond to otherwise comparable entrepreneurs. Investors ask women entrepreneurs prevention-oriented questions and men entrepreneurs promotion-oriented questions, a pattern associated with differential funding outcomes (Kanze et al., 2018); evaluators prefer the same entrepreneurial pitch when it is delivered by a man rather than a woman (Brooks et al., 2014); and investors are associated with lower resource provision to women founders even in large-scale field settings (Ewens and Townsend, 2020). This literature has therefore established a powerful account of entrepreneurial gender inequality: women often secure fewer resources because those who evaluate entrepreneurial promise assess them less positively, more skeptically, or against standards that are more difficult for them to satisfy. In this account, unequal outcomes are tied to an evaluative deficit. The central analytic object is the gendered judgment of entrepreneurial quality, and the key mechanism is the movement from that judgment to resource provision.

The architecture of that account contains an important and generally reasonable assumption. In the settings that anchor much of the demand-side literature, the actor who evaluates is also the

actor who allocates. The investor who assesses the founder also supplies capital; the competition judge who rates the pitch also helps determine the award; the resource holder who interprets entrepreneurial promise also controls access to the resource. Because evaluation and allocation are located in the same party or in the same tightly organized evaluative sequence, unfavorable evaluation is expected to be associated with worse entrepreneurial outcomes, while favorable evaluation is generally assumed to help. Even work that complicates a uniformly negative account of gendered evaluation often retains this architecture. Women can benefit when communal or gender-homophilous audiences value their projects (Greenberg and Mollick, 2017), and women may be advanced at lower-stakes stages while facing disadvantage at higher-stakes stages (Botelho and Poskanzer, 2026), but these arguments still largely concern settings in which evaluation is translated into reward inside a recognizable allocative process. The theoretical issue I take up is what happens when entrepreneurial activity is organized in a setting where reception and visibility are no longer linked through a single evaluator-allocator, but are instead recursively associated through audience behavior, human resharing, and an engagement-optimizing recommender.

Algorithmically mediated platforms complicate this connection because they make visibility itself the scarce entrepreneurial resource. For creators, visibility is not merely a signal of status or recognition; it is the condition under which audiences accumulate, sponsorships become valuable, and monetization becomes possible. Digital entrepreneurship is therefore not only a setting in which formal entry barriers are lower, but also a setting in which subsequent opportunity depends on being surfaced repeatedly to audiences at scale (Cutolo and Kenney, 2021). Attention-based theory is useful here because it treats attention as scarce, structured, and distributed through institutional arrangements rather than as a simple reflection of underlying quality (Ocasio, 1997). Category-based research makes a related point in a different domain: actors can be penalized through the omission of attention even when they are not explicitly evaluated negatively (Zuckerman, 1999). In platform entrepreneurship, this omission of attention is especially consequential because attention is not only observed after the fact; it is continuously reallocated. Prior audience responses become inputs into subsequent visibility through both human resharing and an engagement-optimizing

recommender.

This setting therefore requires a more precise separation among three constructs that are often treated as if they moved together. Reception refers to how audiences respond to a creator in the immediate interactional environment, including the positivity, profanity, and intensity of comments. Approval and commitment refer to audience attachment after exposure, such as engagement and subscribing. Visibility refers to whether a creator is seen, surfaced, and viewed at scale. In traditional entrepreneurial evaluation settings, these dimensions often appear closely aligned because the same evaluator supplies the resource after making the evaluation. On algorithmically mediated platforms, however, more positive, less profane, softer treatment can coexist with weaker conversion into visibility if the signals associated with that treatment are less strongly associated with human resharing and engagement-optimizing recommendation. The central theoretical move, then, is not to argue that audiences are irrelevant, nor to replace human evaluation with a simple algorithmic account. It is to ask which audience signals are associated with visibility once attention is structurally allocated by a recursive system in which audiences respond, those responses become behavioral traces, and later audiences are reached through a combination of human and algorithmic amplification.

Competing Accounts of Audience Treatment

The first plausible account follows directly from role incongruity and the broader literature on gendered devaluation. Entrepreneurship is strongly associated with assertiveness, ambition, risk taking, and authority, attributes that are culturally coded as more masculine than feminine. When women occupy roles associated with these expectations, audiences may respond with skepticism, scrutiny, or hostility. The demand-side entrepreneurship literature gives this expectation considerable force: if women founders are asked more prevention-oriented questions (Kanze et al., 2018), if identical entrepreneurial pitches are preferred when delivered by men (Brooks et al., 2014), and if investors are associated with lower provision of resources to women (Ewens and Townsend, 2020), then platform audiences might reproduce the same evaluative structure in a more distributed

form. Online contexts may further intensify this pattern because public comment environments lower social friction, make incivility visible, and can reward performative disagreement. From this perspective, platform entrepreneurship should not relax the demand-side evaluation problem identified in the entrepreneurship literature; it should relocate it into a larger, more diffuse audience. Instead of a small number of investors or judges evaluating women less positively, many audience members may do so publicly and repeatedly.

This account implies a direct continuity between familiar entrepreneurial gatekeeping and audience reception on platforms. The evaluators have changed, but the expected direction of gendered treatment has not. If women are perceived as less congruent with entrepreneurial authority, or if audiences hold them to different standards of competence, credibility, or ambition, women creators should receive more negative and more intense responses. Profanity and high-arousal reactions would then be symptoms of unfavorable treatment, and lower visibility would be the platform analogue of the familiar demand-side mechanism. The disadvantage would originate at the point of reception: women would be evaluated less positively, and that unfavorable reception would be associated with worse visibility. Under this account, platform mediation changes the scale and publicity of evaluation, but not the underlying logic that gender inequality is produced through harsher treatment.

Harsher-treatment prediction. If gender disadvantage in platform entrepreneurship follows the unfavorable-evaluation logic that dominates demand-side research, women's channels should be associated with more negative, more profane, and higher-arousal audience treatment than men's channels. This prediction treats audience reception as an expanded form of the gatekeeping process already theorized in gender and entrepreneurship: the audience is larger and less centralized, but the underlying disadvantage remains harsher treatment.

A genuinely competing account comes from benevolent and ambivalent sexism. Ambivalent sexism theory argues that hostile and benevolent forms of sexism are complementary rather than opposing ideologies: hostile sexism defends men's structural power through antipathy toward women who violate gender roles, while benevolent sexism preserves gender hierarchy through

subjectively positive, protective, and paternalistic treatment (Glick and Fiske, 2001; Bareket and Fiske, 2023). The important point is that the positive surface of benevolent sexism does not make it egalitarian. It can affirm women while also positioning them as needing protection, less suited for challenge, or more appropriately received through politeness and restraint. In organizational settings, this positive treatment is not innocuous. Benevolent sexism is associated with withholding challenging developmental assignments from women (King et al., 2012), with gendered forms of inflated or less candid feedback (Jampol and Zayas, 2021), and with interpretations that make sexism appear less objectionable because it is experienced as flattering or well intentioned (Hopkins-Doyle et al., 2019). Entrepreneurial settings are not exempt from these dynamics; benevolent sexism is present in startup evaluation and can be associated with gendered consequences even when its surface form is positive (Nguyen et al., 2024).

This perspective suggests a different form of gendered audience treatment. Women need not be met with more hostility for gender inequality to be present. They may instead receive more positive, less profane, softer treatment, not because audiences are free of bias, but because gendered norms of protection, politeness, and paternalism shape the register in which women are publicly received. Such treatment is easily misread as the absence of disadvantage because its surface valence is positive. Yet benevolent sexism theory directs attention to what positive treatment can do: it can lower challenge, mute candor, reduce confrontational intensity, and preserve gendered expectations under the guise of favorable regard. In a public comment environment, this means that women may draw reactions that are less hostile and less profane, while also being less affectively charged, less conflictual, and less likely to generate the behavioral intensity that platform attention systems reward. The consequential gender difference would therefore lie not in audiences rejecting women, but in audiences responding to women through a lower-intensity evaluative register that is itself gendered.

This account also clarifies why reception should not be collapsed into approval. More positive, less profane, softer treatment is a form of audience response, not necessarily evidence that women are being more strongly endorsed or that they face no penalty. Benevolent sexism often

works precisely because it preserves inequality through reactions that appear courteous, supportive, or protective. The surface positivity of the response can obscure its material implications, especially when the outcome of interest depends not on whether the initial audience sounds positive, but on whether that audience generates signals that lead the creator to be surfaced to subsequent audiences. In this sense, benevolent treatment is not merely a more pleasant version of the same evaluative process. It is a different mechanism of gendered reception, one that may produce disadvantage through softness rather than rejection.

Benevolent-treatment prediction. If gendered audience reception follows the logic of benevolent sexism, women's channels should be associated with more positive, less profane, softer treatment than men's channels. This prediction does not define such reception as approval or as the absence of penalty. It treats more positive and lower-intensity reception as a gendered bias in its own right, one whose consequences depend on whether that form of reception is associated with subsequent visibility.

These two predictions are intentionally in tension. The harsher-treatment account is more consistent with the dominant demand-side view that women's entrepreneurial disadvantage is associated with unfavorable evaluation. The benevolent-treatment account is more consistent with the possibility that women face disadvantage even when audience reception appears more positive. The point is not to resolve the tension theoretically before the analysis, but to specify what each account would imply about the affective and behavioral signature of audience response. In both accounts, gendered reception matters; what differs is whether women's disadvantage is expected to be associated with more negative, higher-arousal treatment or with more positive, lower-arousal treatment. The next theoretical step is therefore to ask how these different forms of reception become associated with visibility once visibility is allocated through attention systems rather than through a single evaluator's resource decision.

Reception Must Be Converted into Visibility

The next step is to theorize why these two forms of reception should be differently associated with visibility. Research on virality and affective transmission shows that diffusion is not simply a function of positive evaluation. High-arousal content is more likely to be shared, while low-arousal positive content is less likely to spread (Berger and Milkman, 2012). Moral-emotional language is associated with greater diffusion in social networks (Brady et al., 2017), out-group animosity is associated with stronger engagement than positive affect (Rathje et al., 2021), and engagement-based ranking is associated with the amplification of divisive or emotionally charged content beyond what users report preferring (Milli et al., 2025). These findings imply a distinction between being treated positively and being amplified. Positive reception may signal that an audience finds a creator agreeable, credible, or pleasant, but the signals most strongly associated with spread are often intensity, conflict, arousal, and negativity. The affective signature that reads as more positive in a comment section is therefore not necessarily the behavioral signature that expands visibility.

Human resharing is one part of this translation. Audience members do not merely evaluate content; through liking, replying, sharing, embedding, recommending, and discussing, they participate in allocating attention to later audiences. The motivation to reshare is shaped not only by whether something is liked but by whether it captures attention, provokes emotion, supplies conversational currency, or invites response. Negative and high-arousal reactions are especially useful for these purposes because they create urgency and social information. They give audiences something to contest, defend, mock, warn against, or pass along as notable. By contrast, positive low-arousal reactions may be sincere but behaviorally quiet. They can mark satisfaction without generating the impulse to forward, debate, or return repeatedly. Thus, even before considering the recommender, audiences can convert different forms of reception into visibility at different rates. A creator who receives more positive, less profane, softer treatment may be treated better in the immediate interactional environment while receiving fewer of the intense audience behaviors that expand exposure through human resharing.

The engagement-optimizing recommender is the second part of the translation. I do not

observe the recommender’s decision rule directly, and I do not treat it as an actor with gendered intent. Rather, the recommender is the inferred allocative context in which observable reception is associated with subsequent visibility. Public technical accounts of YouTube’s recommender describe optimization around expected watch time and implicit behavioral signals (Covington et al., 2016). Research on algorithmic allocation further shows that gender inequality can be associated with optimization objectives that are not themselves gendered (Lambrecht and Tucker, 2019), and that ranking systems can amplify some kinds of content relative to chronological baselines (Huszár et al., 2022). The relevant theoretical point is therefore not that an algorithm dislikes women, nor that it directly observes gender. It is that engagement optimization creates a reward structure in which behavioral intensity can matter more than the positivity of reception. If high-arousal or negative reactions are more predictive of continued watching, replying, returning, or sharing, then a recommender organized around engagement can be associated with greater visibility for content that elicits those reactions, even when users do not explicitly prefer such content (Milli et al., 2025).

Dual amplification changes the meaning of audience reception. If human resharing and the engagement-optimizing recommender are both more strongly associated with high-arousal and negative reactions than with positive, low-arousal reactions, then more positive, less profane, softer treatment need not be associated with more views. I term this possibility a “negative arousal bonus”: the visibility premium associated with the behavioral intensity of negative or profane reception, relative to the weaker amplification associated with positivity alone. The term is useful because it identifies the conversion mechanism without implying that negativity is intrinsically better, that audiences consciously reward hostility, or that the recommender directly observes gender. It simply names the possibility that the signals most associated with being surfaced are not the same signals that read as positive treatment. In an attention market, the conversion problem is therefore not whether an audience response is positive in tone, but whether it generates the behavioral traces through which visibility is allocated.

This conversion problem is especially important because amplification is recursive. Visibility exposes a creator to additional audiences; those audiences respond; their responses become

part of the next round of behavioral signals; and human resharing and the engagement-optimizing recommender are again associated with which creators become visible to still later audiences. Reception is therefore not a one-time judgment that precedes a single allocation decision. It is repeatedly folded back into the conditions under which later reception occurs. A creator who receives high-arousal, conflictual, or profane responses may receive more behavioral intensity, which is associated with greater visibility, which in turn exposes the creator to more audiences capable of generating still more intensity. A creator who receives more positive, less profane, softer treatment may convert exposed viewers into engaged followers, yet receive fewer of the signals that expand exposure to new viewers. Small differences in the affective form of reception can therefore become associated with larger differences in cumulative visibility.

The recursive reception-amplification loop also clarifies why approval and commitment are not permanently insulated from visibility. In the short run, a creator can have equal or higher engagement among those who see the content and still receive fewer views if the visibility channel brings fewer new audiences. Over time, however, visibility shapes the size of the audience pool from which future engagement and subscribers can be drawn. If a creator is surfaced less often, the stock of potential future followers grows more slowly, even if conversion after exposure is strong. Thus the theory expects the clearest initial disadvantage to appear in views, because views are closest to the visibility allocation process. But it also leaves open the possibility that repeated visibility differences may later be associated with commitment outcomes as well. The key distinction is temporal and structural: softer reception can coexist with parity or strength in approval among exposed audiences, while still being associated with weaker expansion to audiences not yet reached.

The Visibility Paradox

These arguments yield a two-dimensional paradox. On the reception dimension, the benevolent-treatment account implies that women may be associated with more positive, less profane, softer treatment. On the allocation dimension, negativity and arousal research implies that

this very form of treatment is less strongly associated with visibility than more intense or negative reception. Attention-based theory clarifies why the paradox is organizationally consequential: reception matters only insofar as it is converted into the scarce resource of attention (Ocasio, 1997). When the conversion regime privileges intensity, the audience response that appears more positive may be less useful for visibility. The same response can therefore look advantageous at the level of immediate interaction and disadvantageous at the level of cumulative attention.

Visibility paradox proposition. If the benevolent-treatment account holds, and if positivity is amplified less than intensity through human resharing and the engagement-optimizing recommender, then women's channels should be associated with a disadvantage in visibility even as approval and commitment remain at parity or higher. Because reception and amplification are recursive, this visibility disadvantage may also become associated with downstream commitment differences over time.

The proposition is deliberately conditional. If the harsher-treatment prediction is supported, the platform setting would extend the familiar demand-side account: women would be associated with worse reception and worse visibility. If the benevolent-treatment prediction is supported, the platform setting would reveal a different mechanism: disadvantage associated not with rejection, but with the weaker visibility conversion of more positive, less profane, softer treatment. In that case, gendered inequality would not require audiences to evaluate women negatively, nor would it require an engagement-optimizing recommender to contain a gendered objective. It would arise from the association between gendered reception norms and an attention structure in which the signals attached to intensity are more strongly associated with being seen.

This formulation also specifies the boundary condition for the gender and entrepreneurship literature. When evaluation and allocation are coupled, more positive reception should generally improve outcomes because the evaluator who responds favorably is positioned to supply the resource. When algorithmically mediated platforms intermediate between reception and visibility, however, more positive reception must first be converted into attention through human resharing and the engagement-optimizing recommender. The relevant inequality is therefore not only whether

women are treated more or less favorably at the point of reception, but whether the form of treatment they receive is associated with subsequent visibility. Platform entrepreneurship thus reveals a route to gendered disadvantage in which favorable-sounding reception can be structurally costly. The paradox is that softer treatment may be associated with lower visibility precisely because the attention system rewards the forms of audience intensity that softer treatment lacks.

DATA AND METHODS

Research Design and Setting

I use a mediation design with temporal ordering to examine whether gendered audience language mediates the association between creator gender and subsequent viewership, and I assess the credibility of that interpretation with four complementary strategies. First, an eight-month two-wave design uses first-wave audience language to predict second-wave viewership conditional on first-wave viewership. Second, I report formal Cinelli-Hazlett sensitivity analyses for unmeasured confounding (Cinelli and Hazlett, 2020). Third, I estimate every focal specification with and without content-category fixed effects. Fourth, I use a supply-side falsification strategy that measures creator-controlled language and packaging across four public surfaces of the channel, allowing me to test whether what creators produce, rather than how audiences respond, drives the observed gap. Together, these strategies make the direction of the mediation claim temporally explicit, bound how much unmeasured confounding the estimates could withstand, and rule out category composition and measurable creator-side content choices as rival explanations.

YouTube is an appropriate setting for studying gendered evaluative inequality in digital work because it separates audience evaluation from outcome allocation in a way that is observable at scale. Creators operate in a platform architecture in which an algorithmic intermediary translates engagement signals into visibility, so the conversion of audience response into views is a central feature of the setting rather than a nuisance of measurement (Covington et al., 2016). Viewers typically encounter videos through search, subscription feeds, external links, or recommendation,

watch some portion of the content, and only then may scroll to the comment section. Passive viewing is common, and commenting is a publicly visible, post-viewing evaluative act rather than a complete record of all audience response (Khan, 2017; Nonnecke and Preece, 2000). This makes comments useful for my purpose: they capture the affective texture of the audience response that is visible to other users and potentially legible to platform systems, even though they do not represent the full viewing audience.

The setting also makes the focal comparison theoretically tractable. The creators in this study are individually identifiable, and creator gender is legible to audiences through profile images, channel names, biographical information, and content. This condition matters because gendered evaluative schemas require gender to be socially available (Ridgeway, 2009). At the same time, YouTube provides public, channel-level measures of cumulative views, subscribers, video counts, channel age, comment activity, titles, descriptions, transcripts when available, and thumbnails. These data allow me to observe the outcome, the audience language proposed as the mediating channel, and a broad set of creator-controlled content and packaging choices. The economic stakes are also meaningful. The sampling frame consists of established creators above the rough subscriber threshold at which YouTube monetization becomes accessible through the Partner Program, so viewership is not merely a status metric. It is tied to advertising revenue, sponsorship opportunities, and future discoverability. The setting therefore allows me to study a management question about digital entrepreneurship: how gendered audience treatment is converted, or not converted, into platform-mediated visibility.

Sample and Data Collection

I obtained the sampling frame from Infludata, a commercial influencer analytics platform, under an academic license. I specified five criteria: YouTube as the platform, English-language channels, United States location, male or female gender assignment in Infludata's database, and a recent post within 90 days of data collection. I used Infludata's gender field only as a screening device to locate likely individual creator accounts rather than faceless brand accounts, which

typically receive no gender classification. Infludata delivered 14,700 records. After deduplication and removal of records without a usable channel identifier, the file contained 14,169 unique channels.

I then applied two verification steps. First, trained research assistants reviewed channels to confirm that the account represented an individual creator rather than a team, organization, media property, or non-identifiable account. Second, research assistants independently coded creator gender using multiple profile sources, including profile images, channel names, creator content, biographical information, and linguistic cues. Interrater reliability on a 5% validation subsample was $\kappa = 0.96$.¹ Channels for which individual creator status or creator gender could not be assigned with confidence were excluded. These screens yielded an analytic sample of 9,232 U.S.-based YouTube creators. The sample is therefore a sample of established, identifiable individual creators indexed by Infludata, not a census of all YouTube accounts. Very new creators, creators below Infludata’s coverage threshold, non-individual accounts, and creators whose gender is not legible from public materials fall outside the scope of the analysis.

Channel metadata were collected through the YouTube Data API v3 under the YouTube Researcher Program. First-wave channel metrics were retrieved between May 20 and June 2, 2025; a second wave was collected on February 2, 2026, approximately eight months later, supporting the lagged-dependent-variable specifications described below. These data include cumulative channel views, subscriber counts, video counts, channel creation dates, video-quality indicators, and other public channel characteristics. The dependent variable is based on YouTube’s channel-level cumulative view count, which aggregates views across the creator’s public video portfolio. This portfolio-level measure is appropriate for the argument because the paper studies total platform visibility for entrepreneurial creators, although public API data do not permit a clean decomposition of lifetime channel views by format.

For each channel, I also collected creator-controlled text and audience comments. Video titles were sampled from up to 12 videos per channel: three oldest, three newest, three most popular, and up to three sponsored videos when such videos were identifiable. Channel descriptions

¹The $\kappa = 0.96$ reliability estimate applies only to human research-assistant coding of *creator* gender. It does not describe any algorithmic or imputed measure of commenter gender.

were collected alongside titles. Comments were sampled through the YouTube Data API by retrieving each channel's uploads playlist, randomly shuffling available videos, and collecting up to 200 comments per video until reaching 1,000 comments or exhausting available videos. Videos with disabled comments were skipped. This procedure produced a large channel-level corpus of public comments, which I aggregate to the creator level for the mediation analyses. The main mediation sample uses listwise deletion for missing comment-language measures. The full analytic sample is 9,232 channels, and the mediator-complete subsamples range from 9,097 to 9,100 channels depending on the specific language measure used. Missingness on the comment-language requirement is assessed in the diagnostics reported in Appendix Table A14.

Measures

The primary dependent variable is cumulative viewership, measured as the natural log of total channel views. I focus on views because viewership is the core metric through which YouTube creators obtain advertising revenue, visibility, and downstream commercial opportunities (Cunningham and Craig, 2021; Duffy, 2017). Views also matter because recommender systems learn from behavioral engagement signals, so accumulated viewing is both an outcome of prior visibility and an input into future visibility (Covington et al., 2016; Salganik et al., 2006). I use the logged transformation because cumulative views are highly right-skewed. In secondary analyses, I examine subscriber count and engagement rate as alternative outcomes. I do not include either as a headline control. Subscriber count is jointly determined with views and is itself an outcome of audience commitment. Engagement rate is also excluded from the control set because its numerator includes commenting behavior and its denominator includes views, placing it on the causal pathway rather than in the set of pre-treatment confounders.

The independent variable is creator gender, coded as an indicator equal to one for women creators and zero for men creators, based on the research-assistant coding described above. The empirical comparison is therefore between publicly identifiable women and men creators in the established-creator sampling frame. I do not interpret this measure as capturing gender identity in all

of its complexity. It captures the binary gender presentation legible to audiences and codable from public channel materials, which is the theoretically relevant dimension for the audience-treatment mechanism studied here.

The mediators are two dimensions of audience comment language: valence and arousal. I anchor these measures in the dimensional-affect framework, which distinguishes the positive or negative character of emotional language from its activation or intensity (Bradley and Lang, 1999; Warriner et al., 2013). This distinction is central to the argument because prior work on online diffusion shows that arousal can matter above and beyond valence (Berger and Milkman, 2012). I operationalize valence with the LIWC-22 Tone summary variable and arousal with the LIWC-22 Swear category (Tausczik and Pennebaker, 2010; Boyd et al., 2022). Tone captures the balance of positive and negative emotional language, with higher values indicating more positive language. Swear captures profanity. I interpret profanity as an arousal or intensity signal rather than as a simple negative-valence measure because taboo words are emotionally activating and can appear in both positive and negative contexts (Janschewitz, 2008; Stephens and Robertson, 2020). Both comment measures are computed at the comment level and then averaged to the channel level. Because LIWC variables are word-proportion measures, the comment measures are length-normalized by construction. The models also control for total comment volume, and robustness specifications add mean comment length.

LIWC is appropriate for this study for three reasons. First, it is a transparent, dictionary-based instrument with a long validation record in social science research (Tausczik and Pennebaker, 2010; Boyd et al., 2022). Second, it has been used extensively in management research, including studies of CEO letters, executive news coverage, earnings calls, investor communications, interviews, and crowdfunding narratives (Gamache et al., 2015; Love et al., 2017; Malhotra et al., 2018; Pan et al., 2018; Anglin et al., 2018). Third, using the same dictionary architecture across audience comments and creator-controlled surfaces allows me to compare audience response with creator content using parallel measures. LIWC summary variables and single-category dictionary measures are not Likert-scale indices, so internal-consistency statistics such as Cronbach's alpha

are not the relevant reliability criterion. Instead, I report construct-validity checks. In particular, I benchmark LIWC Tone against VADER compound sentiment (Hutto and Gilbert, 2014), a rule-based sentiment instrument designed for social media text, and use this cross-method comparison as a convergent-validity check in Appendix Table A15.

The main control variables are channel age, logged video count, an HD video-quality indicator, a non-subscriber-trailer indicator, and logged total comment volume. Channel age captures tenure and accumulated opportunity to acquire views. Video count captures production volume and catalog size. The HD indicator captures a basic production-quality threshold. The non-subscriber-trailer indicator captures whether the creator has configured a distinct channel trailer for non-subscribers, which I interpret as a public signal of channel strategy and audience conversion effort. Total comment volume accounts for the scale of the comment corpus from which the mediators are estimated. I report models with and without YouTube content-category fixed effects using the platform’s granular topic classifications. The fixed-effect specifications compare women and men within content categories, addressing horizontal sorting as a rival explanation.

Supply-Side Falsification Instrument

A central threat to the interpretation is that audience comment language may simply mirror creator-side content. If women produce different videos, titles, descriptions, speech, or thumbnails, and those creator-controlled choices independently drive views, then comment tone and profanity could be proxies for supply-side strategy rather than audience treatment. I address this concern with a four-surface supply-side falsification instrument that measures creator-controlled language and packaging across the public surfaces through which creators present their work.

The first surface is video titles. I apply LIWC-22 to sampled video titles and compute channel-level measures of title tone and title profanity. The second surface is video descriptions. I apply the same LIWC-22 measures to description text, again aggregated to the channel level. The third surface is spoken content. Where transcripts are available, I apply LIWC-22 to video transcripts to measure the tone and profanity of the creator’s spoken language. The fourth surface is

thumbnail packaging. For thumbnails, I measure visual packaging features such as face presence, facial expression, color features, and text-overlay characteristics. I also apply OCR to thumbnail overlay text and compute LIWC tone and profanity on that text. The thumbnail measures are treated as creator packaging choices, not as measures of creator gender. Apparent gender presentation in thumbnails is therefore not used as a control for the gender independent variable.

I use these four surfaces in two ways. First, I enter creator-side language and packaging measures as controls and rival mediating channels to test whether creator-controlled content explains the gender-viewership association that the comment mediators otherwise capture. Second, I use the same surfaces to test whether gender differences in audience comment language are reducible to gender differences in creator content. This is a falsification strategy rather than a claim that all supply-side quality is observed. Some important platform data, including impressions, click-through rates, watch time, retention, and viewer demographics, are owner-only YouTube Analytics measures and cannot be observed through the public Data API. The supply-side instrument therefore gives a broad public-surface test of content and packaging, while the unobservable platform measures remain scope conditions.

I also include two additional rule-out strategies. To address the possibility that comment language reflects who comments rather than how audiences treat creators, I examine commenter-gender composition and report specifications that add commenter gender as a rival channel in Appendix Table A8. To address the possibility that the audience-language result is a LIWC artifact, I replicate the comment-valence channel using VADER sentiment on the comment side while retaining LIWC on the creator-content side. This cross-method check is reported in Appendix Table A15.

Estimation Strategy

I estimate the main models with ordinary least squares and heteroskedasticity-robust standard errors (White, 1980). The unit of analysis is the creator channel. The baseline specifications estimate the association between creator gender and logged cumulative views, first without controls,

then with the five channel controls, and then with content-category fixed effects. The main progression appears in Table 3. I use the fixed-effect models to show how the estimate changes when comparisons are made within YouTube content categories rather than across the full category distribution. I also report category-clustered and wild-cluster-bootstrap specifications in the robustness appendix because the category structure is a natural source of dependence across channels.

I then estimate mediation models in which comment tone and comment profanity enter as audience-treatment mediators. The mediation analysis follows the standard decomposition logic of Baron and Kenny (1986), with bootstrap confidence intervals for indirect effects (Preacher and Hayes, 2008). I estimate the indirect paths for tone and profanity separately and jointly, and I report the bootstrap mediation results in Table 4. I interpret the mediation estimates as evidence about whether the association between creator gender and views is statistically carried by the proposed audience-treatment channel.

The temporal-ordering specifications use the two-wave structure of the data. First-wave comment tone and profanity predict second-wave logged views, conditional on first-wave logged views and the same control variables. This lagged-dependent-variable design asks whether audience language measured before the second outcome measurement predicts subsequent viewership after accounting for the viewership stock already accumulated at the first measurement. The design helps address reverse ordering, but it does not eliminate all unmeasured confounding. I report these models in Appendix Table A1.

I use three additional sets of models to address rival explanations. First, I assess horizontal sorting with category fixed effects and an Oaxaca-Blinder decomposition, reported in Appendix Table A16. Second, I test supply-side alternatives by adding creator-controlled title, description, transcript, thumbnail visual, and thumbnail overlay-text measures, both as control blocks and as rival channels in the mediation framework. These analyses are summarized in Table 5 and reported in full in the appendix supply-side tables. Third, I report Cinelli-Hazlett sensitivity analyses that quantify how strong an omitted confounder would have to be, in partial- R^2 terms, to explain away the focal estimates (Cinelli and Hazlett, 2020). These sensitivity analyses are reported in Appendix

Table A3. This structure follows the logic that the main threat should be named directly and then bounded empirically.

Robustness Checks

The robustness checks are organized around measurement, specification, and rival-mechanism concerns. For the dependent variable, I re-estimate the models using subscriber count and engagement rate as alternative outcomes, while keeping them out of the main control set for the causal-ordering reasons described above. For the mediator measurement, I test alternative transformations of the right-skewed profanity measure, including logarithmic, inverse-hyperbolic-sine, winsorized, and binary specifications. I also add mean comment length, because LIWC scores are proportions and short comments can be noisier. For construct validity, I use VADER sentiment as an independent measure of comment valence. For model dependence, I report specifications with and without category fixed effects, category-clustered standard errors, and wild-cluster-bootstrap inference.

I also probe sample and scope conditions. I repeat the main specifications on channels with longer posting histories, examine the consequences of including channels with ambiguous gender coding, and report missing-data diagnostics for the mediator-complete sample. I use commenter-gender composition to distinguish treatment norms from the gender composition of the commenting audience. I use the four-surface supply-side instrument to test whether the creator's own titles, descriptions, spoken transcripts, thumbnail visuals, or OCR'd thumbnail text account for either the view gap or the audience-language gap. The full set of rival explanations and robustness checks is summarized in Table 5, with detailed specifications reported in the appendix.

RESULTS

Descriptive Patterns and the Core Puzzle

Women creators receive substantially fewer views even though the observable audience response to their channels is more positive and lower-arousal. Table 1 reports descriptive statistics for the analytic sample of 9,232 channels, 34.5% of which are run by women. Viewership is highly dispersed: the standard deviation of log views is 2.1, so channels one standard deviation apart differ by roughly a factor of eight. The bivariate correlations preview the empirical pattern: the *woman* indicator correlates negatively with log views ($r = -0.09$), positively with comment tone ($r = 0.27$), and negatively with comment profanity, while comment tone correlates negatively with log views ($r = -0.28$). Thus, before any model-based adjustment, women receive fewer views, receive warmer comments, and warmer comment sections are associated with smaller audiences.

[Insert Table 1 about here]

Table 2 sharpens the puzzle by showing that the view gap is not accompanied by worse audience interaction. Comments on women's channels average 53.5 on the LIWC Tone scale, compared with 43.1 on men's channels. Profanity in those comments is less than half as common, 0.19 versus 0.43 percent of comment words. Engagement rate, computed from the YouTube API as likes plus comments divided by views, multiplied by 100, and averaged over the twelve sampled videos per channel, is also higher for women's channels, 3.89% versus 3.24%. Women's channels also accumulate more total comments. Yet viewership is 0.408 log points lower for women, approximately 33% fewer views, while the raw subscriber difference is only about 10%. The disadvantage is therefore concentrated in views rather than in observable engagement from viewers who arrive.

[Insert Table 2 about here]

This descriptive pattern is consistent with benevolently biased treatment, not with a simple account in which audiences respond less to women. Women's channels receive more positive, less

profane, and more engaged audience interaction, but those forms of treatment do not translate into comparable viewership. The remainder of the analyses therefore examine whether the gender view gap is associated with the affective signature of audience comments: lower negativity and lower arousal, which should travel less readily through both human resharing and the recommendation algorithm.

The Viewership Gap and the Comment-Language Channel

The raw gender view gap is large and statistically precise. Column 1 of Table 3 reports the unconditional association between creator gender and log views: the coefficient on *woman* is -0.4074 ($p < .001$, $N = 9,229$), approximately 33% fewer views. This estimate describes the full observed difference in accumulated views between women’s and men’s channels. It is the benchmark against which the subsequent columns separate production differences, topic sorting, creator-side language, and audience treatment.

[Insert Table 3 about here]

Observable channel characteristics explain part, but not most, of the raw gap. Column 2 adds channel age, log video count, HD quality, the non-subscriber-trailer indicator, and log total comments. With these controls included, the coefficient on *woman* falls to -0.2434 , approximately 22% fewer views. Video count accounts for much of the reduction from the raw estimate, consistent with women having fewer videos and younger channels. Log total comments works in the opposite direction as a suppressor, because women’s channels receive disproportionately many comments relative to their views. The remaining gap is therefore not simply a difference in channel age, production volume, HD availability, trailer strategy, or comment volume.

Horizontal category sorting explains little of the controlled gap. Column 3 adds YouTube category fixed effects, so women and men are compared within the same content category while still holding constant channel age, log video count, HD quality, the non-subscriber-trailer indicator, and log total comments. The coefficient on *woman* is -0.2167 . Category fixed effects remove only

about 11% of the controlled gap, leaving roughly 89% within categories. This result rules out the original sorting account as the main explanation: women are not receiving fewer views primarily because they cluster in lower-view content categories.

Creator-side title language does not materially account for the view gap. Column 4 adds title-LIWC tone and title-LIWC swear, measured from creators' own video titles, to the model with channel age, log video count, HD quality, the non-subscriber-trailer indicator, log total comments, and category fixed effects. The coefficient on *woman* moves only from -0.2167 to -0.2066 ($N = 8,502$). This modest attenuation indicates that women's own title language, although gendered in tone and profanity, is not the channel through which most of the view gap operates.

Audience comment profanity explains a smaller but theoretically consistent portion of the gap. Column 5 adds comment-LIWC profanity to the column 4 specification, again holding constant channel age, log video count, HD quality, the non-subscriber-trailer indicator, log total comments, category fixed effects, and title-LIWC controls. The coefficient on *woman* falls to -0.1788 ($p < .001$). Comment profanity is positively associated with views, $+0.1329$ ($p < .001$), meaning that channels receiving more profane comments also receive more views. Because women receive less profanity in their comments, this positive association makes lower-arousal treatment a disadvantage in view accumulation.

Comment tone accounts for the remaining gender association. Column 6 adds comment-LIWC tone to the model with channel age, log video count, HD quality, the non-subscriber-trailer indicator, log total comments, category fixed effects, title-LIWC controls, and comment profanity. The coefficient on *woman* drops to -0.0288 and is reduced to statistical non-significance ($p = .44$). In the same model, comment tone enters at -0.4381 ($p < .001$), while comment profanity remains positive at $+0.0484$ ($p = .006$). Thus, more positive comment tone is associated with fewer views, more profane comment language is associated with more views, and jointly these audience-treatment measures mediate the gender view gap that remains after channel characteristics, topic, title language, and comment volume are held constant. Figure 1 plots the *woman* coefficient and its 95% confidence interval across these six specifications.

[Insert Figure 1 about here]

The same progression holds without category indicators. Appendix Table A5 reports the six-column sequence without category fixed effects. The endpoint coefficient on *woman* is +0.0153 and statistically indistinguishable from zero. This parallel specification matters because it shows that the attenuation is not an artifact of absorbing category-level variation before estimating the comment-language channel. Whether category fixed effects are included or omitted, the gender coefficient is reduced to statistical non-significance once comment tone and comment profanity enter the model.

Quantifying the Mediated Share

Bootstrap mediation estimates confirm that comment language accounts for the gender view gap in the main specification. Table 4 reports 5,000-replication percentile bootstrap estimates using the model with category fixed effects, channel age, log video count, HD quality, the non-subscriber-trailer indicator, log total comments, and the comment-language mediators ($N = 9,097$). The total association between the *woman* indicator and log views is -0.221 . The indirect effect through comment tone is -0.177 [$-0.204, -0.152$], and the indirect effect through comment profanity is -0.015 [$-0.026, -0.005$]. The total indirect effect is -0.192 [$-0.220, -0.167$]. The direct effect is -0.029 and statistically indistinguishable from zero [$-0.107, +0.043$]. The share mediated by comment language is 87.0% [63.9%, 128.0%]. A generalized structural equation model reproduces the indirect effects to four decimal places.

[Insert Table 4 about here]

Figure 2 depicts the estimated mediation paths from creator gender through comment tone and comment profanity to log views.

[Insert Figure 2 about here]

The mediated share is not dependent on the inclusion of category fixed effects. In the corresponding model without category fixed effects, the mediated share is 106.1% [82.8%, 141.3%]. The confidence interval includes 100%, and the direct effect is statistically indistinguishable from zero. This result complements Table 3: the comment-language channel is not merely absorbing differences across categories. Women are differentially exposed to a platform-wide language gradient within categories, and that gradient accounts for the view gap in the mediation estimates.

A broader parallel mediation model shows that audience treatment dominates creator-side content and packaging channels when the channels compete directly. In the integrated specification with category fixed effects, channel age, log video count, HD quality, the non-subscriber-trailer indicator, and log total comments, comment treatment accounts for 89.6% of the gender view gap, with a bootstrap indirect effect of -0.2201 [-0.254 , -0.188]. Title and description content, transcript content, and thumbnail packaging are all non-negative channels, either null or statistically significant suppressors. The direct effect after all channels enter is -0.0746 [-0.159 , $+0.008$], statistically indistinguishable from zero. This partition is consistent with the negative arousal bonus: more negative, higher-arousal audience response travels more readily through human resharing and the recommendation algorithm than positive, lower-arousal treatment.

The serial content-to-comment path exists, but it is smaller than the treatment path not reducible to creator content. Women produce warmer creator-side content, warmer creator-side content is associated with warmer comments, and warmer comments are associated with fewer views. In the composite tone model, the serial path from creator tone to comment tone to views accounts for 17.8% of the gap [-0.047 , -0.033]. The treatment path not operating through creator content accounts for 65.7% [-0.171 , -0.123]. The direct effect after the tone paths enter is statistically indistinguishable from zero. Thus, creator-side tone contributes to the warmer treatment women receive, but the larger association remains audience treatment that is not reducible to measured creator content.

The per-surface serial paths show the same pattern at a finer level. Descriptions account for 11.3% of the gap, titles for 9.8%, transcripts for 7.2%, and thumbnail overlay text for 2.0%.

Each path is statistically significant, but none is large enough to make creator content the main mechanism. The analogous profanity serial path accounts for less than 1% of the gap. Profanity has the expected signs, since women use and receive less profanity and profanity in comments is associated with more views, but the magnitude is too small to carry the result. The demand-side mechanism is therefore principally a tone channel, with profanity as a secondary arousal component.

A Platform-Wide Gradient

The association between comment language and views is platform-wide, not confined to women's channels. Appendix Table A4 drops the gender indicator and estimates the same language-to-views specification on all channels. Comment tone predicts log views at -0.43 , while comment profanity predicts log views at $+0.20$. Channels with more positive comment sections receive fewer views, and channels with more profane comment sections receive more views, regardless of creator gender. Women are disadvantaged because they are differentially exposed to this general gradient: their channels attract more positive and less profane comments.

This platform-wide gradient is central to the interpretation of the gender coefficient. The models do not imply that women's channels are treated differently by the view-production process once comment language is held constant. Instead, women receive a different kind of audience treatment, and that treatment is attached to a lower-view position in the platform's amplification environment. The same affective signature that marks benevolently biased treatment, more positive and lower-arousal response, is also the signature associated with weaker propagation through human resharing and the recommendation algorithm.

Temporal Ordering

Comment language measured earlier predicts subsequent views in the same direction as the cross-sectional models. Appendix Table A1 uses comment tone and swearing measured in the first wave to predict viewership measured eight months later, conditional on first-wave viewership, channel age, log video count, HD quality, the non-subscriber-trailer indicator, and

log total comments, on the mediator-complete two-wave sample of $N = 9,055$. First-wave tone predicts subsequent views at -0.0343 ($p < .0001$) when entered alone and at -0.0306 ($p = .0001$) when profanity is included. The profanity coefficient is $+0.0137$ ($p = .097$). Comment language measured before the outcome window therefore carries the same signs as in the main cross-sectional models, net of accumulated baseline viewership and channel characteristics.

The two-wave model does not show a new gender divergence over the eight-month window. In the same sample, the creator-gender coefficient is statistically indistinguishable from zero. This null result is substantively informative because the dependent variable is subsequent accumulated viewership conditional on first-wave views. It suggests that the gender gap documented in the main tables resides primarily in the accumulated stock of views rather than in a newly emerging divergence over the eight-month window. The temporal evidence therefore supports the ordering of the language-to-views association: earlier reception predicts later views conditional on baseline views.

Attrition does not explain the two-wave result. Attrition is unrelated to creator gender in a probit model ($p = .29$), with a marginal effect of -0.2 percentage points on a 99% retention base. The temporal test is therefore not driven by women being differentially lost from the follow-up sample. This check is important because the temporal specification is used to assess ordering: if attrition were gendered, the lagged result could reflect sample composition rather than the persistence of the comment-language gradient.

A third wave of channel statistics, collected approximately twelve months after the first wave, reproduces the ordering at a longer horizon. Predicting this later viewership from first-wave comment language, conditional on first-wave viewership and the same controls, first-wave tone enters at -0.042 ($p < .001$, $N = 9,082$), marginally stronger than the eight-month estimate, and the association survives content-category fixed effects (-0.039 , $p < .001$). First-wave profanity is again positive and not statistically significant ($+0.014$, $p = .12$), and the creator-gender coefficient is statistically indistinguishable from zero once the language mediators are included. The temporal-ordering result therefore replicates at the twelve-month horizon, the longer separation between the

language measure and the outcome strengthening rather than weakening the ordering evidence (companion to Appendix Table A1).

Alternative Explanations

Table 5 lists the alternative explanations considered here and the empirical pattern that distinguishes each from the comment-language mechanism. I take them in turn.

[Insert Table 5 about here]

Horizontal sorting into different content categories does not account for the gender view gap. The fixed-effects columns of Table 3 show that adding YouTube category fixed effects moves the controlled coefficient only from -0.2434 to -0.2167 . An Oaxaca-Blinder decomposition reaches the same conclusion: the composition component is statistically indistinguishable from zero ($p = .17$). Because roughly 89% of the controlled gap remains within categories, the view gap attaches to women creators within the same content areas rather than to women's overrepresentation in lower-view categories.

Production volume explains part of the raw gap, but a large residual remains. Women have younger channels and fewer videos, and roughly one-third of the raw gap reflects those production differences. This is visible in the movement from the raw estimate of -0.4074 to the controlled estimate of -0.2434 when channel age, log video count, HD quality, the non-subscriber-trailer indicator, and log total comments enter the model. The residual survives at $p < .001$. Thus, production is a meaningful descriptive difference, but it does not eliminate the view gap among channels with comparable age, output, HD quality, trailer strategy, and comment volume.

Creator-controlled content and packaging do not carry the view gap. The strongest supply-side test uses four public creator-controlled surfaces: titles, descriptions, transcripts, and thumbnail packaging. Across these surfaces, the gender difference in content is clear: women's creator-side content is warmer and lower-arousal in titles, descriptions, transcripts, and packaging. The view-predictive leg largely fails. No single creator surface explains more than about 9% of the view gap,

and in the integrated parallel mediation, comment treatment accounts for about 90% of the gap while title and description content, transcript content, and thumbnail packaging are net suppressors. Creator content and packaging therefore do not explain why women receive fewer views.

Within-channel video-level models reach the same conclusion from a different source of variation. Within a creator, posting a lower-tone or higher-profanity video is not significantly associated with more views. The pooled association between creator content language and views appears between creators rather than within creators. This matters because a strong supply-side account would predict that when the same creator posts a more negative or higher-arousal video, that video should receive more views. The within-channel evidence does not support that pattern.

Comment treatment is not merely a proxy for creator content. With all four creator surfaces included, 66% to 76% of the gender difference in comment treatment remains. The comment-tone gap falls from +11.49 to +7.58, the VADER compound gap from +0.176 to +0.124, and the comment-swear gap from -0.254 to -0.193 . A broader LIWC sweep reaches the same conclusion. No single content metric fully explains the comment-tone gap, and a kitchen-sink model explains only 32.4%, leaving a residual woman coefficient of +0.435 ($p < .001$). Audience treatment is therefore not simply a readout of what creators write, say, or show.

The gap is specific to views rather than subscribers. The raw subscriber gap is -0.103 ($p = .008$), but when channel age, log video count, HD quality, the non-subscriber-trailer indicator, and log total comments are included, the coefficient falls to -0.027 and is statistically indistinguishable from zero. The contrast between views and subscribers helps localize the disadvantage. Women do not show a comparable controlled gap in this commitment-related outcome, even though they receive substantially fewer views.

An engagement-quality explanation runs in the opposite direction. Table 2 shows that women's channels are higher on every reported engagement measure, including the API-derived engagement rate of 3.89% versus 3.24%. If the view gap reflected lower-quality engagement from audiences who arrive, the descriptive pattern should show weaker interaction on women's channels. It instead shows stronger interaction and more benevolently biased comment language.

Engagement-rate differences therefore sharpen the puzzle rather than resolving it.

Commenter-gender composition does not account for the treatment pattern. Male-coded commenters average a tone of 42.9 on men's channels and 53.0 on women's channels; female-coded commenters average 44.2 on men's channels and 54.4 on women's channels. Profanity follows the same pattern: 0.440 versus 0.223 for male-coded commenters and 0.389 versus 0.171 for female-coded commenters. Commenters of both genders write more positively and less profanely on women's channels. The gendered comment pattern is therefore a treatment norm, not merely a compositional difference in who comments.

Commenter-gender composition also fails as a competing mediation channel. In the integrated mediation, commenter-gender composition is a suppressor rather than an explanatory channel: women attract a more female audience, and a more female audience is associated with more views. Adding commenter-gender composition leaves the comment-treatment channel intact, 92.3% $[-0.2624, -0.1948]$. This result addresses the possibility that the language channel is simply capturing audience composition. The composition channel is statistically real, but it runs in the opposite direction from an explanation of women's fewer views.

A method-specific sentiment artifact does not explain the result. Replacing LIWC comment tone with VADER compound sentiment reproduces the pattern using an independent sentiment measure. The cross-method serial path remains significant at 15.9%, while the treatment-not-via-content path accounts for 92.8% and the direct effect is statistically indistinguishable from zero. Because the comment-side measure is no longer LIWC, this result reduces concern that the mechanism is an artifact of shared dictionary construction between creator content and audience comments.

Off-platform multi-homing does not account for the gap. Women link out to external platforms more than men, 20.4% versus 15.6% of channels. Conditioning on outbound linking moves the controlled gender coefficient only from -0.2434 to -0.2449 when channel age, log video count, HD quality, the non-subscriber-trailer indicator, and log total comments are included. If women's lower YouTube views were explained by greater diversion of audience activity to other

platforms, the coefficient should attenuate when outbound linking is held constant. It does not.

A backlash or role-congruity account does not fit the heterogeneity pattern. A backlash account would expect the tone gradient to bind differently for women depending on the gender-typing of the content category. The three-way interaction of the *woman* indicator, comment tone, and an indicator for female-typed categories is 0.026 ($p = .73$). The tone-by-category pattern applies to creators of both genders. The amplification gradient therefore appears to attach to the affective signal itself, not to a woman-specific backlash dynamic in particular categories.

Standard diagnostics do not indicate that the results are artifacts of influential observations, collinearity, missingness, comment length, or the skewed profanity distribution. The maximum Cook's distance is 0.008, and the maximum VIF is 1.10. Missingness on the comment-language measures is unrelated to creator gender ($p = .66$). The mediation results are robust to controlling for mean comment length. The positive profanity association is also robust across five alternative codings of the right-skewed swear distribution. These diagnostics support the stability of the main estimates.

Sensitivity to Unmeasured Confounding

The final question is how much unmeasured confounding the estimates could withstand. Appendix Table A3 reports sensitivity benchmarks following Cinelli and Hazlett (2020). To reduce the controlled gender gap to zero, an unmeasured confounder would need to explain 7.4% of the residual variance of both creator gender and log views. The strongest observed controls, video count and channel age, explain about 0.5% on the gender side. A hypothetical confounder as strongly related to creator gender as video count would leave the gap at -0.096 ($t = -3.8$). The headline gender association therefore exceeds the confounding strength exhibited by the measured controls.

The tone pathway is more robust to unmeasured confounding than the profanity pathway. Appendix Table A3 shows that overturning the tone pathway requires a confounder explaining 21% of the residual variance of both comment tone and log views. The corresponding benchmark for

the profanity pathway is 4.4%. This difference is consistent with the mediation estimates: tone is the primary mediator, while profanity is a smaller secondary channel. The sensitivity analysis therefore supports the main interpretation while preserving the distinction between the stronger tone pathway and the weaker profanity pathway.

DISCUSSION

The starting assumption of this paper was the one that has organized much of the demand-side literature on gender and entrepreneurship: women's entrepreneurial disadvantage is theorized to run through unfavorable evaluation by resource holders, and because the party that evaluates is usually also the party that allocates, favorable evaluation is generally assumed to help. The findings deny that assumption. Women creators received more positive, less profane, softer treatment, drew comparable or higher engagement, and showed no controlled subscriber gap, yet they received substantially fewer views, about 33% fewer in the raw data and about 22% fewer after holding constant channel age, log video count, HD quality, the non-subscriber-trailer indicator, and log total comments. The mechanism that accounts for this denial is not unfavorable evaluation but the association between gendered reception and dual amplification: human resharing and the inferred engagement-optimizing recommender are more closely associated with negative, high-arousal signals than with positive, lower-intensity ones, so the softer treatment women receive is statistically converted into lower visibility.

Unfavorable Evaluation Is Not Necessary for Disadvantage

The closest contribution is to research on gender and entrepreneurship. Prior work has shown that women entrepreneurs often receive fewer resources because investors, evaluators, and other gatekeepers assess them less favorably, ask them different questions, or apply gendered standards under uncertainty (Kanze et al., 2018; Brooks et al., 2014; Ewens and Townsend, 2020; Botelho and Abraham, 2017). My findings affirm the existence of disadvantage but qualify the mechanism through which it can appear. Women creators were not observably rejected by au-

diences: their channels had higher engagement, more positive comments, less profanity, and no controlled subscriber gap (-0.027 , ns), while the view gap remained large and statistically precise. This advances the literature by identifying a boundary condition on the unfavorable-evaluation account. Negative evaluation is not necessary for gendered entrepreneurial disadvantage when an intermediary helps translate evaluative signals into visibility.

This boundary condition also qualifies democratization accounts of digital entrepreneurship. Digital platforms lower entry barriers, reduce dependence on traditional investors and publishers, and allow creators to reach audiences without prior approval (Cutolo and Kenney, 2021). But lowering access barriers need not remove inequality. In this setting, inequality is shifted from entry to amplification: women can enter, attract engagement, and convert viewers into subscribers at parity once channel characteristics are held constant, yet still receive fewer cumulative views. The result therefore complements work showing that platform-dependent entrepreneurs face new forms of structural dependence, because visibility is an allocated resource even when participation appears open.

The findings further clarify the relationship between audience support and resource conversion. In some entrepreneurial settings, favorable audience response is directly tied to resource allocation, as in crowdfunding contexts where audiences both evaluate and fund (Greenberg and Mollick, 2017). In other settings, staged evaluation can give women early support but withhold later rewards (Botelho and Poskanzer, 2026). YouTube differs because the audience response is only one input into a larger visibility system. The audience evaluates, attends, comments, and shares, while the inferred engagement-optimizing recommender helps determine whether prior reception is surfaced to later audiences. This structure means that positive treatment can coexist with worse outcomes not because audience response is irrelevant, but because the signals associated with that response are not equally associated with visibility.

Softer Treatment Can Be Costly

The second contribution is to research on benevolent and ambivalent sexism. Prior work has emphasized that ostensibly positive treatment can preserve inequality by protecting women from challenge, muting resistance, inflating feedback, or channeling women away from developmental opportunities (Glick and Fiske, 2001; King et al., 2012; Jampol and Zayas, 2021; Hopkins-Doyle et al., 2019; Bareket and Fiske, 2023). My findings extend that logic to a new consequence: unequal visibility in an attention market. Women received comment tone of 53.5 compared with 43.1 for men, and profanity of 0.19 versus 0.43 percent of comment words. This more positive, less profane, softer treatment is not the absence of bias. It is a gendered form of reception that is associated with material disadvantage when attention is allocated through signals that privilege intensity.

This extension matters because the harm observed here does not operate through the classic pathway of lower ratings or explicit exclusion. The treatment women receive looks, on its surface, kinder. Yet kinder treatment is associated with fewer views because the more positive and less profane comment environment is less aligned with the signals associated with human resharing and the inferred engagement-optimizing recommender. In this respect, the study adds a visibility-based cost to the known costs of benevolent sexism. Softer treatment can be materially disadvantaging not only because it withholds challenge or candor, but because it is less convertible into the attention that digital entrepreneurs need.

Amplification Need Not Have a Gendered Objective

The third contribution is to research on platform attention and algorithmic amplification. Prior work shows that engagement-based systems can generate unequal outcomes without explicitly discriminatory objectives and that ranking systems can amplify some forms of content more than others (Lambrecht and Tucker, 2019; Huszár et al., 2022; Covington et al., 2016; Milli et al., 2025). My evidence is consistent with this tradition but adds a gendered mechanism. When gender is dropped from the model, the language-to-views gradient remains platform-wide: comment tone is associated with fewer views, about -0.43 , and profanity is associated with more views, about

+0.20. Women are differentially exposed to that general gradient because their channels receive more positive and less profane comments.

This point is important for locating the mechanism. The three-way interaction of creator gender, comment tone, and female-typed category is not statistically significant ($\beta = 0.026$, $p = .73$), indicating that the gradient attaches to the affective signal rather than to women creators as such. The evidence therefore does not require a gendered algorithmic objective or an individually discriminating actor. It is sufficient that gendered treatment norms are associated with different reception signals and that dual amplification, through human resharing and the inferred engagement-optimizing recommender, is associated with negative and high-arousal signals. The novel bridge is precisely this joining of benevolent-sexism treatment norms to algorithmic amplification: gendered reception and engagement-optimizing amplification together author the disadvantage.

Scope Conditions

The argument applies most directly to established creators in an algorithmically mediated attention market. The sampling frame consists of identifiable U.S., English-language YouTube creators indexed above roughly 2,000 subscribers, a rough subscriber threshold at which monetization becomes accessible, although it does not map exactly to the YouTube Partner Program gate. These creators are not casual entrants at the first moment of posting. They are established enough that views plausibly matter for income, sponsorship, and future discoverability.

Two structural conditions are central. First, visibility must be shaped by engagement-optimizing recommendation and by human resharing, such that evaluative signals can be differentially surfaced. Second, creator gender must be legible to audiences, because the argument concerns gendered reception. The study is also bounded by platform, country, language, and outcome: it examines YouTube, U.S. and English-language creators, and a cross-sectional stock of cumulative channel views rather than a pure flow of newly allocated impressions. Because reception and amplification are recursive, I do not claim that the disadvantage is visibility-only in an absolute sense. The cross-sectional evidence localizes the observed gap to views, but over time lower visibility may

also be associated with commitment, subscribers, posting incentives, and creator exit.

Limitations

The principal limitation is that the design relies on temporal ordering rather than longitudinal panel variation, and is therefore observational rather than causal. I do not observe an exogenous change in comment tone, profanity, recommendation weighting, or audience composition. The credibility of the interpretation rests on three legs. First, the eight-month lagged-dependent-variable test shows that first-wave tone is associated with later views conditional on first-wave views and the five channel controls ($\beta = -0.031$, $p = .0001$), and the twelve-month replication is slightly stronger ($\beta = -0.042$, $p < .001$), surviving category fixed effects at -0.039 ($p < .001$). Second, the Cinelli-Hazlett sensitivity analysis indicates that an omitted confounder would need to explain about 7.4% of the residual variance of both creator gender and views to nullify the gap, while the tone pathway would require about 21%; the profanity pathway is weaker, at about 4.4%, consistent with profanity being a secondary channel. Third, the results are estimated with and without category fixed effects.

A second limitation concerns owner-only platform measures. No public-API study can observe click-through rates and thumbnail impressions, watch time and retention, or viewer demographic composition. Viewer demographics are distinct from the commenter gender that I do observe. I partially proxy the packaging channel through thumbnails, and those analyses indicate that women's thumbnails are more positive in expression and that packaging does not account for the gap, but thumbnail features are not click-through rates. The dependent variable also blends Shorts and long-form views in cumulative channel totals. I treat this as a portfolio-level measure of visibility, but the public data do not allow the gap to be decomposed by format.

A third limitation concerns selection into the analytic frame. The sample is male-skewed, about 65% men, and it excludes very early-stage creators, non-individual accounts, and creators whose gender is not publicly legible. Surviving women in this established-creator frame are likely positively selected, since they cleared the same threshold in a tougher environment; this selection

should bias the observed gap toward zero, making the estimated visibility gap conservative. This issue is distinct from within-sample missingness, which is small, 1.46%, and unrelated to creator gender ($p = .66$). Residual confounds also remain after the four-surface supply-side rule-out. Unmeasured creator-side quality could still matter, although transcript language is now measured and 66% to 76% of the comment-tone gap survives full four-surface content controls. Residual reverse causality is mitigated, not eliminated, by the lagged design.

Implications

The practical implication is that advice focused on changing women's content targets the wrong locus. No single creator-controlled surface explains more than about 9% of the view gap, and in the integrated parallel mediation comment treatment accounts for about 90% while titles, descriptions, transcripts, and thumbnail packaging are null or suppressing channels. This does not mean creator strategy is irrelevant, but it means that telling women to fix their content does not address the main association observed here. The locus is the ranking architecture and the broader amplification environment, including human resharing and the inferred engagement-optimizing recommender, not simply individual creator behavior.

The mechanism also has implications beyond YouTube. Any setting in which an opaque algorithmic intermediary sits between an evaluative act and a resource allocation can weaken the assumption that positive reception yields positive outcomes. Hiring platforms, promotion systems, creator marketplaces, and content distribution systems all potentially separate how producers are received from how visible or resourced they become. The stakes are material. Holding revenue per view fixed, a visibility gap maps roughly proportionally onto an advertising-revenue gap, and independent non-U.S. evidence finds that the gender gap appears in creator earnings, not only in views (Verwiebe et al., 2025). I therefore treat visibility not as a cosmetic metric but as a resource through which inequality can accumulate.

Future Directions

The most direct next test is an early-stage or cold-start longitudinal panel. My sample observes established creators, not the first period in which a new channel is being classified, recommended, and discovered. A cold-start design could test whether early-stage women face distinct algorithmic treatment before the platform has rich behavioral histories, whether softer reception emerges immediately or only after audiences sort in, and whether early reception predicts subsequent visibility trajectories differently for women and men.

A second direction is to test whether the recursive loop reaches commitment over time. The current cross-sectional evidence localizes the observed gap to views, while subscribers show no comparable controlled gap. But if lower visibility reduces the pool of potential subscribers and reshapes creator incentives, the same process may eventually be associated with commitment, posting frequency, sponsorship opportunities, and persistence. Longer panels could examine whether early more positive, less profane reception is associated with long-run viewership paths, subscriber accumulation, and creator exit.

A third direction is to exploit exogenous variation in recommendation weighting. Platform redesigns, policy changes, or natural experiments that alter the relative weight placed on watch time, comments, shares, or other engagement signals would allow direct causal tests of the amplification mechanism. Comparative work across platforms would also be valuable because platforms differ on at least one of the two structural conditions: some rely more heavily on recommendation, others on follower networks; some make creator gender more legible, others less so. Such variation would clarify when softer treatment is merely different reception and when it becomes an amplified visibility disadvantage.

CONCLUSION

This paper began from a tension in the study of gender, entrepreneurship, and platform-mediated work. Demand-side accounts have shown that women entrepreneurs are often disad-

vantaged because evaluators receive them less favorably. But in attention markets, evaluation and allocation are not held by the same actor. Audiences respond, engage, comment, and share, while an inferred engagement-optimizing recommender helps allocate visibility to later audiences. In that setting, favorable reception need not be associated with favorable outcomes. The pattern here is that women creators receive more positive, less profane, softer treatment, draw comparable or higher engagement and subscriber commitment, and yet are still associated with lower visibility. The disadvantage is therefore not well described as simple audience rejection, nor as evidence that women produce less compelling content.

The more unsettling implication is that inequality can be authored jointly by gendered reception and dual amplification, through human resharing and the inferred engagement-optimizing recommender, without any actor intending the result. Softer treatment is not merely kinder treatment when visibility depends on the signals that attention markets amplify. It is a gendered form of reception that is associated with fewer views even as the subscriber gap is reduced to statistical non-significance. For scholars of organizations, the point is broader than YouTube: when algorithmic intermediaries stand between evaluation and allocation, positive reception can lose its protective force. In algorithmically mediated attention markets, women can be received more favorably and committed to just as strongly, yet made less visible, because gendered reception and engagement-optimizing amplification jointly author the disadvantage, with no actor intending it.

REFERENCES

- Anglin, A. H., Short, J. C., Drover, W., Stevenson, R. M., McKenny, A. F., and Allison, T. H. (2018). The power of positivity? the influence of positive psychological capital language on crowdfunding performance. *Journal of Business Venturing*, 33(4):470–492.
- Bareket, O. and Fiske, S. T. (2023). A systematic review of the ambivalent sexism literature: Hostile sexism protects men’s power; benevolent sexism guards traditional gender roles. *Psychological Bulletin*, 149(11-12):637–698.
- Baron, R. M. and Kenny, D. A. (1986). The moderator-mediator variable distinction in social psychological research: Conceptual, strategic, and statistical considerations. *Journal of Personality and Social Psychology*, 51(6):1173–1182.
- Berger, J. and Milkman, K. L. (2012). What makes online content viral? *Journal of Marketing Research*, 49(2):192–205.
- Botelho, T. L. and Abraham, M. (2017). Pursuing quality: How search costs and uncertainty magnify gender-based double standards in a multistage evaluation process. *Administrative Science Quarterly*, 62(4):698–730.
- Botelho, T. L. and Poskanzer, E. (2026). Getting in the door versus winning it all: How gendered outcomes change across evaluation stages in entrepreneurship. *Administrative Science Quarterly*. Forthcoming; confirm volume, issue, pages, and DOI at proof stage.
- Boyd, R. L., Ashokkumar, A., Seraj, S., and Pennebaker, J. W. (2022). The development and psychometric properties of liwc-22. Technical report, University of Texas at Austin.
- Bradley, M. M. and Lang, P. J. (1999). Affective norms for English words (ANEW): Instruction manual and affective ratings. Technical Report Technical Report C-1, NIMH Center for the Study of Emotion and Attention, University of Florida, Gainesville, FL.

- Brady, W. J., Wills, J. A., Jost, J. T., Tucker, J. A., and Van Bavel, J. J. (2017). Emotion shapes the diffusion of moralized content in social networks. *Proceedings of the National Academy of Sciences*, 114(29):7313–7318.
- Brooks, A. W., Huang, L., Kearney, S. W., and Murray, F. E. (2014). Investors prefer entrepreneurial ventures pitched by attractive men. *Proceedings of the National Academy of Sciences*, 111(12):4427–4431.
- Cinelli, C. and Hazlett, C. (2020). Making sense of sensitivity: Extending omitted variable bias. *Journal of the Royal Statistical Society Series B (Statistical Methodology)*, 82(1):39–67.
- Covington, P., Adams, J., and Sargin, E. (2016). Deep neural networks for youtube recommendations. In *Proceedings of the 10th ACM Conference on Recommender Systems*, pages 191–198.
- Cunningham, S. and Craig, D. (2021). *Creator Culture: An Introduction to Global Social Media Entertainment*. NYU Press, New York.
- Cutolo, D. and Kenney, M. (2021). Platform-dependent entrepreneurs: Power asymmetries, risks, and strategies in the platform economy. *Academy of Management Perspectives*, 35(4):584–605.
- Duffy, B. E. (2017). *(Not) Getting Paid to Do What You Love: Gender, Social Media, and Aspirational Work*. Yale University Press, New Haven, CT.
- Ewens, M. and Townsend, R. R. (2020). Are early stage investors biased against women? *Journal of Financial Economics*, 135(3):653–677.
- Gamache, D. L., McNamara, G., Mannor, M. J., and Johnson, R. E. (2015). Motivated to acquire? the impact of CEO regulatory focus on firm acquisitions. *Academy of Management Journal*, 58(4):1261–1282.
- Glick, P. and Fiske, S. T. (2001). An ambivalent alliance: Hostile and benevolent sexism as complementary justifications for gender inequality. *American Psychologist*, 56(2):109–118.

- Greenberg, J. and Mollick, E. (2017). Activist choice homophily and the crowdfunding of female founders. *Administrative Science Quarterly*, 62(2):341–374.
- Hopkins-Doyle, A., Sutton, R. M., Douglas, K. M., and Calogero, R. M. (2019). Flattering to deceive: Why people misunderstand benevolent sexism. *Journal of Personality and Social Psychology*, 116(2):167–192.
- Huszár, F., Ktena, S. I., O’Brien, C., Belli, L., Schlaikjer, A., and Hardt, M. (2022). Algorithmic amplification of politics on Twitter. *Proceedings of the National Academy of Sciences*, 119(1):e2025334119.
- Hutto, C. J. and Gilbert, E. (2014). VADER: A parsimonious rule-based model for sentiment analysis of social media text. In *Proceedings of the International AAAI Conference on Web and Social Media*, volume 8, pages 216–225.
- Jampol, L. and Zayas, V. (2021). Gendered white lies: Women are given inflated performance feedback compared with men. *Personality and Social Psychology Bulletin*, 47(1):57–69.
- Janschewitz, K. (2008). Taboo, emotionally valenced, and emotionally neutral word norms. *Behavior Research Methods*, 40(4):1065–1074.
- Kanze, D., Huang, L., Conley, M. A., and Higgins, E. T. (2018). We ask men to win and women not to lose: Closing the gender gap in startup funding. *Academy of Management Journal*, 61(2):586–614.
- Khan, M. L. (2017). Social media engagement: What motivates user participation and consumption on YouTube? *Computers in Human Behavior*, 66:236–247.
- King, E. B., Botsford, W., Hebl, M. R., Kazama, S., Dawson, J. F., and Perkins, A. (2012). Benevolent sexism at work: Gender differences in the distribution of challenging developmental experiences. *Journal of Management*, 38(6):1835–1866.

- Lambrecht, A. and Tucker, C. (2019). Algorithmic bias? An empirical study of apparent gender-based discrimination in the display of STEM career ads. *Management Science*, 65(7):2966–2981.
- Love, E. G., Lim, J., and Bednar, M. K. (2017). The face of the firm: The influence of CEOs on corporate reputation. *Academy of Management Journal*, 60(4):1462–1481.
- Malhotra, S., Reus, T. H., Zhu, P., and Roelofsen, E. M. (2018). The acquisitive nature of extraverted CEOs. *Administrative Science Quarterly*, 63(2):370–408.
- Milli, S., Carroll, M., Wang, Y., Pandey, S., Zhao, S., and Dragan, A. D. (2025). Engagement, user satisfaction, and the amplification of divisive content on social media. *PNAS Nexus*, 4(3):pgaf062.
- Nguyen, N., Hideg, I., Engel, Y., and Godart, F. (2024). Benevolent sexism and the gender gap in startup evaluation. *Entrepreneurship Theory and Practice*, 48(2):506–546.
- Nonnecke, B. and Preece, J. (2000). Lurker demographics: Counting the silent. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, pages 73–80. ACM.
- Ocasio, W. (1997). Towards an attention-based view of the firm. *Strategic Management Journal*, 18(S1):187–206.
- Pan, L., McNamara, G., Lee, J. J., Haleblan, J., and Devers, C. E. (2018). Give it to us straight (most of the time): Top managers’ use of concrete language and its effect on investor reactions. *Strategic Management Journal*, 39(8):2204–2225.
- Preacher, K. J. and Hayes, A. F. (2008). Asymptotic and resampling strategies for assessing and comparing indirect effects in multiple mediator models. *Behavior Research Methods*, 40(3):879–891.
- Rahman, H. A. (2021). The invisible cage: Workers’ reactivity to opaque algorithmic evaluations. *Administrative Science Quarterly*, 66(4):945–988.
- Rathje, S., Van Bavel, J. J., and van der Linden, S. (2021). Out-group animosity drives engagement on social media. *Proceedings of the National Academy of Sciences*, 118(26):e2024292118.

- Ridgeway, C. L. (2009). Framed before we know it: How gender shapes social relations. *Gender & Society*, 23(2):145–160.
- Salganik, M. J., Dodds, P. S., and Watts, D. J. (2006). Experimental study of inequality and unpredictability in an artificial cultural market. *Science*, 311(5762):854–856.
- Stephens, R. and Robertson, O. (2020). Swearing as a response to pain: Assessing hypoalgesic effects of novel “swear” words. *Frontiers in Psychology*, 11:723.
- Tausczik, Y. R. and Pennebaker, J. W. (2010). The psychological meaning of words: Liwc and computerized text analysis methods. *Journal of Language and Social Psychology*, 29(1):24–54.
- Verwiebe, R., Osorio-Krauter, C., Fritsch, N.-S., Weißmann, S., Philipp, A., and Buder, C. (2025). Working on an algorithm-controlled platform as a content creator: What explains unequal earnings on YouTube? *European Societies*.
- Warriner, A. B., Kuperman, V., and Brysbaert, M. (2013). Norms of valence, arousal, and dominance for 13,915 English lemmas. *Behavior Research Methods*, 45(4):1191–1207.
- White, H. (1980). A heteroskedasticity-consistent covariance matrix estimator and a direct test for heteroskedasticity. *Econometrica*, 48(4):817–838.
- Yang, T., Kacperczyk, A., and Naldi, L. (2024). The motherhood wage penalty and female entrepreneurship. *Organization Science*, 35(1):27–51.
- Zuckerman, E. W. (1999). The categorical imperative: Securities analysts and the illegitimacy discount. *American Journal of Sociology*, 104(5):1398–1438.

Table 1: Descriptive statistics and pairwise correlations

Panel A: Descriptive statistics (full sample)

	N	Mean	SD	Min	Max
ln(Views)	9,230	15.092	2.124	7.275	24.178
Woman creator	9,232	0.345	0.475	0.000	1.000
Comment tone	9,100	46.730	17.980	0.000	99.000
Comment profanity	9,100	0.346	0.539	0.000	9.235
Title tone	8,502	49.671	33.430	1.000	99.000
Title profanity	9,230	0.070	0.421	0.000	10.340
Engagement rate (%)	9,230	3.464	3.337	0.000	104.080
Subscribers	9,230	184,483	877,292	2,010	31,400,000
Channel age (years)	9,231	9.7	4.9	0.3	19.9
ln(Video count)	9,230	6.301	1.168	1.609	14.527
HD quality	9,232	0.970	0.171	0.000	1.000
Non-subscriber trailer	9,232	0.593	0.491	0.000	1.000
ln(Total comments)	9,230	4.090	2.014	0.000	10.825

Panel B: Pairwise correlations

	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
1. ln(Views)								
2. Woman creator	-0.09							
3. Comment tone	-0.28	0.27						
4. Comment profanity	0.12	-0.21	-0.31					
5. Channel age (years)	0.16	-0.09	0.12	0.02				
6. ln(Video count)	0.41	-0.10	-0.14	0.04	0.19			
7. HD quality	0.04	0.01	0.02	0.00	-0.01	-0.06		
8. Non-subscriber trailer	0.16	-0.06	-0.06	0.02	0.11	0.16	0.00	
9. ln(Total comments)	0.50	0.03	-0.05	0.01	-0.03	-0.16	0.06	0.02

Notes: Random sample drawn from the Infludata sampling frame (N varies by item owing to item-level missingness; the channel sample is 9,232). Infludata supplies the sampling frame; creator gender was hand-coded by the author, and Infludata's gender field was used only to screen for individual (non-brand) creators when building the frame. All other measures are computed from the YouTube Data API or from text collected through it. ln(Views) is the natural log of total cumulative channel views (the dependent variable); Woman creator is a 0/1 indicator. Comment tone and Comment profanity are the LIWC-22 Tone and Swear measures from audience comments; Title tone and Title profanity are the same measures from video titles (all channel-level means, unstandardized here). Engagement rate is (likes + comments) divided by views $\times 100$, computed from the YouTube Data API across each channel's twelve sampled videos (reported descriptively only — excluded from the models as a post-treatment ratio); Subscribers is the raw subscriber count. The five regression controls are channel age (years), ln(video count), HD quality, non-subscriber trailer, and ln(total comments); the non-subscriber trailer is a content-strategy indicator. Panel B reports pairwise correlations (each on the pairwise-complete sample); correlations among the content-language measures appear in the content-language appendix table.

Table 2: Descriptive statistics by creator gender

	Men	Women	Difference	<i>p</i> -value	N
ln(Views)	15.233	14.825	0.408***	0.0000	9,230
ln(Subscribers)	10.156	10.054	0.103**	0.0091	9,230
Comment tone	43.143	53.544	-10.401***	0.0000	9,100
Comment profanity	0.427	0.194	0.233***	0.0000	9,100
Engagement rate (%)	3.240	3.888	-0.649***	0.0000	9,230
Channel age (years)	10.022	9.097	0.925***	0.0000	9,231
ln(Video count)	6.383	6.145	0.238***	0.0000	9,230
HD quality	0.969	0.972	-0.003	0.4729	9,232
Non-subscriber trailer	0.614	0.554	0.060***	0.0000	9,232
ln(Total comments)	4.041	4.181	-0.140**	0.0015	9,230

Notes: Two-sample *t*-tests comparing men and women creators; Difference = men – women; stars on the difference. ln(Views) is the natural log of total cumulative channel views; ln(Subscribers) is the natural log of subscriber count. Median views: men 3,088,760 vs. women 1,982,898; median subscribers: men 17,500 vs. women 16,600. Comment tone and Comment profanity are LIWC-22 Tone and Swear from audience comments (channel-level means, unstandardized): comments on women’s channels are more positive in tone and contain less profanity. Engagement rate is (likes + comments) divided by views $\times$ 100, computed from the YouTube Data API across each channel’s twelve sampled videos; women are higher (reported descriptively only — excluded from the models as a post-treatment ratio). The five regression controls are channel age (years), ln(video count), HD quality, non-subscriber trailer, and ln(total comments). * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table 3: Creator gender and channel viewership: baseline gap, category fixed effects, content controls, and comment-language mediators

	(1) Raw gap	(2) Channel controls	(3) + Category FE	(4) + Title language	(5) + Comment profanity	(6) + Comment tone
Woman creator	-0.4074*** (0.0458)	-0.2434*** (0.0332)	-0.2167*** (0.0381)	-0.2066*** (0.0400)	-0.1788*** (0.0384)	-0.0288 (0.0372)
Comment profanity (z)					0.1329*** (0.0191)	0.0484** (0.0175)
Comment tone (z)						-0.4381*** (0.0192)
Title tone (z)				-0.0260 (0.0174)		
Title profanity (z)				-0.0237 (0.0178)		
Channel age (years)		0.0310*** (0.0032)	0.0310*** (0.0032)	0.0310*** (0.0033)	0.0314*** (0.0032)	0.0444*** (0.0031)
ln(video count)		0.8621*** (0.0146)	0.8626*** (0.0148)	0.8707*** (0.0156)	0.8682*** (0.0148)	0.8294*** (0.0145)
HD quality		0.4009*** (0.0857)	0.3376*** (0.0853)	0.3924*** (0.0889)	0.3392*** (0.0870)	0.3510*** (0.0885)
Non-subscriber trailer		0.2824*** (0.0322)	0.2756*** (0.0317)	0.2614*** (0.0332)	0.2711*** (0.0316)	0.2387*** (0.0305)
ln(total comments)		0.6075*** (0.0085)	0.5993*** (0.0085)	0.6007*** (0.0089)	0.6090*** (0.0083)	0.5972*** (0.0081)
Category fixed effects (40)	No	No	Yes	Yes	Yes	Yes
Observations	9,229	9,229	9,229	8,502	9,097	9,097
R-squared	0.008	0.507	0.528	0.531	0.538	0.568

Notes: Dependent variable: ln(total cumulative channel views). Robust standard errors in parentheses. Column 1 includes no controls. Columns 2–6 include the five channel controls shown: channel age (years), ln(video count), HD quality, non-subscriber trailer, and ln(total comments). Columns 3–6 add indicator variables for 40 content categories (coefficients suppressed). Title tone (z) and Title profanity (z) are standardized LIWC-22 Tone and Swear measured on video titles; column 4 is estimated on the subsample with nonmissing title-language measures. Comment profanity (z) and Comment tone (z) are standardized LIWC-22 Swear and Tone from audience comments; column 5 adds comment profanity, and column 6 adds comment tone. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table 4: Bootstrap mediation of the creator-gender view gap through comment language


	Estimate	95% CI (percentile)
<i>Panel A: With 40-category fixed effects (main specification)</i>		
Indirect effect via comment tone	−0.1767	[−0.2037, −0.1525]
Indirect effect via comment profanity	−0.0153	[−0.0265, −0.0046]
Total indirect effect	−0.1920	[−0.2196, −0.1668]
Direct effect (Woman creator)	−0.0288	[−0.1071, +0.0426]
Total effect (Woman creator)	−0.2208	[−0.2997, −0.1477]
Share of total effect mediated	0.870	[0.639, 1.280]
<i>Panel B: Without category fixed effects</i>		
Indirect effect via comment tone	−0.2353	[−0.2638, −0.2085]
Indirect effect via comment profanity	−0.0322	[−0.0473, −0.0177]
Total indirect effect	−0.2675	[−0.2966, −0.2400]
Direct effect (Woman creator)	+0.0153	[−0.0538, +0.0804]
Total effect (Woman creator)	−0.2522	[−0.3207, −0.1888]
Share of total effect mediated	1.061	[0.828, 1.413]

Notes: Product-of-coefficients indirect effects of Woman creator on ln(total cumulative channel views) through standardized comment tone and comment profanity (LIWC-22 Tone and Swear from audience comments), bootstrapped with B = 5,000 replications (seed 20260610); 95% confidence intervals are percentile intervals. N = 9,097 in every model. All models include the five channel controls: channel age (years), ln(video count), HD quality, non-subscriber trailer, and ln(total comments); Panel A additionally absorbs 40 content-category indicators in every equation. Share mediated = total indirect effect divided by total effect; values above one arise when the remaining direct effect is slightly above zero. A generalized structural equation model with robust standard errors reproduces the Panel B indirect effects to four decimal places (delta-method confidence intervals nearly identical).

Table 5: Alternative explanations for the gender gap in channel viewership: tests and verdicts


Rival explanation	How tested	Key result	Verdict	Detail
Horizontal sorting into low-view categories	40-category fixed effects on the controlled model; Oaxaca–Blinder decomposition	$\beta = -0.243 \rightarrow -0.217$ with category FE; category FE removes only $\sim 11\%$ of the gap ($\sim 89\%$ within-category). Oaxaca composition component ns ($p = 0.17$)	Ruled out	Table A-1
Supply-side production volume	Incremental decomposition, adding the five controls one at a time	Video count is the heaviest contributor (~ 14 points at its step: $-0.348 \rightarrow -0.209$); about one-third of the raw gap reflects production volume. Controlled residual $\beta = -0.243$ ($\sim 22\%$ fewer views) remains, $p < .001$	Partial; residual survives	Table 3
Creator’s own content (title + description language)	LIWC tone and profanity from video titles and channel descriptions added to the controlled model (common $N = 8,247$)	$\beta = -0.253 \rightarrow -0.228$; the full content channel accounts for $\sim 10\%$ of the controlled gap; coefficient remains significant at $p < .001$	Ruled out (as primary)	Table A-2
Subscriber base / audience commitment	Parallel models with $\ln(\text{subscribers})$ as the outcome	Subscriber gap: raw -0.103 ($p = .008$); with the five controls -0.027 (ns, $p = 0.36$). The view gap persists strongly	Ruled out; gap is view-specific	In text
Engagement quality (audiences value women’s content less)	Engagement-rate comparisons by creator gender	Women are higher on every engagement measure (mean 3.89% vs. 3.24%; newest uploads 7.60 vs. 6.60; all $p < .0001$)	Ruled out; opposite direction	Table 2
Commenter gender composition	Mediation re-estimated with tone and profanity measured separately for commenters coded as men and as women	The warmer, lower-arousal treatment of women’s channels holds across commenter gender; the mediation pattern is unchanged	Ruled out (composition); treatment norm holds	Tables A-9, A-9b
Off-platform multi-homing (self-promotion elsewhere)	Channel-description external-link indicator; creator gender $\times$ multi-homing interaction	Woman-creator β : $-0.2434 \rightarrow -0.2449$ with the indicator; interaction ns. Women link out <i>more</i> (20.4% vs. 15.6%, $p < .0001$)	Ruled out; opposite direction	Table A-11
Backlash / role congruity (penalty tied to gender-typed categories)	Woman creator $\times$ tone $\times$ female-typed-category triple interaction	Triple interaction $\beta = 0.026$, $p = 0.73$ (ns); the tone $\times$ female-typed pattern applies to creators of both genders	Ruled out	Table A-7
Other comment-language dimensions	Bootstrapped parallel decomposition, all ten LIWC components	Only tone and profanity carry negative indirect effects; positive-emotion, affect, and prosocial language run in women’s favor; the rest are null	Ruled out; rivals mask rather than explain	Table A-5

Notes: Dependent variable in all view models: $\ln(\text{total cumulative channel views})$; robust standard errors. Unless noted, models include the five controls: channel age (years), $\ln(\text{video count})$, HD quality, non-subscriber trailer, and $\ln(\text{total comments})$. Anchor estimates: raw creator-gender coefficient -0.408 ; with the five controls -0.243 ;



DV: $\ln(\text{total cumulative channel views})$. Robust standard errors; whiskers are 95% CIs. $N = 9,229$ (columns 1–3), 8,502 (column 4), 9,097 (columns 5–6). Columns 2–6 include the five channel controls (channel age in years, $\ln$ video count, HD, non-subscriber trailer, $\ln$ total comments); columns 3–6 add 40 content-category indicators. Comment profanity enters in column 5; comment tone joins in column 6. *** $p < 0.001$. Source: do-file 83.

Figure 1: The creator-gender coefficient on log views across the model progression. The gap is reduced to statistical non-significance once comment tone and comment profanity enter the model, holding constant channel age, log video count, HD quality, the non-subscriber-trailer indicator, log total comments, category fixed effects, and title language. Whiskers are 95% confidence intervals; $N \approx 9,097$.



Path estimates from the 40-category fixed-effects specification (do-file 86; N = 9,097; robust standard errors; all equations include the five channel controls and absorb 40 content-category indicators). Indirect effects (bootstrap, B = 5,000): via tone -0.177 [$-0.204, -0.152$]; via profanity -0.015 [$-0.026, -0.005$]; total -0.192 ; share of the total effect mediated 87% [64%, 128%]. ** $p < 0.01$, *** $p < 0.001$.

Figure 2: Mediation of the creator-gender view gap through comment tone (valence) and comment profanity (arousal), with category fixed effects and the five channel controls. Path coefficients are standardized; the direct effect of creator gender is statistically indistinguishable from zero once both comment-language mediators are included.

ONLINE APPENDIX

Supplementary tables. Each is generated by a numbered Stata do-file on the primary dataset ($N = 9,232$; analytic N varies by item owing to listwise deletion on the comment-language mediators).

Table A1: Temporal Ordering, Lagged Dependent Variable

	(1) Baseline	(2) + T1 Tone	(3) + T1 Tone + Swear
Woman	-0.0157 (0.0134)	0.0104 (0.0133)	0.0140 (0.0137)
T1 Comment Tone (z)		-0.0343*** (0.0074)	-0.0306*** (0.0077)
T1 Comment Swear (z)			0.0137 (0.0082)
T1 ln(Views)	0.9600*** (0.0061)	0.9579*** (0.0055)	0.9575*** (0.0055)
Observations	9,146	9,055	9,055
R^2	0.931	0.937	0.937

Standard errors in parentheses

Dependent variable: ln(views T2). T1 ln(views) included as lagged control.

T1 = May 2025, T2 = February 2026.

Comment language measured at T1; T2 viewership measured February 2026.

Controls: channel age, ln(video count), HD quality, non-subscriber trailer, ln(total comments).

Robust standard errors in parentheses.

* $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$

Table A2: Temporal ordering at a twelve-month horizon: first-wave comment language predicting third-wave log views.

	(1) L0	(2) L1	(3) L2	(4) L3	(5) L0 FE	(6) L1 FE	(7) L2 FE
woman	-0.0287* (0.0145)	-0.0017 (0.0149)	0.0022 (0.0152)		-0.0046 (0.0160)	0.0122 (0.0162)	0.0145 (0.0164)
z_tone_mean		-0.0418*** (0.0081)	-0.0379*** (0.0085)	-0.0377*** (0.0081)		-0.0390*** (0.0090)	-0.0368*** (0.0093)
z_swear_mean			0.0141 (0.0090)	0.0140 (0.0089)			0.0100 (0.0093)
ln_viewcount	0.9472*** (0.0067)	0.9422*** (0.0065)	0.9418*** (0.0065)	0.9418*** (0.0065)	0.9435*** (0.0069)	0.9393*** (0.0065)	0.9390*** (0.0066)
Category fixed effects	No	No	No	No	Yes	Yes	Yes
N	9,082	8,994	8,994	8,994	9,082	8,994	8,994
r2	0.914	0.918	0.918	0.918	0.915	0.918	0.918

Standard errors in parentheses

* $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$

Table A3: Sensitivity of the key estimates to unmeasured confounding

	Estimate (SE)	Partial R^2	RV ($q = 1$)	RV ($q = 1, \alpha = .05$)
<i>Panel A: Focal estimates</i>				
Woman creator (gap model, M1)	-0.243 (0.033)	0.006	0.074	0.055
Comment tone (z) (M3)	-0.404 (0.018)	0.054	0.211	0.195
Comment profanity (z) (M3)	+0.075 (0.018)	0.002	0.044	0.024
<i>Panel B: Benchmarks, partial R^2 of the observed controls</i>				
	Gap model (M1)		Mediation model (M3)	
	Woman creator	ln(video count)	Comment tone	ln(video count)
ln(video count)	0.0054	0.2744	0.0199	0.2642
Channel age (days)	0.0052	0.0099	,	,
ln(total comments)	0.0004	0.3569	0.0065	0.3771

Notes: Sensitivity quantities follow Cinelli and Hazlett (2020), computed from robust t -statistics. The robustness value RV ($q = 1$) is the minimum share of residual variance of *both* the focal variable and the outcome that an unmeasured confounder must explain to drive the estimate to zero; RV ($q = 1, \alpha = .05$) is the share required to render the estimate statistically indistinguishable from zero at the 5% level. Panel A: the controlled gap is from the five-control view model ($N = 9,229$; controls: channel age (days), ln(video count), HD quality, non-subscriber trailer, and ln(total comments)); the tone and profanity rows are from the mediation model adding standardized comment tone and profanity ($N = 9,097$); robust standard errors in parentheses. Panel B benchmarks the RVs against the observed controls: every observed control's partial R^2 with creator gender is ≈ 0.005 or less, an order of magnitude below the gap's RV of 0.074. The tone pathway is the most robust (a confounder must explain $\sim 21\%$ of the residual variance of both comment tone and views); the profanity pathway is the most sensitive (RV = 0.044), consistent with its secondary role in the mediation. A package cross-check (sensemakr) reproduces the gap RV (0.0741 on classical t); its bound analysis implies a confounder as strongly related to creator gender as ln(video count), with outcome-side partial $R^2 = 0.42$, would still leave the gap at -0.096 ($t = -3.8$). Dashes: benchmark not computed in the source log.

Table A4: Comment tone and profanity predict viewership platform-wide (no creator-gender indicator)

	(1) Tone	(2) Profanity	(3) Both
Comment tone (z)	-0.426*** (0.016)		-0.403*** (0.017)
Comment profanity (z)		0.196*** (0.020)	0.074*** (0.017)
Observations	9,097	9,097	9,097
R^2	0.549	0.519	0.550

Notes: Dependent variable: $\ln(\text{total cumulative channel views})$. Robust standard errors in parentheses. All models include the five controls: channel age (days), $\ln(\text{video count})$, HD quality, non-subscriber trailer, and $\ln(\text{total comments})$. Comment tone (z) and Comment profanity (z) are standardized LIWC-22 Tone and Swear from audience comments. No creator-gender indicator is included: these models show that the comment-language pathways operate platform-wide, for creators of both genders, rather than only through creator gender. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table A5: Viewership progression without category fixed effects

	(1) Raw gap	(2) Channel controls	(3) + Title language	(4) + Comment profanity	(5) + Comment tone
Woman creator	-0.4074*** (0.0458)	-0.2434*** (0.0332)	-0.2414*** (0.0350)	-0.1753*** (0.0340)	0.0153 (0.0335)
Comment profanity (z)				0.1793*** (0.0196)	0.0751*** (0.0176)
Comment tone (z)					-0.4045*** (0.0178)
Title tone (z)			-0.0165 (0.0171)		
Title profanity (z)			-0.0024 (0.0184)		
Channel age (years)		0.0310*** (0.0032)	0.0312*** (0.0034)	0.0318*** (0.0032)	0.0460*** (0.0031)
ln(video count)		0.8621*** (0.0146)	0.8686*** (0.0154)	0.8639*** (0.0145)	0.8142*** (0.0143)
HD quality		0.4009*** (0.0857)	0.4553*** (0.0892)	0.4036*** (0.0868)	0.4379*** (0.0876)
Non-subscriber trailer		0.2824*** (0.0322)	0.2668*** (0.0337)	0.2787*** (0.0319)	0.2428*** (0.0308)
ln(total comments)		0.6075*** (0.0085)	0.6090*** (0.0090)	0.6144*** (0.0083)	0.5987*** (0.0081)
Category fixed effects (40)	No	No	No	No	No
Observations	9,229	9,229	8,502	9,097	9,097
R-squared	0.008	0.507	0.510	0.521	0.550

Notes: Dependent variable: ln(total cumulative channel views). Robust standard errors in parentheses. This table reports the Table 3 progression without category fixed effects; gap estimates are slightly larger in magnitude without the category indicators. Column 1 includes no controls; columns 2–5 include the five channel controls shown: channel age (years), ln(video count), HD quality, non-subscriber trailer, and ln(total comments). Column 3 is estimated on the subsample with nonmissing title-language measures; column 4 adds comment profanity and column 5 adds comment tone. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table A6: The creator-content channel: title and description language account for about one-tenth of the controlled gap

	(1) Baseline	(2) +Titles	(3) +Descriptions	(4) +Both
Woman creator	-0.2530*** (0.0347)	-0.2434*** (0.0353)	-0.2300*** (0.0354)	-0.2281*** (0.0357)
Title tone (z)		-0.0238 (0.0172)		-0.0067 (0.0181)
Title profanity (z)		0.0002 (0.0179)		0.0020 (0.0205)
Description tone (z)			-0.0553** (0.0173)	-0.0531** (0.0182)
Description profanity (z)			-0.0053 (0.0132)	-0.0061 (0.0156)
Share of controlled gap accounted for	,	3.8%	9.1%	9.8%
Observations	8,247	8,247	8,247	8,247
R^2	0.515	0.516	0.516	0.516

Notes: Dependent variable: $\ln(\text{total cumulative channel views})$. Robust standard errors in parentheses. All models include the five controls: channel age (days), $\ln(\text{video count})$, HD quality, non-subscriber trailer, and $\ln(\text{total comments})$, and are estimated on the common sample with nonmissing title- and description-language measures ($N = 8,247$). Title and Description tone/profanity (z) are standardized LIWC-22 Tone and Swear measured on video titles and channel descriptions, the creator's own content language, as distinct from audience comments. Adding the full content channel (titles and descriptions together) moves the creator-gender coefficient from -0.2530 to -0.2281 , accounting for about one-tenth of the controlled gap; the coefficient remains significant at $p < .001$. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table A7: Decomposition of the creator-gender view gap across candidate comment-language mediators

	Received by women (<i>a</i>)	Predicts views (<i>b</i>)	Indirect effect (<i>a</i> × <i>b</i>)	95% CI (percentile)
Comment tone	0.4034***	−0.5800***	−0.2339	[−0.2708, −0.2008]
Comment profanity	−0.3164***	0.0401*	−0.0127	[−0.0243, −0.0014]
Moral language	−0.1524***	−0.0125	0.0019	[−0.0038, 0.0066]
Positive-emotion words	0.4446***	0.0812**	0.0361	[0.0114, 0.0595]
Affect (all emotion words)	0.3633***	0.0930**	0.0338	[0.0116, 0.0575]
Social words	0.4067***	0.0168	0.0068	[−0.0136, 0.0310]
Polite language	0.2663***	−0.0315	−0.0084	[−0.0267, 0.0085]
Prosocial language	0.2070***	0.0723*	0.0150	[0.0029, 0.0291]
Anger words	−0.0920***	−0.0116	0.0011	[−0.0026, 0.0047]
Conflict language	−0.0477*	−0.0235	0.0011	[−0.0005, 0.0038]
Total indirect effect (all ten)			−0.1592	[−0.1923, −0.1285]
Direct effect (Woman creator)			−0.0616	[−0.1397, 0.0106]
Total effect (Woman creator)			−0.2208	[−0.2997, −0.1477]
Share of total via tone and profanity			1.117	[0.819, 1.647]

Notes: Parallel multiple-mediator decomposition (Preacher and Hayes, 2008) of the creator-gender gap in ln(total cumulative channel views) across the ten LIWC-22 components measured on audience comments, estimated on the common sample (N = 9,097) with all components nonmissing. Column *a*: standard-deviation difference in the component received by women's channels, from a regression of the standardized component on the Woman creator indicator; column *b*: effect of the standardized component on ln(views) conditional on all ten components simultaneously. Every equation includes the five channel controls (channel age in years, ln video count, HD quality, non-subscriber trailer, ln total comments) and absorbs 40 content-category indicators; robust standard errors for the stars on *a* and *b*. The indirect effect *a* × *b* is each component's contribution to the gap, bootstrapped with B = 5,000 replications (seed 20260610); 95% intervals are percentile intervals. A component can account for the view deficit only if its indirect effect is negative with an interval excluding zero. Components are correlated by construction (LIWC Tone is itself computed from emotional language), which the parallel specification accommodates; intervals reflect that correlation. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table A8: Commenter Gender Decomposition of Tone and Swear Mediation

	(1) Baseline	(2) Aggregate	(3) Men Commenters	(4) Women Commenters	(5) Both
Woman	-0.243*** (0.033)	0.015 (0.033)	-0.022 (0.033)	-0.041 (0.033)	0.004 (0.033)
Comment Tone (z)		-0.404*** (0.018)			
Comment Swear (z)		0.075*** (0.018)			
Tone, men commenters (z)			-0.384*** (0.018)		-0.269*** (0.027)
Swear, men commenters (z)			0.060** (0.019)		0.044* (0.018)
Tone, women commenters (z)				-0.373*** (0.017)	-0.166*** (0.025)
Swear, women commenters (z)				0.052** (0.019)	0.027 (0.018)
Observations	9,229	9,097	9,069	9,045	9,022
R^2	0.507	0.550	0.545	0.543	0.549

Standard errors in parentheses

Dependent variable: $\ln(\text{total cumulative views})$.

Robust standard errors in parentheses.

M-lib = men commenters (liberal coding, 77% assigned).

F-lib = women commenters (liberal coding, 77% assigned).

Controls: channel age, $\ln(\text{video count})$, HD quality, non-subscriber trailer, $\ln(\text{total comments})$.* $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$

Table A9: Commenter Gender Decomposition: Treatment Norms 2×2

	(1) Tone (Men Comm.)	(2) Tone (Women Comm.)	(3) Swear (Men Comm.)	(4) Swear (Women Comm.)
Woman Creator	10.119*** (0.411)	10.232*** (0.418)	-0.217*** (0.013)	-0.218*** (0.014)
Mean (Man Creator)	42.884	44.159	0.440	0.389
Mean (Woman Creator)	53.003	54.391	0.223	0.171
Observations	9,072	9,048	9,072	9,048

Standard errors in parentheses

Robust standard errors in parentheses.

Each column regresses the commenter-gender-specific LIWC measure on creator gender.

Coefficient on Woman Creator = difference in comment patterns toward women vs. men creators.

Liberal gender coding (77% of commenters identified).

* $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$

* $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$

Table A10: The view gap is unchanged by off-platform multi-homing

	(1) Baseline	(2) +Multi-homing	(3) +Interaction
Woman creator	-0.2434*** (0.0332)	-0.2449*** (0.0332)	-0.2598*** (0.0367)
Multi-homing (0/1)		0.0303 (0.0410)	-0.0001 (0.0518)
Woman creator $\times$ Multi-homing			0.0785 (0.0847)
Observations	9,229	9,229	9,229
R^2	0.507	0.507	0.507

Notes: Dependent variable: $\ln(\text{total cumulative channel views})$. Robust standard errors in parentheses. All models include the five controls: channel age (days), $\ln(\text{video count})$, HD quality, non-subscriber trailer, and $\ln(\text{total comments})$. Multi-homing (0/1) indicates that the channel description links to at least one other social platform. Adding the multi-homing indicator leaves the creator-gender coefficient essentially unchanged ($-0.2434 \rightarrow -0.2449$), and the creator-gender $\times$ multi-homing interaction is not statistically significant. Women link off-platform *more* than men (any social link: 20.4% vs. 15.6%, $p < .0001$), so off-platform self-promotion cannot account for the gap. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table A11: Comment-language mediation is robust to controlling for comment length

	(1) M1	(2) M1 +Length	(3) M3	(4) M3 +Length
Woman creator	-0.2524*** (0.0331)	-0.2851*** (0.0325)	0.0152 (0.0335)	-0.0441 (0.0333)
ln(mean words per comment)		-0.5482*** (0.0313)		-0.3982*** (0.0311)
Comment tone (z)			-0.4054*** (0.0178)	-0.3550*** (0.0179)
Comment profanity (z)			0.0746*** (0.0176)	0.0599*** (0.0172)
Observations	9,096	9,096	9,096	9,096
R^2	0.514	0.534	0.550	0.560

Notes: Dependent variable: ln(total cumulative channel views). Robust standard errors in parentheses. All models include the five controls: channel age (days), ln(video count), HD quality, non-subscriber trailer, and ln(total comments); coefficients on the controls are suppressed. All four models are estimated on the common sample with nonmissing comment-length measures ($N = 9,096$). ln(mean words per comment) is the natural log of the channel-level mean comment length. Women receive shorter comments than men (19.1 vs. 20.7 words, $p < .0001$), so comment length acts as a suppressor: adding it makes the controlled gap slightly larger (column 2). The mediation result is unchanged: with length controlled, adding comment tone and profanity still reduces the creator-gender coefficient to statistical non-significance (column 4). Comment tone (z) and Comment profanity (z) are standardized LIWC-22 Tone and Swear from audience comments. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table A12: Profanity mediation is robust to the distribution of the profanity measure

Transformation	a-path (gender $\rightarrow$ profanity)	b-path (profanity $\rightarrow$ views)	Indirect ($a \times b$)	Bootstrap 95% CI	Woman creator (full model)
Raw (canonical)	-0.429*** (0.019)	0.075*** (0.018)	-0.032	[-0.049, -0.018]	0.015 (0.033)
$\ln(1 + x)$	-0.524*** (0.019)	0.101*** (0.017)	-0.053	[-0.072, -0.036]	0.029 (0.034)
Inverse hyperbolic sine	-0.511*** (0.019)	0.097*** (0.017)	-0.050	[-0.068, -0.032]	0.027 (0.034)
Winsorized (p99)	-0.486*** (0.019)	0.090*** (0.017)	-0.044	[-0.061, -0.028]	0.023 (0.034)
Any profanity (binary)	-0.165*** (0.022)	0.085*** (0.016)	-0.014	[-0.021, -0.008]	0.011 (0.033)
Observations	9,097 in every model				

Notes: Robustness of the profanity (arousal) mediation pathway to the skewed distribution of channel-level comment profanity (skew = 4.33, kurtosis = 37.5). Each row re-estimates the mediation with the profanity measure transformed as labeled and then standardized; all coefficients are per one standard deviation of the transformed measure. The a-path regresses transformed profanity on the creator-gender indicator plus the five controls: channel age (days), $\ln$ (video count), HD quality, non-subscriber trailer, and $\ln$ (total comments); the b-path is the profanity coefficient in the full model of $\ln$ (total cumulative channel views) on creator gender, standardized comment tone, transformed profanity, and the five controls; robust standard errors in parentheses. Bootstrap 95% confidence intervals for the indirect effect are percentile intervals (500 replications); bias-corrected intervals are nearly identical and every interval excludes zero. The final column reports the creator-gender coefficient in the full model, which is statistically non-significant in every variant. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table A13: Content-language measures: descriptive statistics and pairwise correlations

<i>Panel A: Descriptive statistics</i>					
	N	Mean	SD	Min	Max
Title tone	8,502	49.671	33.430	1.000	99.000
Title profanity	9,230	0.070	0.421	0.000	10.340
Description tone	8,874	60.918	25.659	1.000	99.000
Description profanity	9,120	0.034	0.245	0.000	11.180
Comment tone	9,100	46.730	17.980	0.000	99.000
Comment profanity	9,100	0.346	0.539	0.000	9.235

<i>Panel B: Pairwise correlations</i>					
	(1)	(2)	(3)	(4)	(5)
1. Title tone					
2. Title profanity	-0.06				
3. Description tone	0.35	-0.04			
4. Description profanity	-0.03	0.41	-0.04		
5. Comment tone	0.28	-0.05	0.26	-0.05	
6. Comment profanity	-0.12	0.11	-0.10	0.09	-0.31

Notes: All measures are LIWC-22 Tone and Swear scores (channel-level means, unstandardized). Title and Description measures are computed on creator-written video titles and channel/video descriptions (creator-controlled content surfaces); Comment measures are computed on audience comments (audience treatment). N varies by surface owing to item-level missingness (titles N = 8,502; descriptions and comments as shown). Panel B reports pairwise correlations on pairwise-complete samples. Standardized versions of these measures enter the models in Table 3 and the content-channel appendix table.

Table A14: Listwise-deletion diagnostics: missingness on the comment-language mediators is unrelated to creator gender

Diagnostic	Value
Full sample (channels)	9,232
Mediator-complete analytic sample	9,100
Dropped to listwise deletion	132 (1.43%)
In-sample rate, men's channels	0.986
In-sample rate, women's channels	0.985
Gender vs. in-sample, p (unadjusted)	0.65
Gender vs. in-sample, p (five channel controls)	0.57

Notes: The analytic sample for the mediation models drops channels missing either comment-language mediator (LIWC Tone or LIWC Swear). The share dropped is small and the probability of being in-sample is unrelated to creator gender, both unconditionally and conditional on the five channel controls (channel age, log video count, HD quality, the non-subscriber-trailer indicator, and log total comments). The analytic sample therefore does not bias the gender comparison.

Table A15: Convergent validity of the LIWC comment-language measures against VADER sentiment

Measure pair	Pearson r	Spearman ρ	N
LIWC Tone vs. VADER compound (convergent, valence)	0.82	0.85	9,100
LIWC Tone vs. LIWC Swear (discriminant)	-0.31	-0.38	9,100

Notes: Channel-level correlations between the LIWC-22 comment-language measures and VADER, a rule-based sentiment instrument designed for social-media text, computed on the audience-comment corpus. LIWC Tone and VADER compound, an independent sentiment instrument, load on the same valence construct, supporting the use of LIWC Tone as a valence measure; LIWC Tone and LIWC Swear are empirically distinct, consistent with treating valence and arousal as separate dimensions. All reported correlations are significant at $p < 0.001$.

Table A16: Oaxaca-Blinder Decomposition of Gender Gap in Views

	(1) Topic Categories	(2) Channel Characteristics	(3) Combined
overall			
Mean, men	15.233*** (0.028)	15.233*** (0.028)	15.233*** (0.028)
Mean, women	14.825*** (0.037)	14.825*** (0.037)	14.825*** (0.037)
Difference	0.408*** (0.046)	0.407*** (0.046)	0.407*** (0.046)
Endowments (explained)	0.063 (0.047)	0.163*** (0.033)	0.248*** (0.046)
Coefficients (unexplained)	0.449*** (0.061)	0.243*** (0.033)	0.248*** (0.045)
Interaction	-0.105 (0.061)	0.001 (0.009)	-0.088* (0.043)
N	9230	9229	9229

Standard errors in parentheses

Dependent variable: $\ln(\text{total cumulative views})$.

Robust standard errors in parentheses.

Column (1) decomposes using 40 topic category indicators.

Column (2) uses channel characteristics.

Column (3) uses both topic categories and channel characteristics.

Endowments = portion attributable to differences in observable characteristics.

Coefficients = differential returns (unexplained component).

* $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$